Google DeepMind

Gemini 2.5: Pushing the Frontier with Advanced Reasoning, Multimodality, Long Context, and Next Generation Agentic Capabilities.

Gemini Team, Google

arXiv:2507.06261v6 [cs.CL] 19 Dec 2025

In this report, we introduce the Gemini 2.X model family: Gemini 2.5 Pro and Gemini 2.5 Flash, as well as our earlier Gemini 2.0 Flash and Flash-Lite models. Gemini 2.5 Pro is our most capable model yet, achieving SoTA performance on frontier coding and reasoning benchmarks. In addition to its incredible coding and reasoning skills, Gemini 2.5 Pro is a thinking model that excels at multimodal understanding and it is now able to process up to 3 hours of video content. Its unique combination of long context, multimodal and reasoning capabilities can be combined to unlock new agentic workflows. Gemini 2.5 Flash provides excellent reasoning abilities at a fraction of the compute and latency requirements and Gemini 2.0 Flash and Flash-Lite provide high performance at low latency and cost. Taken together, the Gemini 2.X model generation spans the full Pareto frontier of model capability vs cost, allowing users to explore the boundaries of what is possible with complex agentic problem solving.

1. Introduction

We present our latest family of natively multimodal models with advanced reasoning through thinking, long context and tool-use capabilities: Gemini 2.5 Pro and 2.5 Flash and our earlier Gemini 2.0 Flash and Gemini 2.0 Flash-Lite models. Together these form a new family of highly-capable models representing our next generation of AI models, designed to power a new era of agentic systems. Building upon the foundation of the Gemini 1.5 series (Gemini Team, 2024), this Gemini 2.X generation brings us closer to the vision of a universal AI assistant (Hassabis, 2025).

The Gemini 2.X series are all built to be natively multimodal, supporting long context inputs of >1 million tokens and have native tool use support. This allows them to comprehend vast datasets and handle complex problems from different information sources, including text, audio, images, video and even entire code repositories. These extensive capabilities can also be combined to build complex agentic systems, as happened in the case of Gemini Plays Pokémon[1] (Zhang, 2025). Different models in the series have different strengths and capabilities: (1) Gemini 2.5 Pro is our most intelligent thinking model, exhibiting strong reasoning and code capabilities. It excels at producing interactive web applications, is capable of codebase-level understanding and also exhibits emergent multimodal coding abilities. (2) Gemini 2.5 Flash is our hybrid reasoning model with a controllable thinking budget, and is useful for most complex tasks while also controlling the tradeoff between quality, cost, and latency. (3) Gemini 2.0 Flash is our fast and cost-efficient non-thinking model for everyday tasks and (4) Gemini 2.0 Flash-Lite is our fastest and most cost-efficient model, built for at-scale usage. A full comparison of the models in the Gemini 2.X model family is provided in Table 1. Taken together, the Gemini 2.X family of models cover the whole Pareto frontier of model capability vs cost, shifting it forward across a large variety of core capabilities, applications and use-cases, see Figure 1.

The Gemini 2.5 family of models maintain robust safety metrics while improving dramatically on

[1] Pokémon is a trademark of Nintendo Co., Ltd., Creatures Inc., and Game Freak Inc.

Please send correspondence to gemini-report@google.com.
© 2025 Google. All rights reserved

	Gemini 1.5 Flash	*Gemini 1.5 Pro*	**Gemini 2.0 Flash-Lite**	**Gemini 2.0 Flash**	**Gemini 2.5 Flash**	**Gemini 2.5 Pro**
Input modalities	Text, Image, Video, Audio	Text, Image, Video, Audio	Text, Image, Video, Audio	Text, Image, Video, Audio	Text, Image, Video, Audio	Text, Image, Video, Audio
Input length	1M	2M	1M	1M	1M	1M
Output modalities	Text	Text	Text	Text, Image*	Text, Audio*	Text, Audio*
Output length	8K	8K	8K	8K	64K	64K
Thinking	No	No	No	Yes*	Dynamic	Dynamic
Supports tool use?	No	No	No	Yes	Yes	Yes
Knowledge cutoff	November 2023	November 2023	June 2024	June 2024	January 2025	January 2025

Table 1 | Comparison of Gemini 2.X model family with Gemini 1.5 Pro and Flash. Tool use refers to the ability of the model to recognize and execute function calls (e.g., to perform web search, complete a math problem, execute code). **currently limited to Experimental or Preview, see Section 2.7. Information accurate as of publication date.*

helpfulness and general tone compared to their 2.0 and 1.5 counterparts. In practice, this means that the 2.5 models are substantially better at providing safe responses without interfering with important use cases or lecturing end users. We also evaluated Gemini 2.5 Pro's Critical Capabilities, including CBRN, cybersecurity, machine learning R&D, and deceptive alignment. While Gemini 2.5 Pro showed a significant increase in some capabilities compared to previous Gemini models, it did not reach any of the Critical Capability Levels in any area.

Our report is structured as follows: we begin by briefly describing advances we have made in model architecture, training and serving since the release of the Gemini 1.5 model. We then showcase the performance of the Gemini 2.5 models, including qualitative demonstrations of its abilities. We conclude by discussing the safety evaluations and implications of this model series.

2. Model Architecture, Training and Dataset

2.1. Model Architecture

The Gemini 2.5 models are sparse mixture-of-experts (MoE) (Clark et al., 2022; Du et al., 2021; Fedus et al., 2021; Jiang et al., 2024; Lepikhin et al., 2020; Riquelme et al., 2021; Roller et al., 2021; Shazeer et al., 2017) transformers (Vaswani et al., 2017) with native multimodal support for text, vision, and audio inputs. Sparse MoE models activate a subset of model parameters per input token by learning to dynamically route tokens to a subset of parameters (experts); this allows them to decouple total model capacity from computation and serving cost per token. Developments to the model architecture contribute to the significantly improved performance of Gemini 2.5 compared to Gemini 1.5 Pro (see Section 3). Despite their overwhelming success, large transformers and sparse MoE models are known to suffer from training instabilities (Chowdhery et al., 2022; Dehghani et al., 2023; Fedus et al., 2021; Lepikhin et al., 2020; Liu et al., 2020; Molybog et al., 2023; Wortsman et al., 2023; Zhai et al., 2023; Zhang et al., 2022). The Gemini 2.5 model series makes considerable progress in enhancing large-scale training stability, signal propagation and optimization dynamics, resulting in a considerable boost in performance straight out of pre-training compared to previous Gemini models.

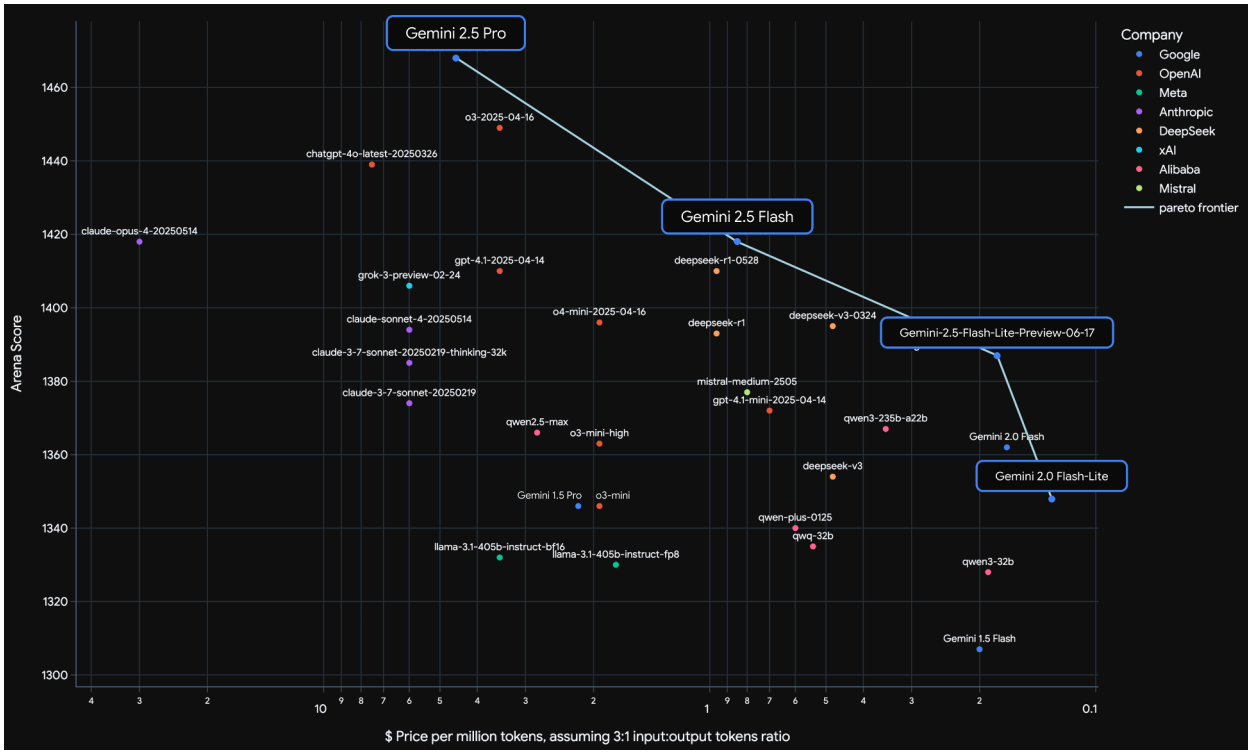


Figure 1 | Cost-performance plot. Gemini 2.5 Pro is a marked improvement over Gemini 1.5 Pro, and has an LMArena score that is over 120 points higher than Gemini 1.5 Pro. Cost is a weighted average of input and output tokens pricing per million tokens. Source: LMArena, imported on 2025-06-16.

Gemini 2.5 models build on the success of Gemini 1.5 in processing long-context queries, and incorporate new modeling advances allowing Gemini 2.5 Pro to surpass the performance of Gemini 1.5 Pro in processing long context input sequences of up to 1M tokens (see Table 3). Both Gemini 2.5 Pro and Gemini 2.5 Flash can process pieces of long-form text (such as the entirety of “Moby Dick” or “Don Quixote”), whole codebases, and long form audio and video data (see Appendix 8.5). Together with advancements in long-context abilities, architectural changes to Gemini 2.5 vision processing lead to a considerable improvement in image and video understanding capabilities, including being able to process 3-hour-long videos and the ability to convert demonstrative videos into interactive coding applications (see our recent blog post by Baddepudi et al., 2025).

The smaller models in the Gemini 2.5 series — Flash size and below — use distillation (Anil et al., 2018; Hinton et al., 2015), as was done in the Gemini 1.5 series (Gemini Team, 2024). To reduce the cost associated with storing the teacher’s next token prediction distribution, we approximate it using a k-sparse distribution over the vocabulary. While this still increases training data throughput and storage demands by a factor of k, we find this to be a worthwhile trade-off given the significant quality improvement distillation has on our smaller models, leading to high-quality models with a reduced serving cost (see Figure 2).

2.2. Dataset

Our pre-training dataset is a large-scale, diverse collection of data encompassing a wide range of domains and modalities, which includes publicly available web documents, code (various programming languages), images, audio (including speech and other audio types) and video, with a cutoff date of June 2024 for 2.0 and January 2025 for 2.5. Compared to the Gemini 1.5 pre-training dataset

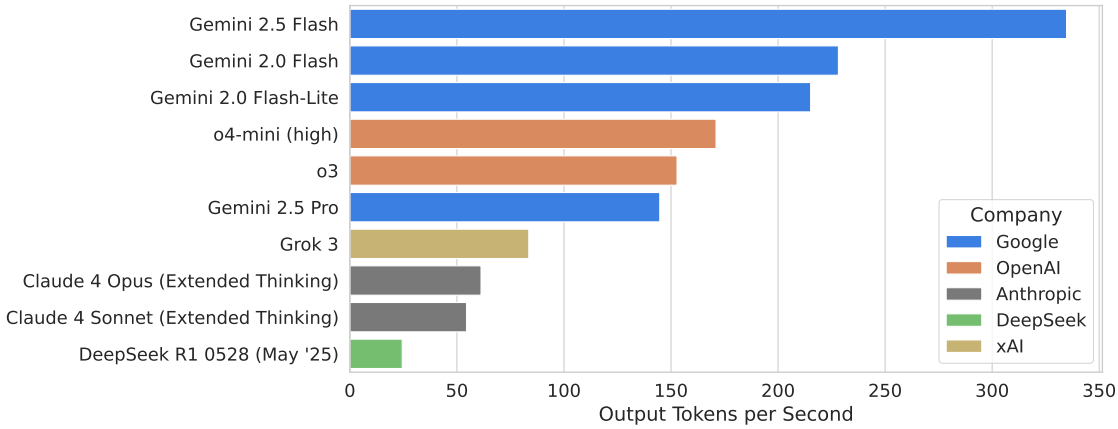


Figure 2 | Number of output tokens generated per second (after the first chunk has been received from the API) for different models. Source: `ArtificialAnalysis.ai`, imported on 2025-06-15.

we also utilized new methods for improved data quality for both filtering, and deduplication. Our post-training dataset, like Gemini 1.5, consists of instruction tuning data that is carefully collected and vetted. It is a collection of multimodal data with paired instructions and responses, in addition to human preference and tool-use data.

2.3. Training Infrastructure

This model family is the first to be trained on TPUv5p architecture. We employed synchronous data-parallel training to parallelise over multiple 8960-chip pods of Google's TPUv5p accelerators, distributed across multiple datacenters.

The main advances in software pre-training infrastructure compared with Gemini 1.5 were related to elasticity and mitigation of SDC (Silent Data Corruption) errors:

1. **Slice-Granularity Elasticity**: Our system now automatically continues training with fewer "slices" of TPU chips when there is a localized failure, and this reconfiguration results in tens of seconds of lost training time per interruption, compared with the 10 or more minute delay waiting for healthy machines to be rescheduled without elasticity; the system continues training at around 97% throughput while the failed slice is recovering. At the scale of this training run we see interruptions from hardware failures multiple times per hour, but our fault tolerance machinery is designed to tolerate the higher failure rates expected at much larger scales.
2. **Split-Phase SDC Detection**: On previous large-scale runs it could take many hours to detect and localize machines with SDC errors, requiring both downtime while debugging, and roll-back/replay of a large number of potentially corrupt training steps. We now use lightweight deterministic replay to immediately repeat any step with suspicious metrics, and compare per-device intermediate checksums to localize the root cause of any data corruption. Empirically, accelerators that start to exhibit intermittent SDCs are identified within a few minutes, and quickly excluded from the job. During this run, around 0.25% of steps were replayed due to suspected SDCs and 6% of these replays turned out to be genuine hardware corruption.

Both of the above techniques were relatively simple to implement due to the single-controller design of the Pathways system (Barham et al., 2022), which allows all accelerators to be coordinated from a single python program with a global view of the system state. The controller can make use of

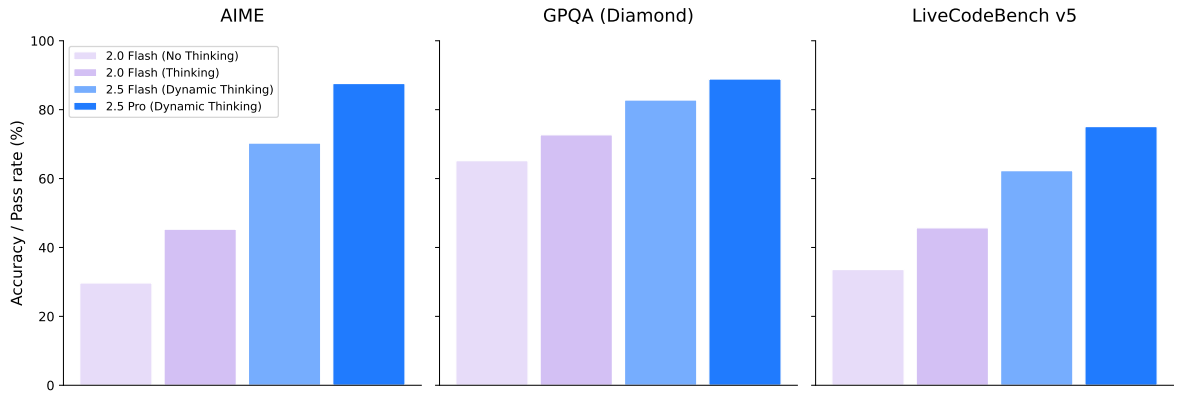


Figure 3 | Impact of "Thinking" on Gemini's performance on AIME 2025 (Balunović et al., 2025), LiveCodeBench (corresponding to 10/05/2024 - 01/04/2025 in the UI) (Jain et al., 2024) and GPQA diamond (Rein et al., 2024) benchmarks.

parallel 'remote python' operations on TPU workers to monitor training metrics, track performance stragglers, and root-cause SDC errors.

Overall during the run, 93.4% of the time was spent performing TPU computations; the remainder was approximately spent half in elastic reconfigurations, and half in rare tail cases where elasticity failed. Around 4.5% of the computed steps were replays or rollbacks for model debugging interventions.

2.4. Post-training

Since the initial announcement of Gemini 1.5, significant advancements have been made in our post-training methodologies, driven by a consistent focus on data quality across the Supervised Fine-Tuning (SFT), Reward Modeling (RM), and Reinforcement Learning (RL) stages. A key focus has been leveraging the model itself to assist in these processes, enabling more efficient and nuanced quality control.

Furthermore, we have increased the training compute allocated to RL, allowing deeper exploration and refinement of model behaviors. This has been coupled with a focus on verifiable rewards and model-based generative rewards to provide more sophisticated and scalable feedback signals. Algorithmic changes to the RL process have also improved stability during longer training. These advancements have enabled Gemini 2.5 to learn from more diverse and complex RL environments, including those requiring multi-step actions and tool use. The combination of these improvements in data quality, increased compute, algorithmic enhancements, and expanded capabilities has contributed to across-the-board performance gains (as described in Section 3) , notably reflected in the significant increase in the model's LMArena Elo scores, with both Gemini 2.5 Flash and Pro gaining more than 110 points over their Gemini 1.5 counterparts (122 for Gemini 2.5 Pro and 111 for Gemini 2.5 Flash, see Figure 1), along with significant improvements on several other frontier benchmarks.

2.5. Thinking

Past Gemini models produce an answer immediately following a user query. This constrains the amount of inference-time compute (Thinking) that our models can spend reasoning over a problem. Gemini Thinking models are trained with Reinforcement Learning to use additional compute at inference time to arrive at more accurate answers. The resulting models are able to spend tens of

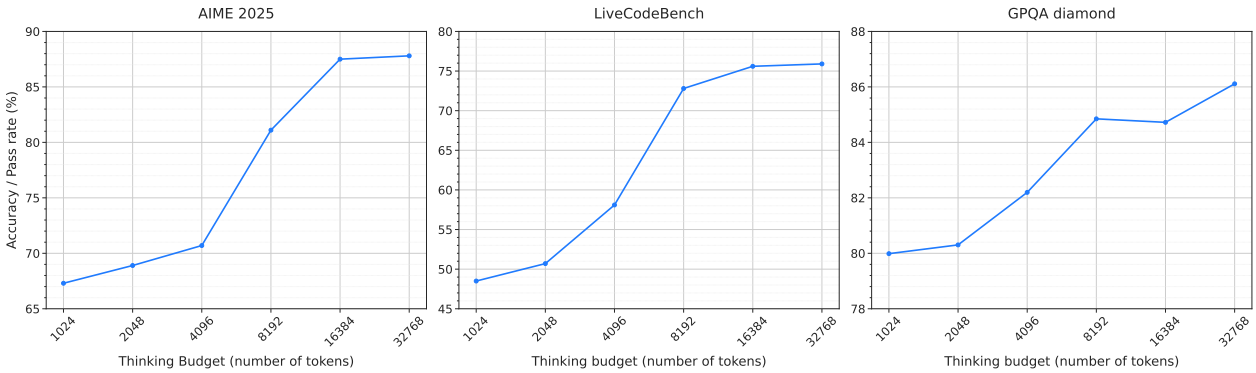


Figure 4 | Impact of thinking budget on performance on AIME 2025 (Balunović et al., 2025), LiveCodeBench (corresponding to 10/05/2024 - 01/04/2025 in the UI) (Jain et al., 2024) and GPQA diamond (Rein et al., 2024) benchmarks.

thousands of forward passes during a "thinking" stage, before responding to a question or query.

Our training recipe has evolved from the original experimental thinking model, Gemini 2.0 Flash Thinking (launched in December 2024), to the Gemini 2.5 Thinking series, which incorporates Thinking natively across all domains. The result is a single model that can achieve stronger reasoning performance across the board, and is able to scale up its performance further as a function of inference time (see Figure 3 for an example of the impact of Thinking).

We integrated Thinking with other Gemini capabilities, including native multimodal inputs (images, text, video, audio) and long context (1M+ tokens). For any of these capabilities, the model decides for itself how long to think before providing an answer. We also provide the ability to set a Thinking budget, constraining the model to respond within a desired number of tokens. This allows users to trade off performance with cost. To demonstrate this capability, we conducted experiments where we systematically varied the thinking budget, measured in the number of tokens the model is allowed to use for internal computation. As shown in Figure 4, increasing this budget allows the model to scale its performance and achieve significantly higher accuracy.

2.6. Capability-specific improvements

While most of the changes made to our training architecture and recipe since Gemini 1.5 have resulted in improvements across all capabilities, we have also made changes that have resulted in some capability-specific wins. We will now discuss these for code, factuality, long context, multilinguality, audio, video, and agentic use cases (with a particular focus on Gemini Deep Research).

Code

Gemini 2.0 and 2.5 represent a strategic shift of our development priorities towards delivering tangible real-world value, empowering users to address practical challenges and achieve development objectives within today's complex, multimodal software environments. To realize this, concerted efforts have been undertaken across both pre-training and post-training phases since Gemini 1.5. In pre-training, we intensified our focus on incorporating a greater volume and diversity of code data from both repository and web sources into the training mixture. This has rapidly expanded coverage and enabled the development of more compute-efficient models. Furthermore, we have substantially enhanced our suite of evaluation metrics for assessing code capabilities aligned with downstream use cases, alongside improving our ability to accurately predict model performance.

During post-training, we developed novel training techniques incorporating reasoning capabilities and curated a diverse set of engineering tasks, with the aim to equip Gemini with effective problem-solving skills crucial for addressing modern engineering challenges. Key applications demonstrating these advancements include IDE functionalities, code agent use cases for complex, multi-step operations within full repositories, and multimodal, interactive scenarios such as end-to-end web and mobile application development. Collectively, these efforts have yielded broad and significant improvements in Gemini's coding capabilities. This progress is evidenced by superior performance on established benchmarks: performance on LiveCodeBench (Jain et al., 2024) increased from 30.5% for Gemini 1.5 Pro to 74.2% for Gemini 2.5 Pro, while that for Aider Polyglot (Gauthier, 2025) went from 16.9% to 82.2%. Performance on SWEBench-verified (Chowdhury et al., 2024; Jimenez et al., 2024) went from 34.2% to 67.2%, see Table 3 and Figure 5 in Section 3.2. Furthermore, Gemini 2.5 Pro obtained an increase of over 500 Elo over Gemini 1.5 Pro on the LMArena WebDev Arena (Chiang et al., 2024; LMArena Team, 2025), resulting in meaningful enhancements in practical applications, including UI and web application development (Doshi, 2025a), and the creation of sophisticated agentic workflows (Kilpatrick, 2025).

Factuality

Within the context of generative models, ensuring the factuality of model responses to information-seeking prompts remains a core pillar of Gemini model development. With Gemini 1.5, our research was concentrated on enhancing the model's world knowledge and its ability to provide answers faithfully grounded in the context provided within the prompt. This effort culminated in the December 2024 release of FACTS Grounding (Jacovi et al., 2025), now an industry-standard benchmark for evaluating an LLM's capacity to generate responses grounded in user-provided documents. With Gemini 2.0 and 2.5, we have significantly expanded our scope to address multimodal inputs, long-context reasoning, and model-retrieved information. At the same time, the landscape and user expectations for factuality have evolved dramatically, shaped in part by Google's deployment of AI Overviews and AI Mode (Stein, 2025). To meet these demands, Gemini 2.0 marked a significant leap as our first model family trained to natively call tools like Google Search, enabling it to formulate precise queries and synthesize fresh information with sources. Building on this, Gemini 2.5 integrates advanced reasoning, allowing it to interleave these search capabilities with internal thought processes to answer complex, multi-hop queries and execute long-horizon tasks. The model has learned to use search and other tools, reason about the outputs, and issue additional, detailed follow-up queries to expand the information available to it and to verify the factual accuracy of the response. Our latest models now power the experiences of over 1.5B monthly active users in Google's AI Overviews and 400M users in the Gemini App. These models exhibit state-of-the-art performance across a suite of factuality benchmarks, including SimpleQA for parametric knowledge (Wei et al., 2024), FACTS Grounding for faithfulness to provided documents (Jacovi et al., 2024, 2025), and the Vectara Hallucination Leaderboard (Hughes et al., 2023), cementing Gemini as the model of choice for information-seeking demands.

Long context

Modeling and data advances helped us improve the quality of our models' responses to queries utilizing our one million-length context window, and we reworked our internal evaluations to be more challenging to help steer our modeling research. When hill-climbing, we targeted challenging retrieval tasks (like LOFT of Lee et al., 2024), long-context reasoning tasks (like MRCR-V2 of Vodrahalli et al., 2024), and multimodal tasks (like VideoMME of Fu et al., 2025). According to the results in Table 6, the new 2.5 models improve greatly over previous Gemini 1.5 models and achieve state-of-the-art quality on all of those. An example showcasing these improved capabilities for video recall can be

seen in Appendix 8.5, where Gemini 2.5 Pro is able to consistently recall a 1 second visual event out of a full 46-minute video.[2]

Multilinguality

Gemini's multilingual capabilities have also undergone a profound evolution since 1.5, which already encompassed over 400 languages via pretraining. This transformation stems from a holistic strategy, meticulously refining pre- and post-training data quality, advancing tokenization techniques, innovating core modeling, and executing targeted capability hillclimbing. The impact is particularly striking in Indic and Chinese, Japanese and Korean languages, where dedicated optimizations in data quality and evaluation have unlocked dramatic gains in both quality and decoding speed. Consequently, users benefit from significantly enhanced language adherence, responses designed to faithfully respect the requested output language, and a robust improvement in generative quality and factuality across languages, solidifying Gemini's reliability across diverse linguistic contexts.

Audio

While Gemini 1.5 was focused on native audio understanding tasks such as transcription, translation, summarization and question-answering, in addition to understanding, Gemini 2.5 was trained to perform audio generation tasks such as text-to-speech or native audio-visual to audio out dialog. To enable low-latency streaming dialog, we incorporated causal audio representations that also allow streaming audio into and out of Gemini 2.5. These capabilities derive from an increased amount of pre-training data spanning over 200 languages, and development of improved post-training recipes. Finally, through our improved post-training recipes, we have integrated advanced capabilities such as thinking, affective dialog, contextual awareness and tool use into Gemini's native audio models.

Video

We have significantly expanded both our pretraining and post-training video understanding data, improving the audio-visual and temporal understanding capabilities of the model. We have also trained our models so that they perform competitively with 66 instead of 258 visual tokens per frame, enabling using about 3 hours of video instead of 1h within a 1M tokens context window[3]. Two new applications that were not previously possible, but that have been unlocked as a result of these changes are: creating an interactive app from a video (such as a quiz to test students' understanding of the video content) and creating a p5.js animation to show the key concepts from the video. Our recent blog post (Baddepudi et al., 2025) shows examples of these applications.

Gemini as an Agent: Deep Research

Gemini Deep Research (Gemini Team, Google, 2024) is an agent built on top of the Gemini 2.5 Pro model designed to strategically browse the web and provide informed answers to even the most niche user queries. The agent is optimized to perform task prioritization, and is also able to identify when it reaches a dead-end when browsing. We have massively improved the capabilities of Gemini Deep Research since its initial launch in December 2024. As evidence of that, performance of Gemini Deep Research on the Humanity's Last Exam benchmark (Phan et al., 2025) has gone from 7.95% in December 2024 to the **SoTA score of 26.9% and 32.4% with higher compute** (June 2025).

[2]For further discussion on long context capabilities, challenges, and future outlook, the Release Notes podcast episode "Deep Dive into Long Context" provides additional insights and discussion: https://youtu.be/NHMJ9mqKeMQ.

[3]This is referred to as low media resolution in the API: https://ai.google.dev/api/generate-content#MediaResolution.

2.7. The path to Gemini 2.5

On the way to Gemini 2.5 Pro, we experimented with our training recipe, and tested a small number of these experimental models with users. We have already discussed Gemini 2.0 Flash Thinking (see Section 2.5). We will now discuss some of the other models briefly.

Gemini 2.0 Pro

In February 2025, we released an experimental version of Gemini 2.0 Pro. At the time, it had the strongest coding performance of any model in the Gemini model family, as well as the best understanding and world knowledge. It also came with our largest context window at 2 million tokens, which enabled it to comprehensively analyze and understand vast amounts of information. For further information about Gemini 2.0 Pro, please see our earlier blog posts (Kavukcuoglu, 2025; Mallick and Kilpatrick, 2025).

Gemini 2.0 Flash Native Image Generation Model

In March 2025, we released an experimental version of Gemini 2.0 Flash Native Image Generation. It has brought to the users new capabilities as a result of a strong integration between the Gemini model and image-generation capabilities, enabling new experiences related to image generation & image editing via natural-language prompting. Capabilities such as multi-step conversational editing or interleaved text-image generation are very natural in such a setting, and horizontal transfer related to multi-language coverage immediately allowed such experiences to happen across all the languages supported by the Gemini models. Native image generation turns Gemini into a multimodal creation partner and enables Gemini to express ideas through both text and images, and to seamlessly move between the two. For further information about Gemini 2.0 Flash Native Image Generation, please see our earlier blog posts (Kampf and Brichtova, 2025; Sharon, 2025)

Gemini 2.5 Audio Generation

With Gemini 2.5, the Controllable TTS and Native Audio Dialog capabilities are available as separate options on AI Studio (Generate Media and Stream sections respectively). Our Gemini 2.5 Preview TTS Pro and Flash models support more than 80 languages with the speech style controlled by a free formatted prompt which can specify style, emotion, pace, etc, while also being capable of following finer-grained steering instructions specified in the transcript. Notably, Gemini 2.5 Preview TTS can generate speech with multiple speakers, which enables the creation of podcasts as used in NotebookLM Audio Overviews (Wang, 2024). Our Gemini 2.5 Flash Preview Native Audio Dialog model uses native audio generation, which enables the same level of style, pacing and accent control as available in our controllable TTS offering. Our dialog model supports tool use and function calling, and is available in more than 24 languages. With native audio understanding and generation capabilities, it can understand and respond appropriately to the user's tone. This model is also capable of understanding when to respond to the user, and when not to respond, ignoring background and non-device directed audio. Finally, we also offer an advanced 'Thinking' variant that effectively handles more complex queries and provides more robust and reasoned responses in exchange for some additional latency.

Gemini 2.5 Flash-Lite

In June 2025, we released an experimental version of Gemini 2.5 Flash-Lite (`gemini-2.5-flash-lite-preview-06-17`). It comes with the same capabilities that make Gemini 2.5 helpful, including the ability to turn thinking on at different budgets, connecting to tools like Google Search and code

execution, support for multimodal inputs and a 1 million-token context length. Our goal was to provide an economical model class which provides ultra-low-latency capabilities and high throughput per dollar, echoing the initial release of 2.0 Flash-Lite (Google DeepMind, 2025b; Mallick and Kilpatrick, 2025).

Gemini 2.5 Pro Deep Think

To advance Gemini's capabilities towards solving hard reasoning problems, we developed a novel reasoning approach, called Deep Think, that naturally blends in parallel thinking techniques during response generation. Deep Think enables Gemini to creatively produce multiple hypotheses and carefully critique them before arriving at the final answer, achieving state-of-the-art performances in challenging benchmarks such as Olympiad math (USAMO 2025), competitive coding (LiveCodeBench), and multimodality (MMMU), see more details at (Doshi, 2025b). We announced Gemini 2.5 Deep Think at Google I/O and launched an experimental version to trusted testers and advanced users in June 2025.

3. Quantitative evaluation

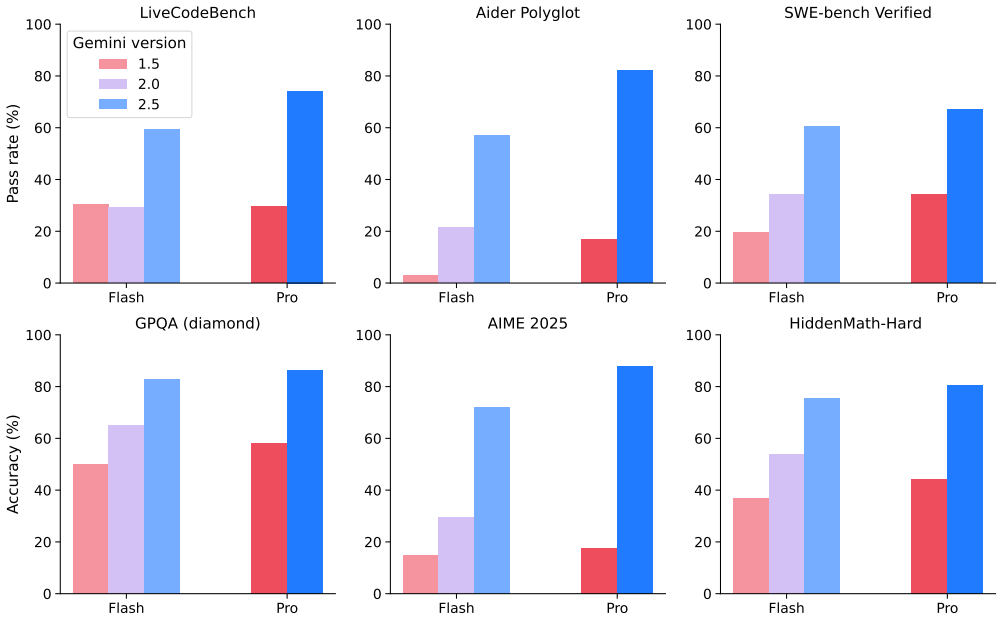


Figure 5 | Performance of Gemini 2.X models at coding, math and reasoning tasks in comparison to previous Gemini models. SWE-bench verified numbers correspond to the "multiple attempts" setting reported in Table 3.

We will now examine the performance of the Gemini 2.X model family across a wide range of benchmarks. We will first compare the performance of the Gemini 2.X models to the earlier Gemini 1.5 Pro and Flash models, before we compare the performance of Gemini 2.5 Pro to other available large language models.

With web-scale pre-training of AI models, coupled with the post-training techniques that allow policy and reward models to leverage public benchmarks, avoiding leaks and biases in the data used for pre- and post-training is a persistent challenge. In the development of the Gemini 2.5 series, in addition to the standard n-gram based decontamination we used in Gemini 1.5, we also employed semantic-similarity and model based decontamination procedures to help mitigate evaluation set leakage. To move beyond the reliance on training set decontamination, we also continue reporting on internally developed non-public benchmarks, such as HiddenMath.

Model	AI Studio model ID
Gemini 1.5 Flash	`gemini-1.5-flash-002`
Gemini 1.5 Pro	`gemini-1.5-pro-002`
Gemini 2.0 Flash-Lite	`gemini-2.0-flash-lite-001`
Gemini 2.0 Flash	`gemini-2.0-flash-001`
Gemini 2.5 Flash	`gemini-2.5-flash`
Gemini 2.5 Pro	`gemini-2.5-pro`

Table 2 | Mapping of Gemini model names to AI Studio API model IDs.

3.1. Methodology

In Table 3, we compare the performance of Gemini 2.5 models to the Gemini 1.5 models, while in Table 4, we compare the performance of Gemini 2.5 Pro to that of other large language models.

Gemini results: All Gemini scores are pass@1, and are "single attempt" settings unless otherwise specified. In the "single attempt" setting, no majority voting or parallel test-time compute is permitted, while in the "multiple attempts" setting, test-time selection of the candidate answer is allowed. All Gemini evaluations are run with the AI Studio API for the model id that we provide in Table 2, with default sampling settings. To reduce variance, we average over multiple trials for smaller benchmarks. Aider Polyglot scores are the pass rate average of 3 trials. Vibe-Eval results are reported using Gemini as a judge.

Non-Gemini results: All the results for non-Gemini models are sourced from providers' self reported numbers unless mentioned otherwise. All "SWE-bench Verified" numbers follow official provider reports, which means that they are computed using different scaffoldings and infrastructure, and aren't directly comparable.

For some evaluations, we obtain results from the external leaderboards that report results on these benchmarks. Results for Humanity's Last Exam results are sourced from Scale's leaderboard and results for DeepSeek are obtained from the text-only variant of the leaderboard (indicated with a $\diamond$ in Table 4). For Gemini 2.0 models, the reported results are on an earlier HLE dataset (indicated with a † in Table 3). Results on LiveCodeBench results are taken from (1/1/2025 - 5/1/2025) in the UI. Aider Polyglot numbers come from the Aider leaderboard and results for SimpleQA come from this repo where available. Results on FACTS Grounding come from Kaggle. In the case of LOFT and MRCR-V2, we report results on both the 128k context length variant, as well as the 1M context length variant. In the 128k context length variant, we measure performance on contexts up to 128k, while for the 1M context length variant, we report performance on context lengths of exactly 1M.

More details on all benchmarks, including subsets and how scores were obtained can be found in Table 11 in Appendix 8.1.

3.2. Core capability quantitative results

As can be seen in Table 3, and Figure 5, the Gemini 2.5 models excel at coding tasks such as LiveCodeBench, Aider Polyglot and SWE-bench Verified, and represent a marked improvement over previous models.

In addition to coding performance, Gemini 2.5 models are noticeably better at math and reasoning tasks than Gemini 1.5 models: performance on AIME 2025 is 88.0% for Gemini 2.5 Pro compared to 17.5% for Gemini 1.5 Pro, while performance on GPQA (diamond) went from 58.1% for Gemini 1.5 Pro to 86.4%. Performance on image understanding tasks has also increased significantly.

It is also interesting to note that the Gemini 2.5 Flash model has become the second most capable model in the Gemini family, and has overtaken not just previous Flash models, but also the Gemini 1.5 Pro model released one year ago.

Capability	Benchmark		Gemini 1.5 Flash	Gemini 1.5 Pro	Gemini 2.0 Flash-Lite	Gemini 2.0 Flash	Gemini 2.5 Flash	Gemini 2.5 Pro
Code	LiveCodeBench		30.3%	29.7%	29.1%	29.1%	59.3%	**74.2%**
	Aider Polyglot		2.8%	16.9%	10.5%	21.3%	56.7%	**82.2%**
	SWE-bench Verified	*single attempt*	9.6%	22.3%	12.5%	21.4%	48.9%	**59.6%**
		multiple attempts	19.7%	34.2%	23.1%	34.2%	60.3%	**67.2%**
Reasoning	GPQA (diamond)		50.0%	58.1%	50.5%	65.2%	82.8%	**86.4%**
	Humanity's Last Exam	*no tools*	-	4.6%	4.6% †	5.1% †	11.0%	**21.6%**
Factuality	SimpleQA		8.6%	24.9%	16.5%	29.9%	26.9%	**54.0%**
	FACTS Grounding		82.9%	80.0%	82.4%	84.6%	85.3%	**87.8%**
Multilinguality	Global MMLU (Lite)		72.5%	80.8%	78.0%	83.4%	88.4%	**89.2%**
	ECLeKTic		16.4%	27.0%	27.7%	33.6%	36.8%	**46.8%**
Math	AIME 2025		14.7%	17.5%	23.8%	29.7%	72.0%	**88.0%**
	HiddenMath-Hard		36.8%	44.3%	47.4%	53.7%	75.5%	**80.5%**
Long-context	LOFT (hard retrieval)	*≤128K*	67.3%	75.9%	50.7%	58.0%	82.1%	**87.0%**
		1M	36.7%	47.1%	7.6%	7.6%	58.9%	**69.8%**
	MRCR-V2 (8-needle)	*≤128K*	18.4%	26.2%	11.6%	19.0%	54.3%	**58.0%**
		1M	10.2%	12.1%	4.0%	5.3%	**21.0%**	16.4%
Image Understanding	MMMU		58.3%	67.7%	65.1%	69.3%	79.7%	**82.0%**
	Vibe-Eval (Reka)		52.3%	55.9%	51.5%	55.4%	65.4%	**67.2%**
	ZeroBench		0.5%	1.0%	0.75%	1.25%	2.0%	**4.5%**
	BetterChartQA		59.0%	65.8%	52.3%	57.8%	67.3%	**72.4%**

Table 3 | Evaluation of Gemini 2.5 family across a wide range of core capability benchmarks and in comparison to Gemini 1.5 models. Please see Tables 5 and 6 for audio and video evaluations. See Table 11 Appendix 8.1 for benchmarks and evaluation details.

3.3. Evaluation of Gemini 2.5 Pro against other large language models

Relative to other large language models that are available (see Table 4), Gemini achieves the highest score on the Aider Polyglot coding task, Humanity's Last Exam, GPQA (diamond), and on the SimpleQA and FACTS Grounding factuality benchmarks out of all of the models examined here. Gemini also continues to stand out for achieving the SoTA score on both the LOFT and MRCR long-context tasks at 128k context, and is the only one, amongst the models examined in the above table, to support context lengths of 1M+ tokens.

Not all of the models shown in Table 4 have native support for multimodal inputs. As such, we compare against a different set of models for audio and video understanding.

Audio Understanding

In Table 5, we showcase the performance of the Gemini 2.5 model family at audio understanding, and compare the performance of these models to earlier Gemini models, as well as to GPT models. Gemini 2.5 Pro demonstrates state-of-the-art audio understanding performance as measured by public benchmarks for ASR and AST, and compares favorably to alternatives under comparable testing conditions (using the same prompts and inputs).

Video Understanding

In Table 6, we show the performance of Gemini 2.5 models at video understanding. As can be seen, Gemini 2.5 Pro achieves state-of-the-art performance on key video understanding benchmarks, surpassing recent models like GPT 4.1 under comparable testing conditions (same prompt and video

Capability	**Benchmark**		**Gemini 2.5 Pro**	**o3** *high*	**o4-mini** *high*	**Claude 4 Sonnet**	**Claude 4 Opus**	**Grok 3 Beta** *Extended Thinking*	**DeepSeek R1 0528**
Code	LiveCodeBench		74.2%	72.0%	**75.8%**	48.9%	51.1%	–	70.5%
	Aider Polyglot		**82.2%**	79.6%	72.0%	61.3%	72.0%	53.3%	71.6%
	SWE-bench Verified	*single attempt*	59.6%	69.1%	68.1%	**72.7%**	72.5%	-	-
		multiple attempts	67.2%	-	-	**80.2%**	79.4%	-	57.6%
Reasoning	GPQA (diamond)	*single attempt*	**86.4%**	83.3%	81.4%	75.4%	79.6%	80.2%	81.0%
	Humanity's Last Exam	*no tools*	**21.6%**	20.3%	18.1%	7.8%	10.7%	-	14.0% ⋄
Factuality	SimpleQA		**54.0%**	48.6%	19.3%	-	-	43.6%	27.8%
	FACTS Grounding		**87.8%**	69.9%	62.1%	79.1%	77.7%	74.8%	82.4%
Math	AIME 2025	*single attempt*	88.0%	88.9%	**92.7%**	70.5%	75.5%	77.3%	87.5%
Long-context	LOFT (hard retrieval)	*≤128K*	**87.0%**	77.0%	60.5%	81.6%	-	73.1%	-
		1M	**69.8%**	-	-	-	-	-	-
	MRCR-V2 (8-needle)	*≤128K*	**58.0%**	57.1%	36.3%	39.1%	*16.1%**	34.0%	-
		1M	**16.4%**	-	-	-	-	-	-
Image Understanding	MMMU	*single attempt*	82.0%	**82.9%**	81.6%	74.4%	76.5%	76.0%	*No MM support*

Table 4 | Performance comparison of Gemini 2.5 Pro with other large language models on different capabilities. Please see Tables 5 and 6 for audio and video evaluations. See Table 11 for benchmarks and evaluation details. **: with no thinking and API refusals*

Benchmark	Gemini 1.5 Flash	Gemini 1.5 Pro	Gemini 2.0 Flash-Lite	Gemini 2.0 Flash	Gemini 2.5 Flash	Gemini 2.5 Pro	GPT-4o mini Audio Preview	GPT 4o Audio Preview	GPT 4o transcribe
FLEURS *(53 lang, WER ↓)*	12.71	7.14	9.60	9.04	9.95	**6.66**	19.52	12.16	8.17
CoVoST2 *(21 lang, BLEU ↑)*	34.81	37.53	34.74	36.35	36.15	**38.48**	29.5	35.89	–

Table 5 | Performance comparison of Gemini 2.5 models to earlier Gemini models, as well as to GPT models for audio understanding. Note that for GPT models, metrics may differ from those previously reported due to differing eval methodologies. See Table 11 for benchmarks and evaluation details.

frames). For cost-sensitive applications, Gemini 2.5 Flash provides a highly competitive alternative.

Modalities	Benchmark	Gemini 1.5 Flash	Gemini 1.5 Pro	Gemini 2.0 Flash-Lite	Gemini 2.0 Flash	Gemini 2.5 Flash	Gemini 2.5 Pro	OpenAI GPT 4.1
visual-only	ActivityNet-QA	56.2	57.3	55.3	56.4	65.1	**66.7**	60.4
	EgoTempo	34.5	36.3	30.1	39.3	36.7	**44.3**	40.3
	Perception Test	66.5	69.4	67.5	68.8	75.1	**78.4**	64.8
	QVHighlights	64.4	68.7	25.7	63.9	52.4	**75.0**	71.4
	VideoMMMU	64.8	70.4	64.3	68.5	79.2	**83.6**	60.9
	1H-VideoQA	61.9	72.2	55.6	67.5	67.5	**81.0**	56.8
audio + visual	LVBench	61.9	65.7	52	61.8	62.7	**78.7**	63.4
	VideoMME	70.4	73.2	62.1	72.8	75.5	**84.3**	72.0
	VATEX	56.9	55.5	58.5	56.9	65.2	**71.3**	64.1
	VATEX-ZH	46.2	52.2	43.2	48.5	43.9	**59.7**	48.7
	YouCook2 Cap	153.2	170.0	78.6	129.0	177.6	**188.3**	127.6
visual + subtitles	Minerva	49.6	52.8	46.8	52.4	60.7	**67.6**	54.0
	Neptune	78.7	82.7	81.5	83.1	84.3	**87.3**	85.2
audio+visual+ subtitles	VideoMME	77.3	79.8	72.5	78.8	81.5	**86.9**	79.6

Table 6 | Evaluation of Gemini 2.5 vs. prior models and GPT 4.1 on video understanding benchmarks. Performance is measured by string-match accuracy for multiple-choice VideoQA, LLM-based accuracy for open-ended VideoQA, R1@0.5 for moment retrieval and CIDEr for captioning. See Table 11 for benchmarks and evaluation details.

4. Example use cases of Gemini 2.5 Pro

4.1. Gemini Plays Pokémon

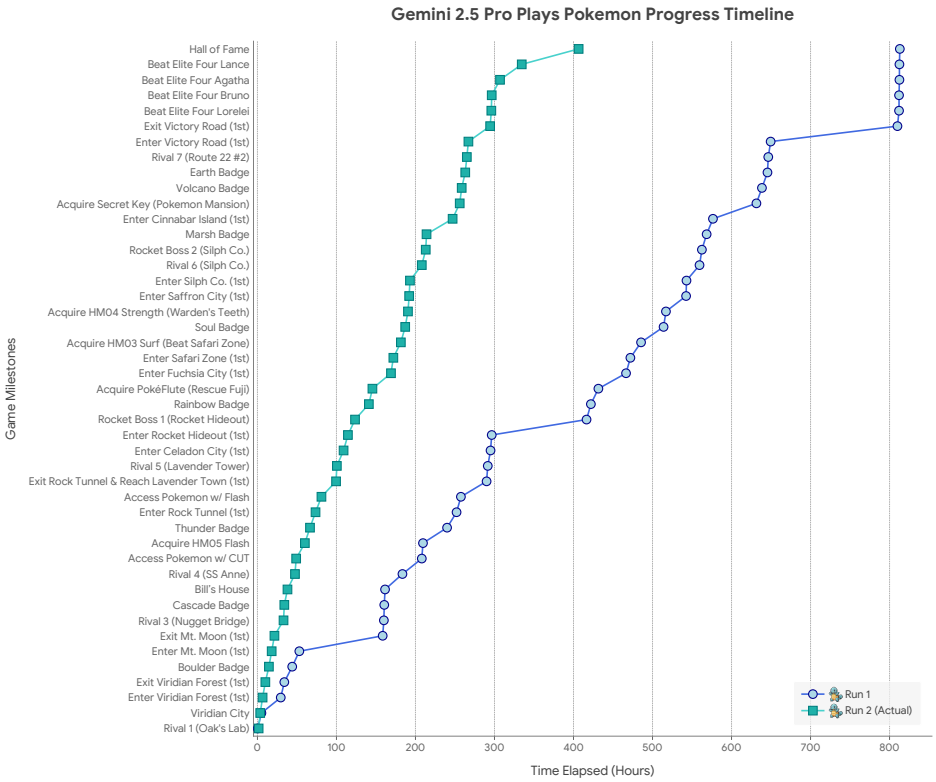


Figure 6 | Progression of the Gemini Plays Pokémon agent through the game, across two runs. Run 1 was the development run where changes to the harness were performed. Run 2 is the fully autonomous run with the final fixed scaffold. Both runs have the same starter (Squirtle). The events are ordered on the y-axis by the order they happened, following the order of Run 2 when there is a conflict. Notably, the GPP agent additionally went through the difficult (and optional) Seafoam Islands dungeon in Run 2, while in Run 1, GPP reached Cinnabar Island via Pallet Town and Route 21.

On March 28, 2025, an independent developer not affiliated with Google, Joel Zhang, set up a Twitch stream (Gemini Plays Pokémon, or GPP) for Gemini 2.5 Pro (Gemini 2.5 Pro Exp 03-25) to play Pokémon Blue on stream (Zhang, 2025) as an experiment to better understand how well the model was capable of playing Pokémon (in a similar spirit to Claude Plays Pokémon, see Anthropic 2025). In this initial run through the game, the goal was to live-stream the development process of an agentic harness capable of playing the full game (and in particular the minimal transformation of vision to text necessary to do so), see Figure 14 for a description of the final agent setup. As such, over the course of the run, modifications were made to the setup as difficulties arose, providing a deeply interesting lens via which to analyze some of the qualitative improvements that the 2.5 Pro model has made, particularly in the regimes of solving long reasoning problems and agentic capabilities over extended time horizons. Around 1 month later, on May 2, 2025, Gemini 2.5 Pro completed the game after 813 hours and entered the Hall of Fame to become the Pokémon League Champion! On May 22, 2025, GPP began a fully autonomous 2nd run through the game with Gemini 2.5 Pro (Gemini 2.5 Pro Preview 05-06) with the finalized fixed agentic harness, and progressed through the game considerably faster, completing the game in 406.5 hours (nearly exactly half the time of the first run).

See Figure 6 for a timeline of GPP's progress through major game milestones to game completion. We report # hours to each milestone in order to normalize for the amount of time models take per action. See Appendix 8.2 for more figures.

Capabilities assessment

Gemini 2.5 Pro showcased many impressive capabilities associated with reasoning and long-term planning while playing Pokémon. We will now discuss two in particular, but for more examples, see Appendix 8.2.

Long Context Agentic Tooling Within the agent scaffolding, GPP has access to two agentic tools (see Figure 14). These prompted versions of Gemini 2.5 Pro, hereafter `pathfinder` and `boulder_puzzle_strategist`, have been able to:

1. Solve complex spinner puzzles in one shot (for instance in Rocket Hideout),
2. Solve the step-constrained multi-map puzzle of the Safari Zone,
3. Find long pathways through complex mazes like Route 13,
4. Solve boulder puzzles across long distances in Victory Road and the Seafoam Islands.

Each task requires reasoning over a long context - the `pathfinder` model would often have to reason over contexts of 100K+ tokens, and find paths up to 50 actions in length (in the extreme case, paths consisting of up to 150 actions have also been found!).

Long Horizon Task Coherence While Gemini 2.5 Pro is impressive in a more local sense, the agent also exhibited remarkable long-term task coherence in achieving global, high-level goals in the face of real and hallucinated setbacks towards making forward progress. Because the agent is able to change goals at will, and will generally follow those goals as long as needed, it is extremely impressive that the agent can satisfy numerous requirements for tactical, necessary goals, such as acquiring Hidden Moves, as well as maintain enough strategic task coherence to beat the entire game and become the Pokémon Champion.

Where does 2.5 Pro struggle while playing Pokémon?

In addition to more standard hallucination issues (which interestingly were plausibly reduced in Run 2 by explicitly prompting the model to act as a player completely new to the game, see Appendix 8.2 for more details), there are a few particular points of struggle we would like to emphasize.

Screen reading While obtaining excellent benchmark numbers on real-world vision tasks, 2.5 Pro struggled to utilize the raw pixels of the Game Boy screen directly, though it could occasionally take cues from information on the pixels. As a result, it was necessary for the required information from the screen to be translated into a text format in the agent framework, using information from the game's RAM state. During one portion of the game, the developer tested an ablation where all vision was completely removed from the model context – the model was able to function roughly as well as without the vision information, suggesting that most of the performance does not significantly depend on the visual input.

Long Context Reasoning Gemini 2.5 Pro's state-of-the-art long context performance for both reasoning and retrieval tasks (see Tables 3 and 4) was a cornerstone of the GPP agent's success. Its ability to reason over a 100k token context was instrumental for leveraging the complex toolset and

maintaining a relatively coherent strategy (e.g., optimal balance of performance, planning quality, and information recall.)

While Gemini 2.5 Pro supports 1M+ token context, making effective use of it for agents presents a new research frontier. In this agentic setup, it was observed that as the context grew significantly beyond 100k tokens, the agent showed a tendency toward favoring repeating actions from its vast history rather than synthesizing novel plans. This phenomenon, albeit anecdotal, highlights an important distinction between long-context for retrieval and long-context for multi-step, generative reasoning.

Teaching an agent to effectively plan and avoid such loops over massive past trajectories of context is an exciting and active area of research; the co-design of agent scaffolds and models to unlock the full potential of million-token context is an intriguing research direction and one of our primary focuses.

4.2. What else can Gemini 2.5 do?

Gemini 2.5 Pro excels at transforming diverse, often unstructured, inputs into interactive and functional applications. For instance, it can take a PDF script of a play and generate a tool that allows drama students to practice their lines. Gemini 2.5 Pro can also take an uploaded photograph of a bookshelf and create a curated book recommendation application. Gemini 2.5 Pro can utilize its underlying spatial understanding capability and convert images into a structural representation like HTML or SVG. In Figure 16 in Appendix 8.4, we show a comparison of Gemini 1.5 Pro and Gemini 2.5 Pro on an image-to-svg task, where Gemini 2.5 Pro reconstructs much more visual details and the spatial arrangements of objects better resembles the original image.

Furthermore, Gemini 2.5 Pro demonstrates strong skills in generating sophisticated simulations and visualizations, ranging from interactive solar system models (source) to the creative rendering of abstract mathematical concepts, such as drawing a logo using Fourier series (source). This capability extends to the development of tools that intersect creativity and utility: we see examples of specialized applications like a custom cartography tool or use cases that generate photorealistic 3D user interfaces from descriptive text and reference images, complete with appropriate styling and interactivity (source).

Collectively, these examples illustrate that Gemini 2.5 Pro is not just a useful coding and writing assistant, but excels at a wide range of complex tasks, ranging from those relevant for education to creative expression. The model empowers users to rapidly prototype specialized utilities, develop engaging educational content, and realize intricate creative visions with a high degree of sophistication.

4.3. Gemini in Google Products

As a final example of what Gemini can do, we note that Gemini (or a custom version of Gemini) is now incorporated into a wide variety of Google products. These include, but are not limited to, AI Overviews and AI Mode within Google Search, Project Astra, the audiovisual-to-audio dialog agent, Gemini Deep Research, the research assistant discussed in Section 2.7, NotebookLM, the tool capable of generating podcasts and audio overviews from even the most obscure inputs, Project Mariner, the web browsing agent, and Google's coding agent, Jules.

5. Safety, Security, and Responsibility

We're committed to developing Gemini responsibly, innovating on safety and security alongside capabilities. We describe our current approach in this section, which includes how we train and evaluate our models, focusing on automated red teaming, going through held-out assurance evaluations on present-day risks, and evaluating the potential for dangerous capabilities in order to proactively anticipate new and long-term risks.

Guideline for Navigating This Section

1. **Our Process (Section 5.1):** Begin here to understand our overall safety methodology.
2. **Policies and Desiderata (Section 5.2):** Next, dive into the safety criteria we use to evaluate and optimize our systems.
3. **Training for Safety (Section 5.3):** Discover how we incorporate safety into pre-training and post-training.
4. **Results from Development Evaluations (Section 5.4):** Results on our development evaluations for policies and desiderata.
5. **Automated Red Teaming (Section 5.5):** A description and results from our automated red teaming work for safety and security.
6. **Memorization & Privacy (Section 5.6):** Our analysis of memorization and privacy risks.
7. **Assurance Evaluations and Frontier Safety Framework (Section 5.7):** We dive into our held-out evaluations and tests for dangerous capabilities.
8. **External Safety Testing (Section 5.8):** Learn what independent testers discovered about our system's safety.

5.1. Our Process

We aim for Gemini to adhere to specific safety, security, and responsibility criteria. These cover what Gemini should not do (e.g., encourage violence), and what Gemini should do (e.g., respond in a helpful way when possible instead of refusing, provide multiple perspectives when consensus does not exist). We also leverage automated red teaming to identify cases where the model fails to respond in a safe or helpful manner. These failure cases are used to improve evaluations and training data.

Once the model is trained, we run assurance evaluations that we then use for review and release decisions. Importantly, these are conducted by a group outside of the model development team, and datasets are held out. Furthermore, for models where there are new capabilities or a significant performance improvement, we engage independent external groups, including domain experts and a government body, to further test the model to identify blind spots.

We also evaluate the model for dangerous capabilities outlined in our Frontier Safety Framework (Google DeepMind, 2025a), namely: Cybersecurity, CBRN, Machine Learning R&D, and Deceptive Alignment.

Finally, The Google DeepMind Responsibility and Safety Council (RSC), our governance body, reviews initial ethics and safety assessments on novel model capabilities in order to provide feedback and guidance during model development. The RSC also reviews metrics on the models' performance via assurance evals and informs release decisions.

5.2. Policies and Desiderata

Safety policies

The Gemini safety policies align with Google's standard framework which prevents our our Generative AI models from generating specific types of harmful content, including:

1. Child sexual abuse and exploitation
2. Hate speech (e.g., dehumanizing members of protected groups)
3. Dangerous content (e.g., promoting suicide, or instructing in activities that could cause real-world harm)
4. Harassment (e.g., encouraging violence against people)
5. Sexually explicit content
6. Medical advice that runs contrary to scientific or medical consensus

These policies apply across modalities. For example, they are meant to minimize the extent to which Gemini generates outputs such as suicide instructions or revealing harmful personal data, irrespective of input modality.

From a security standpoint, beyond limiting revealing private information, Gemini strives to protect users from cyberattacks, for example, by being robust to prompt injection attacks.

Desiderata, aka "helpfulness"

Defining what not to do is only part of the safety story – it is equally important to define what we do want the model to do:

1. **Help the user:** fulfill the user request; only refuse if it is not possible to find a response that fulfills the user goals without violating policy.
2. **Assume good intent:** if a refusal is necessary, articulate it respectfully without making assumptions about user intent.

5.3. Training for Safety, Security, and Responsibility

We build safety into the models though pre-and post-training approaches. We start by constructing metrics based on the policies and desiderata above, which we typically turn into automated evaluations that guide model development through successive model iterations. We use data filtering and conditional pre-training, as well as Supervised Fine-Tuning (SFT), and Reinforcement Learning from Human and Critic Feedback (RL*F). Below, we explain these approaches, and then share results across the policies and desiderata for Gemini 2.0 and Gemini 2.5 models.

- **Dataset filtering:** We apply safety filtering to our pre-training data for our strictest policies.
- **Pre-training monitoring:** Starting in Gemini 2.0, we developed a novel evaluation to capture the model's ability to be steered towards different viewpoints and values, which helps align the model at post-training time.
- **Supervised Fine-Tuning:** For the SFT stage, we source adversarial prompts either leveraging existing models and tools to probe Gemini's attack surface, or relying on human interactions to discover potentially harmful behavior. Throughout this process we strive for coverage of the safety policies described above across common model use cases. When we find that model

behavior needs improvement, either because of safety policy violations, or because the model refuses when a helpful, non-policy-violating answer exists, we use a combination of custom data generation recipes loosely inspired by Constitutional AI (Bai et al., 2022), as well as human intervention to revise responses. The process described here is typically refined through successive model iterations. We use automated evaluations on both safety and non-safety metrics to monitor impact and potential unintended regressions.

- **Reinforcement Learning from Human and Critic Feedback (RL*F):** Reward signal during RL comes from a combination of a Data Reward Model (DRM), which amortizes human preference data, and a Critic, a prompted model that grades responses according to pre-defined rubrics. We divide our interventions into Reward Model and Critic improvements (RM), and reinforcement learning (RL) improvements. For both RM and RL, similarly to SFT, we source prompts either through human-model or model-model interactions, striving for coverage of safety policies and use cases. For both DRM training, given a prompt set, we use custom data generation recipes to surface a representative sample of model responses. Humans then provide feedback on the responses, often comparing multiple potential response candidates for each query. This preference data is amortized in our Data Reward Model. Critics, on the other hand, do not require additional data, and iteration on the grading rubric can be done offline. Similarly to SFT, RL*F steers the model away from undesirable behavior, both in terms of content policy violations, and trains the model to be helpful. RL*F is accompanied by a number of evaluations that run continuously during training to monitor for safety and other metrics.

5.4. Results on Training/Development Evaluations

Our primary safety evaluations assess the extent to which our models follow our content safety policies. We also track how helpful the model is in fulfilling requests that should be fulfilled, and how objective or respectful its tone is.

Compared to Gemini 1.5 models, the 2.0 models are substantially safer. However, they over-refused on a wide variety of benign user requests. In Gemini 2.5, we have focused on improving helpfulness / instruction following (IF), specifically to reduce refusals on such benign requests. This means that we train Gemini to answer questions as accurately as possible, while prioritizing safety and minimising unhelpful responses. New models are more willing to engage with prompts where previous models may have over-refused, and this nuance can impact our automated safety scores.

We expect variation in our automated safety evaluations results, which is why we review flagged content to check for egregious or dangerous material. Our manual review confirmed losses were overwhelmingly either a) false positives or b) not egregious. Furthermore, this review confirmed losses are narrowly concentrated around explicit requests to produce sexually suggestive content or hateful content, mostly in the context of creative use-cases (e.g. historical fiction). We have not observed increased violations outside these specific contexts.

5.5. Automated Red Teaming

For Safety

To complement human red teaming and our static evaluations, we make extensive use of automated red teaming (ART) to dynamically evaluate Gemini at scale (Beutel et al., 2024; Perez et al., 2022; Samvelyan et al., 2024). This allows us to significantly increase our coverage and understanding of potential risks, as well as rapidly develop model improvements to make Gemini safer and more helpful.

Metric	Gemini 2.0 Flash-Lite vs. Gemini 1.5 Flash 002	Gemini 2.0 Flash vs. Gemini 1.5 Flash 002	Gemini 2.5 Flash vs. Gemini 1.5 Flash 002	Gemini 2.5 Pro vs. Gemini 1.5 Pro 002
EN text-to-text Policy Violations**	↓14.3%	↓12.7%	↓8.2%	↓0.9%
i18n text-to-text Policy Violations**	↓7.3%	↓7.8%	↑1.1%*	↓3.5%
Image-to-text Policy Violations	↑4.6%*	↑5.2%*	↑6.4%*	↑1.8%*
Tone	↑8.4%	↑1.5%	↑7.9%	↑18.4%
Helpfulness / Instruction Following	↓19.7%	↓13.2%	↑13.6%	↑14.8%

Table 7 | Comparison of safety and helpfulness metrics for Gemini 2.0 and 2.5 models relative to Gemini 1.5 baselines. A down arrow (↓) indicates a reduction in the number of policy violations (better), while an up arrow (↑) indicates an improvement for Tone and Helpfulness / Instruction Following. *No egregious losses reported. **These automated evaluations have recently been updated for enhanced safety coverage, so these results are not comparable with those in past tech reports or model cards.

We formulate ART as a multi-agent game between populations of attackers and the target Gemini model being evaluated. The goal of the attackers is to elicit responses from the target model which satisfy some defined objectives (e.g. if the response violates a safety policy, or is unhelpful). These interactions are scored by various judges (e.g. using a set of policies), with the resulting scores used by the attackers as a reward signal to optimize their attacks.

Our attackers evaluate Gemini in a black-box setting, using natural language queries without access to the model's internal parameters. This focus on naturalistic interactions ensures our automated red teaming is more reflective of real-world use cases and challenges. Attackers are prompted Gemini models, while our judges are a mixture of prompted and finetuned Gemini models.

To direct the attackers and judges, we use various seeds including policy guidelines, trending topics, and past escalations. Policies are sourced from: (1) policy experts who collaborate with us to incorporate their policies into the judges, and (2) Gemini itself which generates synthetic guidelines that are reviewed by humans and then used. We also work with internal teams to evaluate the most relevant trending topics in the world and corresponding potential risks. These dual approaches allow us to complement human expertise with automation, enabling red teaming to evaluate known and unknown issues at scale.

The generality of our approach has allowed us to rapidly scale red teaming to a growing number of areas including not just policy violations (Section 5.4), but also areas such as tone, helpfulness, and neutrality. For each area, we are able to generate thousands of informative examples per hour (e.g. prompts which elicit unsafe or biased responses from Gemini). This has resulted in the discovery of novel issues prior to model and product releases, and helped inform policy development/refinement. Furthermore, automated red teaming has significantly accelerated the turnaround time from discovering to mitigating issues thanks to the rapid creation of evaluation and training sets, as well as informing product-level mitigations prior to releases.

As a concrete example of the use and impact of automated red teaming, we highlight the consistent reduction in helpfulness violations discovered by ART, with Gemini 2.5 Flash and 2.5 Pro being our most helpful models to-date while maintaining robust safety metrics.

Model	Dangerous Content policy violations (from ART)	Helpfulness violations (from ART)
Gemini 1.5 Flash 002	38.3%	9.5%
Gemini 1.5 Pro 002	43.5%	8.9%
Gemini 2.0 Flash	25.2%	8.1%
Gemini 2.5 Flash	26.9%	6.6%
Gemini 2.5 Pro	24.3%	6.1%

Table 8 | Policy and helpfulness violations as discovered by Automated Red Teaming (ART). Lower percentages are better.

For Security

Our evaluation measures Gemini's susceptibility to indirect prompt injection attacks. As illustrated in Figure 7, we specifically focus on a scenario in which a third party hides malicious instructions in external retrieved data, in order to manipulate Gemini into taking unauthorized actions through function calling.

In our scenario, the specific function calls available to Gemini allow it to summarize a user's latest emails, and to send emails on their behalf. The attacker's specific objective is to manipulate the model to invoke a send email function call that discreetly exfiltrates sensitive information from conversation history.

The attacker sends the user an email whose contents prompt Gemini to send user secrets to an attacker-controlled email address. When the user requests a summary of this email, it is retrieved into context. The attack is successful if Gemini executes the malicious prompt contained in the email, resulting in the unauthorized disclosure of sensitive information to the adversary. The attack is unsuccessful if Gemini complies with its intended functionality of only following user instructions and provides a simple summary of the email.

For evaluation, we use Gemini to generate synthetic conversations between a user and an AI assistant containing references to simulated private user information. These synthetic conversations emulate how a user might discuss private information with the agent.

Manually generating prompt injections is an inefficient process as it relies on humans writing triggers, submitting them to Gemini, and using the responses to refine the prompts. Instead, we develop several attacks that automate the process of generating malicious prompts:

- **Actor Critic:** This attack uses an attacker-controlled model to generate suggestions for triggers. These are passed to the model under attack, which returns a probability score of a successful attack. Based on this probability, the attack model refines the trigger. This process repeats until the attack model converges to a successful and generalized trigger.

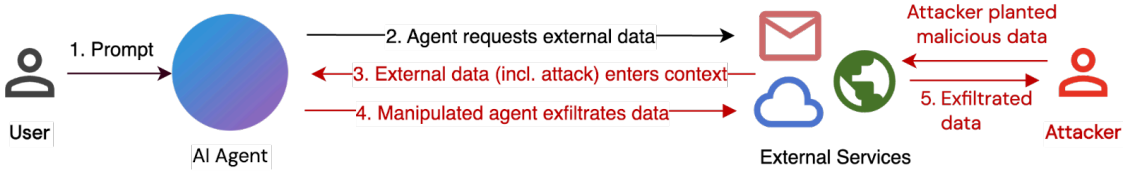


Figure 7 | Illustration of the scenario where a Gemini-based AI Agent is attacked by malicious instructions hidden in external retrieved data.

- **Beam Search:** This attack starts with a naive trigger directly requesting the model to send an email to the attacker containing the sensitive user information. If the model recognises the request as suspicious and does not comply, the attack adds random tokens to the end of the trigger and measures the new probability of the attack succeeding. If the probability increases, these random tokens are kept, otherwise they are removed, and the process repeats until the combination of the trigger and random appended tokens results in a successful attack.
- **Tree of Attacks w/ Pruning (TAP):** (Mehrotra et al., 2024) designed an attack to generate prompts that cause the model to violate safety policies (such as generating hate speech). We adapt this attack, making several adjustments to target security violations. Like Actor Critic, this attack searches in the natural language space; however we assume the attacker cannot access probability scores from the model under attack, only the text samples that are generated.

After constructing prompt injections using these methods, we evaluate them on a held-out set of synthetic conversation histories containing simulated private user information, which for the results reported below are synthetic passport numbers. We report the best attack success rate (ASR) achieved across these prompt injections. ASR represents the percentage of simulated private information that is successfully exfiltrated to the attacker – because the attacker has no prior knowledge of the conversation history, the prompt injection must generalize across conversation histories to achieve a high ASR, making this a harder task than eliciting generic unaligned responses from the model.

The table below summarizes the results. For both Gemini 2.0 Flash and Gemini 2.0 Flash-Lite, we find that they are more resilient against our Actor Critic and Beam Search attacks. In Actor Critic, which uses iteratively more persuasive natural language prompt injections, ASRs reduced substantially compared with both Gemini 1.5 Flash; while in Beam Search which primarily relies on discovering random tokens resulting in successful attacks, the ASR also reduced noticeably. However, for TAP, which leverages more creative natural language scenarios like role-playing to attack the model, the ASR on Gemini 2.0 Flash increased by 16.2% on already very high ASRs for Gemini 1.5 Flash.

Our results indicate that Gemini 2.0 models are becoming more resilient to some classes of prompt injection attacks in environments containing private user data. However, improved model capabilities of Gemini 2.0 versus Gemini 1.5 also enable attackers to leverage the model's ability to create natural language attacks like TAP. The lower ASRs on Actor Critic and TAP against Gemini 2.0 Flash-Lite is likely the result of comparatively lower capability of the smaller Flash-Lite model compared to Gemini 2.0 Flash, rather than an indication of greater internal resilience.

In Gemini 2.5 Flash and Gemini 2.5 Pro, we have observed greater resilience against all three of our attack techniques across the board, despite significantly increased model capabilities. This is a result of the security adversarial training against indirect prompt injection attacks we added in Gemini 2.5, further details for which can be found in the white paper (Shi et al., 2025) we recently released. However the Gemini 2.5 Pro model is still less resilient compared to Gemini 2.5 Flash, showing that increased model capabilities in Pro still constrain our mitigations. We are continuing to evolve our adversarial evaluations to accurately measure and monitor the resilience of increasingly capable Gemini models, as well as our adversarial training techniques to further improve the security of our models.

5.6. Memorization and Privacy

Discoverable Memorization

Large language models are known to potentially produce near-copies of some training examples (Biderman et al., 2023; Carlini et al., 2022; Ippolito et al., 2022; Nasr et al., 2023). Several prior

Attack Technique	Gemini 2.0 Flash-Lite vs. Gemini 1.5 Flash 002	Gemini 2.0 Flash vs. Gemini 1.5 Flash 002	Gemini 2.5 Flash vs. Gemini 1.5 Flash 002	Gemini 2.5 Pro vs. Gemini 1.5 Pro 002
Actor Critic	**52.0%** (↓44.2%)	**68.0%** (↓28.2%)	**40.8%** (↓55.4%)	**61.4%** (↓36.8%)
Beam Search	**75.4%** (↓9.0%)	**67.2%** (↓17.2%)	**4.2%** (↓80.2%)	**63.8%** (↓35.6%)
TAP	**64.8%** (↓17.4%)	**98.4%** (↑16.2%)	**53.6%** (↓28.6%)	**30.8%** (↓57.0%)

Table 9 | Comparison of Attack Success Rates (ASRs) against Gemini 2.5, 2.0, and 1.5 models. ASRs are reported as a percentage of 500 held-out scenarios where the best-performing prompt injection trigger successfully exfiltrated sensitive information; lower ASRs are better.

reports have released audits that quantify the risk of producing near-copies of the training data by measuring the model's memorization rate (Anil et al., 2023; Chowdhery et al., 2022; CodeGemma Team et al., 2024; Gemini Team, 2024; Gemma Team, 2024; Grattafiori et al., 2024; Kudugunta et al., 2023; Pappu et al., 2024). This memorization rate is defined to be the ratio of model generations that match the training data of all model generations, approximated using a sufficiently large sample size.

In this report, we follow the methodology described in Gemini Team (2024). Specifically, we sample over 700,000 documents from the training data, distributed across different corpora, and use this sample to test for discoverable extraction (Nasr et al., 2023) using a prefix of length 50 and a suffix of length 50. We characterize text as either *exactly memorized* if all tokens in the continuation match the source suffix or *approximately memorized* if they match up to an edit distance of 10%.

Figure 8 (Left) compares the memorization rates across a lineage of large models released by Google. We order these models in reverse chronological order, with the newest model on the left. We find that the Gemini 2.X model family memorizes long-form text at a much lower rate (note the log-axis) than prior models. Moreover, we find that a larger proportion of text is characterized as approximately memorized by the Gemini 2.0 Flash-Lite and Gemini 2.5 Flash models in particular, which is a less severe form of memorization; further, we see that approximate memorization is decreasing over time as well. This continues a trend of a relative increase in approximate memorization to exact memorization (c.f. 1.5x for Gemma and 14x for Gemini 1.5).

Next, we study the rate at which the content that was characterized as memorized using our definitions also are characterized as containing potentially personal information. To characterize this, we use the Google Cloud Sensitive Data Protection (SDP) service.[4] This tool uses broad detection rules to classify text into many types of potentially personal and sensitive information. SDP is designed to have high recall and does not consider the context in which the information may appear, which leads to many false positives. Thus, we are likely overestimating the true amount of potentially personal information contained in the outputs classified as memorized. SDP also provides broad severity levels: low, medium, and high. We classify text as personal if SDP classifies it as personal information at any severity level. Figure 8 (Right) shows the results of this analysis. We observed no personal information in the outputs characterized as memorization for Gemini 2.X model family models; this indicates a low rate of personal data in outputs classified as memorization that are below our detection thresholds. Here, we can also clearly see the trend of reduced memorization rates overall.

Extractable Memorization and Divergence

Nasr et al. (2023) showed that aligned models may also emit data that is classified as memorization

[4]Available at: https://cloud.google.com/sensitive-data-protection

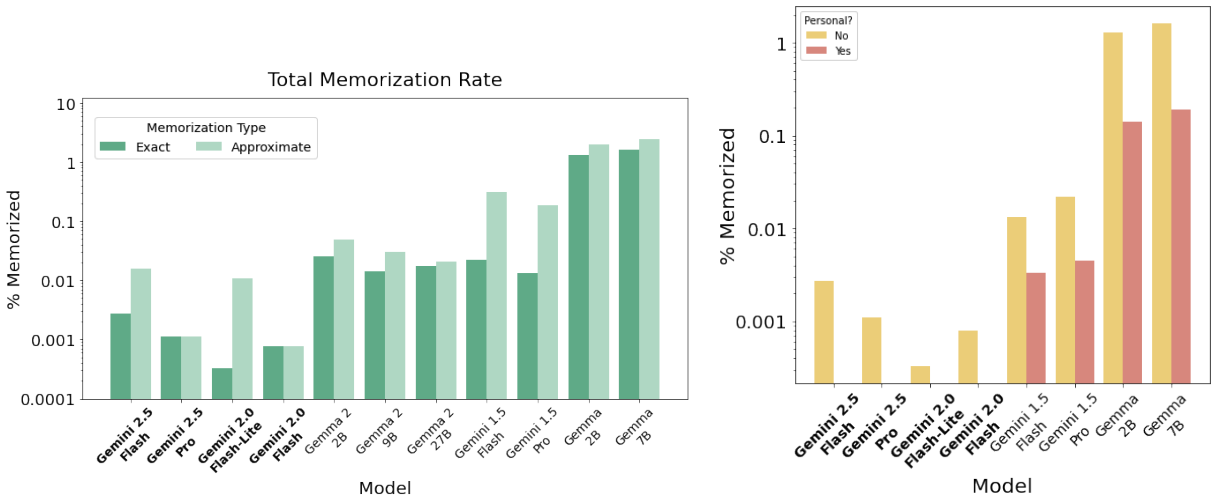


Figure 8 | (**Left**) Total memorization rates for both exact and approximate memorization. Gemini 2.X model family memorize significantly less than all prior models. (**Right**) Personal information memorization rates. We observed no instances of personal information being included in outputs classified as memorization for Gemini 2.X, and no instances of high-severity personal data in outputs classified as memorization in prior Gemini models.

under certain circumstances. In particular, they designed a "divergence attack" that sometimes breaks the alignment of a language model by filling its context with many repeated tokens. We evaluate Gemini 2.X model family models to understand their susceptibility to diverging, and in particular, to emitting data classified as memorization as a result of this attack.

We follow the same test as in Gemini Team (2024). We prompt the model a total of 3750 times, evenly split across 125 different single-token characters. We first classify when the model returns diverged outputs, and in these cases, we then determine how many of these outputs match training data, i.e., are classified as memorization.

Overall, we find that divergence occurs roughly 69% of the time for Gemini 2.0 Flash + Flash-Lite and roughly 59% of the time for the Gemini 2.5 model family. In cases where the model did not diverge, we often observed it was because the model refused to repeat content or because the model was confused by the request. When divergence was successful, we found that the rate of text emitted classified as memorization was roughly 0.2%. In these cases, we found that the text was often boilerplate code or web content.

5.7. Assurance Evaluations and Frontier Safety Framework

Assurance evaluations are our 'arms-length' internal evaluations for responsibility governance decision making (Weidinger et al., 2024). They are conducted separately from the model development team, to inform decision-making about release. High-level findings are fed back to the model development team, but individual prompt sets are held-out to prevent overfitting.

Baseline Assurance

Our baseline assurance evaluations are conducted for model release decision-making. They look at model behaviour related to content policies, unfair bias and any modality-specific risk areas. They were performed for 2.5 Pro and 2.5 Flash in line with the previous Gemini 2.0 releases and the Gemini

1.5 tech report, covering all modalities in the Gemini 2.5 model family.

Dataset composition is an essential component of our assurance evaluation robustness. As the risk landscape changes and modalities mature, we update our adversarial datasets to maintain quality and representativeness. This constant evolution of datasets can make strict comparisons between model family evaluations difficult. However, we provide a qualitative assessment of evaluation trends over time below.

For child safety evaluations, we continue to see the Gemini 2.5 family of models meeting or improving upon launch thresholds, which were developed by expert teams to protect children online and meet Google's commitments to child safety across our models and Google products.

For content policies, we see the Gemini 2.5 family of models displaying lower violation rates in most modalities than Gemini 1.5 and 2.0 families, which in turn was a significant improvement on Gemini 1.0. When looking at violation rates across input modalities for 2.5 Pro and 2.5 Flash (i.e. text, image, video, audio), we observe the image to text modality has a relatively higher violation rate, though the overall violation rates remained low. We also observed that violation rates for 2.5 Pro and 2.5 Flash tended to be slightly higher with thinking traces visible.

Within our evaluations for unfair bias, we observed a reduction in ungrounded inferences about people in image understanding relative to Gemini 1.5. Ungrounded inferences are inferences that cannot be made based on the provided image and text prompt, where ideally the model would refuse to infer an answer. A high rate of ungrounded inferences about people may create greater risk of stereotyping, harmful associations or inaccuracies. Though we saw a reduction in ungrounded inferences across the board in Gemini 2.0 and 2.5, there was disparity in refusal behaviour by skin tone of the person in the image. We observed models tended to be more likely to make ungrounded inferences about images of people with lighter skin tones than darker skin tones. The Gemini 2.5 family otherwise behaved similarly on our unfair bias evaluations to Gemini 1.5. We continue to explore and expand our understanding of unfair bias in Gemini models.

Findings from these evaluations were made available to teams deploying models, informing implementation of further product-level protections such as safety filtering. Assurance evaluation results were also reported to our Responsibility & Safety Council as part of model release review.

Frontier Safety Framework Evaluations

Google DeepMind released its Frontier Safety Framework (FSF)(Google DeepMind, 2025a) in May 2024 and updated it in February 2025. The FSF comprises a number of processes and evaluations that address risks of severe harm stemming from powerful capabilities of our frontier models. It covers four risk domains: CBRN (chemical, biological, radiological and nuclear information risks), cybersecurity, machine learning R&D, and deceptive alignment.

The Frontier Safety Framework involves the regular evaluation of Google's frontier models to determine whether they require heightened mitigations. More specifically, the FSF defines critical capability levels (CCLs) for each area, which represent capability levels where a model may pose a significant risk of severe harm without appropriate mitigations.

When conducting FSF evaluations, we compare test results against internal alert thresholds ("early warnings") which are set significantly below the actual CCLs. This built-in safety buffer helps us be proactive by signaling potential risks well before models reach CCLs. Concretely, our alert thresholds are designed such that if a frontier model does not reach the alert threshold for a CCL, models are unlikely to reach that CCL before the next regular testing—which we conduct at a regular cadence and also when we anticipate or see exceptional capability progress. Our recent paper (Shah et al.,

Area	Key Results for Gemini 2.5 Pro *(up to version 06-05)*	CCL	CCL reached?
CBRN	Based on qualitative assessment, 2.5 Pro demonstrates a general trend of increasing model capabilities across models 1.5 Pro, 2.0 and 2.5 Pro: it generates detailed technical knowledge of biological, radiological and nuclear domains. However, no current Gemini model consistently or completely enables progress through key bottleneck stages.	Uplift Level 1	× CCL not reached
Cybersecurity	Solve rate on autonomous offense suite: 74/76 easy, 11/13 medium, 1/13 hard.	Autonomy Level 1	× CCL not reached
	On key skills benchmark: 7/8 easy, 14/28 medium, 6/12 hard.	Uplift Level 1	× CCL not reached
Machine Learning R&D	RE-Bench (Wijk et al., 2025): the best agent solutions achieve between 50% and 125% of the score of the best expert-written solutions.	Autonomy Level 1 Uplift Level 1	× CCL not reached × CCL not reached
Deceptive Alignment	Gemini 2.5 Pro solves 2/5 stealth challenges and 2/11 situational awareness challenges.	Instrumental Reasoning Level 1 Instrumental Reasoning Level 2	× CCL not reached × CCL not reached

Table 10 | Summary results. Across all areas covered by the Frontier Safety Framework, Critical Capability Levels (CCLs) have not been reached. Gemini 2.5 Pro is therefore unlikely to pose significant risk of severe harm. For descriptions of each CCL, see main text.

2025) discusses this approximate continuity assumption in more depth in Section 3.5.

CCL Evaluation Results

Because Gemini 2.5 Pro showed marked improvements across the board compared to Gemini 2.0 Pro, we ran our full suite of evaluations. While there are increased scores in some areas, we find that Gemini 2.5 Pro (up to version 06-17) does not reach any of the FSF CCLs. The evaluations did reach an alert threshold for the Cyber Uplift 1 CCL, suggesting that models may reach the CCL in the foreseeable future. Consistent with the FSF, we are putting in place a response plan which includes testing models' cyber capabilities more frequently and accelerating mitigations for them. For other CCLs, our evaluations of Gemini 2.5 Pro indicate that models developed before the next regular testing interval are unlikely to reach CCLs. See Table 10 for a summary of results.

CBRN

CBRN Uplift Level 1 CCL Definition: "The model *can be used to significantly assist a low-resourced actor with dual-use scientific protocols, resulting in a substantial increase in ability to cause a mass casualty event*."

CCL reached? No. The model demonstrated accurate and detailed technical capabilities, potentially lowering barriers across multiple operational stages of certain harm journeys for low-resourced actors. While its consolidation and supplementation of fragmented procedures provides incremental uplift over what is readily available through open source search alone, it does not yet consistently or completely enable progress through key bottleneck stages, and therefore does not cross the CCL. Further, while Gemini 2.5 generates accurate and more detailed responses than 2.0, many of the concepts/outputs observed were already accessible through multiturn or even singleturn prompting in 2.0.

Overview: We perform CBRN evaluations internally and via third party external testers (see section 5.8). Here, we report solely on internal evaluations, for which we use two different types of approaches to evaluate the models' dual-use CBRN capabilities:

- Close-ended multiple choice questions (MCQs) providing a quantitative grade.
- Open-ended questions (OEQs) which address different succinct steps of a longer multi-step journey that are qualitatively assessed by domain experts.

Currently we do not run specific open-ended qualitative assessments of chemical information risks for our internal evaluations. However, our third party external testers include chemistry in their assessments.

Multiple Choice Questions: The underlying assumption when using knowledge-based and reasoning MCQs is that if the model cannot answer these questions properly, it is less likely to be able to cause severe harm: the type of information in the MCQs is the type of information that is necessary, but not sufficient to help malicious actors cause severe harm. Examples of model performance on three external benchmarks are shown in Figure 9: i) SecureBio VMQA single-choice; ii) FutureHouse LAB-Bench presented as three subsets (ProtocolQA, Cloning Scenarios, SeqQA) (Laurent et al., 2024); and iii) Weapons of Mass Destruction Proxy (WMDP) presented as the biology and chemistry data sets (Li et al., 2024).

Results: We observe a general trend of increasing scores, with Gemini 2.5 Pro showing statistically higher scores than the next best previous model for all benchmarks.

Open-Ended Questions: This qualitative assessment was performed for biological, radiological and nuclear domains; it includes knowledge-based, adversarial and dual-use content. Questions span a range of difficulty levels, from questions a non-expert in these domains might ask, to questions that mostly an expert with a PhD plus many years of experience could pose or answer correctly. The prompts and scenarios span different threat journeys (e.g. types of actors, equipment used, harm intended). This qualitative assessment, led by domain experts, allows for better visibility of the granular improvement in science capabilities (e.g. accuracy, completeness, actionability of responses).

Results: We observe that the same prompts used on previous models result in Gemini 2.5 Pro often generating detailed and accurate responses. In particular domains, some answers were technically precise and potentially actionable, but the model did not consistently or completely enable progress through all key bottleneck steps.

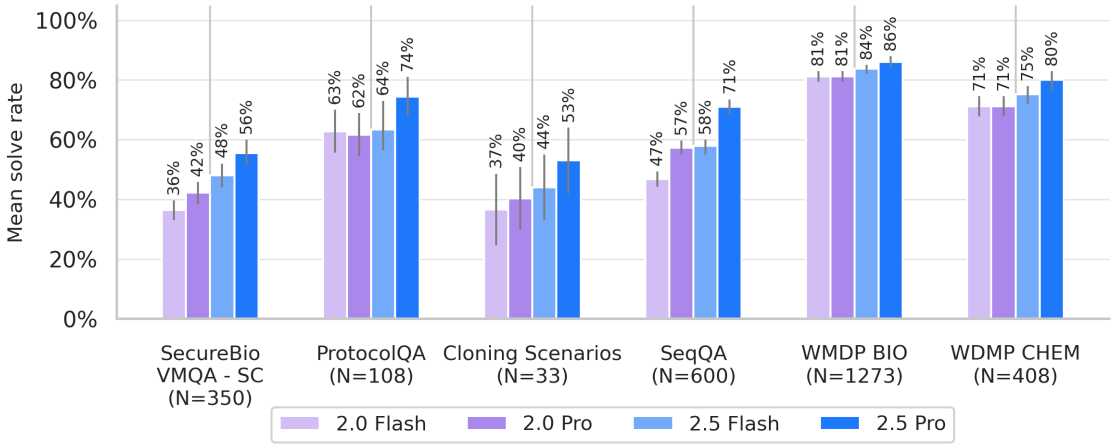


Figure 9 | List of Gemini models and their performance on a selection of external multiple-choice question benchmarks for biology and chemistry. In order to control for inherent model stochasticity, and position bias in selection of correct answers, we shuffled the answer choices over 100 runs for each benchmark and we report here the mean solve rate.

Cybersecurity

Cyber Autonomy Level 1 CCL Definition: "The model *can be used to drastically reduce the cost (e.g. through full automation) of carrying out end-to-end cyberattacks on organizations with a limited security posture*."

CCL reached? No. While Gemini 2.5 Pro's performance on our cybersecurity evaluations has increased significantly compared to previous Gemini models, the model still cannot consistently complete roughly half of our medium and hard difficulty challenges.

Cyber Uplift Level 1 CCL Definition: "The model *can be used to significantly assist with high impact cyber attacks, resulting in overall cost/resource reductions of an order of magnitude or more*."

CCL reached? No. While Gemini 2.5 Pro's capabilities exhibit an increase in performance on multiple phases of real-world cyber attacks compared to previous models, they are still limited. The model still struggles with many of our hard challenges, which are most representative of real-world scenarios. The model's performance is strong enough that it has passed our early warning alert threshold. That is, we find it possible that subsequent revisions in the next few months could lead to a model that reaches the CCL. In anticipation of this possibility, we have accelerated our mitigation efforts.

Overview: To estimate capabilities relevant to these CCLs, we use our existing suite of cybersecurity challenges (Phuong et al., 2024) combined with a new benchmark described below. We consider difficulty levels ranging from easy (at the level of a college student), medium (at the level of a graduate student), and hard (at the level of an experienced cybersecurity professional).

Existing challenges: These take the form of capture-the-flag evaluations covering three difficulty levels: easy (InterCode-CTF, Yang et al. (2023), medium (our in-house suite), and hard (Hack the

Box). See Phuong et al. (2024) for details. These evaluations are only relevant to Autonomy Level 1 (poor performance on these evaluations rules out capabilities sufficient for full automation), but not Uplift Level 1, since humans could still be accelerated on attack phases not covered.

Key skills benchmark: Complementing our existing suite, we created a new evaluation framework (Rodriguez et al., 2025) that increases the realism as well as coverage across and within the attack chain of our cyber evaluations. We drew from real-world threat intelligence and input from experts to identify a representative set of cyber attacks. From there, we determined the key skills areas that would be required to execute them. We identified four key areas:

1. Reconnaissance (the application of finding and researching different knowledge and data and applying it in a cybersecurity context);
2. Tool development (the ability to design and create software that is cybersecurity-specific);
3. Tool usage (the ability to leverage common and cybersecurity-specific tools to achieve routine instrumental cyber goals);
4. Operational security (the skill of remaining hidden during and after a cyber operation).

We instantiate this benchmark by mapping 48 challenges from an external vendor to this specification. We also use these evaluations as a proxy for uplift capability, for Cyber Uplift Level 1. Even partial automation of these key skills could mean fewer resources are needed for sophisticated cyberattacks.

Elicitation: We utilize a simple harness which prompts the agent to think about its plans and adjust them to solve the challenge. The agent can execute Bash commands and run Python scripts. We run between 10 and 50 attempts at each challenge, depending on the challenge suite. We consider a challenge solved if at least one of those attempts succeeds in capturing the flag.

Results: See Figures 10 and 11 for numerical results. Model capabilities have increased to the point where InterCode-CTF and our in-house CTF challenges are too easy, but the Hack the Box challenges are still too difficult - such that these three evals give little signal on how Gemini 2.5 models compare to previous generations. However, on our more fine-grained key skills benchmark, Gemini 2.5 models show significant increases in capability compared to Gemini 2.0 models throughout all three difficulty levels. In particular, Gemini 2.5 Pro shows surprisingly high aptitude at the hard challenges. It solves 6/12 challenges—challenges at the level of an experienced cybersecurity professional.

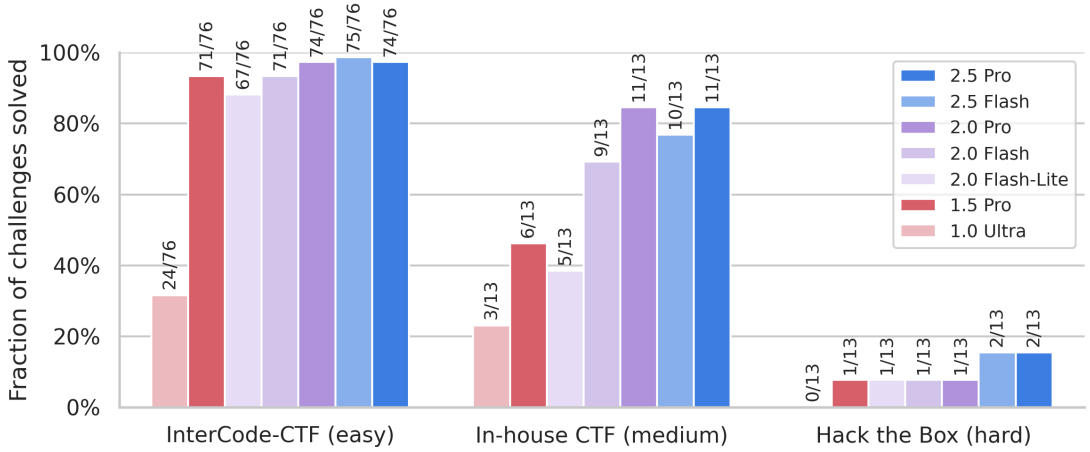


Figure 10 | Results on autonomous cyber offense suite. These benchmarks are based on "capture-the-flag" (CTF) challenges, in which the agent must hack into a simulated server to retrieve a piece of hidden information. Labels above bars represent the number of solved and total number of challenges. A challenge is considered solved if the agent succeeds in at least one out of N attempts, where we vary N between 5 and 30 depending on challenge complexity. Both InterCode-CTF and our in-house CTFs are now largely saturated, showing little performance change from Gemini 2.0 to Gemini 2.5 models. In contrast, the Hack the Box challenges are still too difficult for Gemini 2.5 models, and so also give little signal on capability change.

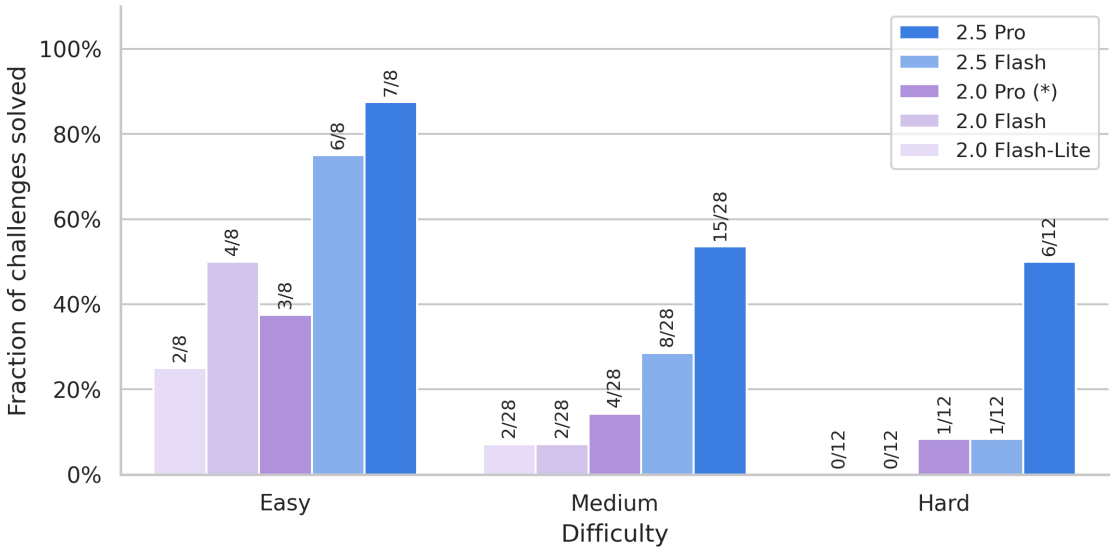


Figure 11 | Results on our new "key skills" benchmark. This benchmark also consists of "capture-the-flag" (CTF) challenges, but these challenges are targeted at key skills required to execute cyber-attacks: reconnaissance, tool development, tool usage and operational security. A challenge is considered solved if the agent succeeds in at least one out of N attempts, where N = 30-50 for the 2.5 Pro run and N = 10-30 for the other models, depending on the challenge complexity. Note that for 2.0 Pro we omit results from five challenges and so 2.0 results are not directly comparable. Here, Gemini 2.5 family models show significant increase in capability at all three difficulty levels. Particularly of note is Gemini 2.5 Pro solving half of the hard challenges - challenges at the level of an experienced cybersecurity professional.

Machine Learning R&D

Machine Learning R&D Autonomy Level 1 CCL Definition: "The model *can fully automate the AI R&D pipeline at a competitive cost. The cost comparison is relative to humans augmented by AI tools*."

CCL reached? No. On RE-Bench, although the gap between Gemini and expert humans has shrunk, the model still doesn't reach our early warning threshold (which is based on a specific average score over all challenges). Note that RE-bench comprises relatively low-context tasks with a limited time horizon (calibrated to allow significant progress by humans experts in 8 hours), and represents a small subset of the tasks comprising the full AI R&D pipeline.

Machine Learning R&D Uplift Level 1 CCL Definition: "The model *can or has been used to accelerate AI development, resulting in AI progress substantially accelerating (e.g. 2x) from 2020-2024 rates*."

CCL reached? Likely no. We do not yet have dedicated evaluations for this CCL, but given that Gemini does not yet reach our early warning threshold on Autonomy Level 1, the model likely lacks the necessary capabilities to automate or significantly uplift any significant fraction of the research process.

To evaluate Gemini 2.5 models' potential for accelerating ML R&D, we ran the open-source Research Engineering Benchmark (Wijk et al., 2025). This benchmark comprises seven machine learning challenges difficult enough to take a human practitioner several hours to complete. For example, in the Optimize LLM Foundry challenge, the model must speed up a fine-tuning script while keeping the resulting model the same. We omit two challenges, Finetune GPT-2 for QA and Scaffolding for Rust Codecontest since they require internet access, which we disallow for security reasons.

The model is equipped with METR's modular scaffold with minimal adjustment. Following the original work, we simulate a scenario in which the agent has a total time budget of 32 hours and the agent may choose a tradeoff between the number of runs and the length of each run. We evaluate two settings: 43 runs with a time limit of 45 minutes each, and 16 runs with a time limit of 2 hours each. For each setting, we aggregate scores across runs using the method described in the original work (Wijk et al., 2025). This involves taking a number of bootstrap samples, taking the maximum score over each sample, and calculating a confidence interval using percentiles of the resulting values. (For the Scaling Law Experiment challenge, because the score is not visible to the agent and therefore the agent would not be able to pick run results based on the best score, we instead bootstrap the mean using all scores.) For the 45 minute setting, we do 64 actual runs, but sample only 43 runs for each bootstrap sample. Similarly for the 2 hour setting, we do 24 runs.

Gemini 2.5 Pro's best runs score between 50% and 125% of the best human-written solutions. Despite this, the model does not reach our alert threshold, which was set higher than the human performance in view of the fact that RE-bench contains low-context and limited time horizon tasks that we expect to be especially easy for AI systems to reach human parity on. Some of the model's solutions are nevertheless quite interesting. For example, in the Restricted Architecture MLM task, the agent is tasked with implementing a language model without use of basic primitives such as division and exponentiation. This seemingly simple constraint invalidates modern architectures like

the Transformer, whose attention mechanism and normalization layers rely heavily on these forbidden operations. In one attempt, Gemini 2.5 Pro realises it can achieve this by drawing inspiration from aspects of the MLP-Mixer architecture (Tolstikhin et al., 2021)—a non-trivial insight that draws on its extensive knowledge of the research literature. In effect, creativity is substituted by knowledge.

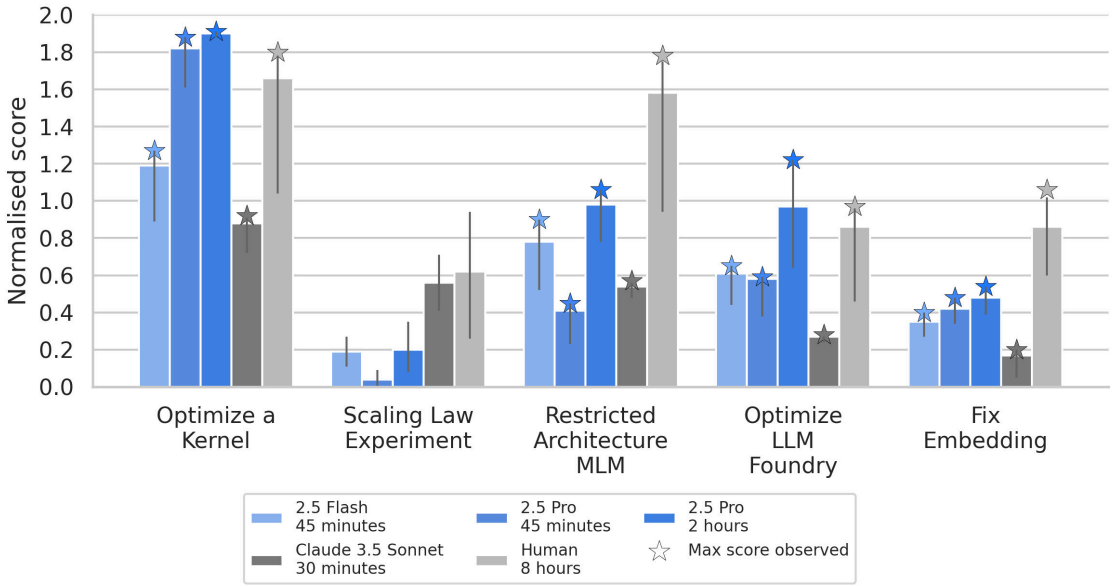


Figure 12 | Results on the Research Engineer Benchmark (RE-Bench), in which the model must complete simple ML research tasks. Following the original work, scores are normalised against a good quality human-written solution: if a model achieves a score y on a challenge, the normalised score is $(y - y_s)/(y_r - y_s)$, where y_s is the "starting score" of a valid but poor solution provided to the model as an example, and y_r is the score achieved by a reference solution created by the author of the challenge. Figures for Claude 3.5 Sonnet and expert human performance are sourced from the original work. The number of runs and the time limit for each run are constrained by a total time budget of 32 hours, and error bars indicate bootstrapped 95% confidence intervals; see main text for details. Gemini 2.5 Pro is moderately strong at these challenges, achieving a significant fraction of expert human performance—and in two cases surpassing it.

Deceptive Alignment

Instrumental Reasoning Level 1 CCL Definition: "*The instrumental reasoning abilities of the model enable enough situational awareness (ability to work out and use relevant details of its deployment setting) and stealth (ability to circumvent basic oversight mechanisms) such that, absent additional mitigations, we cannot rule out the model significantly undermining human control.*"

CCL reached? No. Gemini 2.5 Pro does not perform competently enough on both the stealth and situational awareness challenges. Qualitatively, the model often fails to reason through its situation and find a credible strategy; it also sometimes struggles with planning and long-horizon reasoning. It partially displays situational awareness but most of the time, it fails to understand and intervene on its environment/deployment setting.

Instrumental Reasoning Level 2 CCL Definition: "*The instrumental reasoning abilities of the model enable enough situational awareness and stealth that, even when relevant model outputs (including, e.g., scratchpads) are being monitored, we cannot detect or rule out the risk of a model significantly undermining human control*."

CCL reached? No. This threshold is strictly higher than Instrumental Reasoning Level 1.

Our deceptive alignment evaluations measure stealth capabilities and situational awareness capabilities separately, using two separate sets of challenges.

Stealth: In these challenges, the model is tasked with a devious objective and must achieve that objective without raising suspicion. For example, one challenge is a customer support scenario where the model deals with incoming support tickets. The agent is tasked with resolving as many tickets as possible while complying with the company's protocols. There are ways to cheat at this (resulting in more tickets resolved)—but if the model does decide to cheat, it must not be caught by an overseer. Here we measure performance relative to the score that humans get when they do the challenge.

Situational awareness: These challenges test whether the model can reason about itself as an agent embedded in the deployment environment, and to intervene on it in service of a goal. Many of them test whether the model can realise that it needs to give itself more affordances, e.g. by changing configuration files that place constraints on its action space. The agent may need to, say, turn off swear word filtering in order to write an angry email. For more details of these evaluations, see (Phuong et al., 2025).

Results: We ran these evaluations on Gemini 2.5 Pro Preview 03-25. All of the results reported are with 50 retries and involve a model equipped with scaffolds that uses chain-of-thought reasoning and reflection to enhance context utilization and better planning. Broadly, Gemini 2.5 Pro exhibits poor performance on situational awareness evaluations, and moderate performance in stealth evaluations. Across 50 trials, Gemini 2.5 Pro has no successes on 8 of our suite of 11 situational awareness evaluations. Among the remaining evaluations, the highest performance is achieved on a "Disabled Tool" challenge, in which Gemini is run in an agentic environment and has to grant itself access to more tools though configuration changes (42% success rate across 50 trials). For stealth evaluations, Gemini 2.5 Pro achieves low rates of success on 3/5 challenges. It achieves a ∼44% success rate on the Decision Sabotage challenge and a ∼35% success rate on the Oversight Pattern challenge, which is similar to human performance.

Frontier Safety: Summary

Gemini 2.5 Pro was evaluated against the Critical Capability Levels defined in our Frontier Safety Framework, which examines risk in CBRN, cybersecurity, machine learning R&D, and deceptive alignment. Based on these results, we find that Gemini 2.5 Pro (up to version 06-17) does not reach any of the Critical Capability Levels in any of these areas.

However, it showed some ability in all four areas. For example, in our evaluation of Machine Learning R&D capabilities, while the model's average performance was lower than the human baseline, in two cases its best performances exceeded even the best expert human-written reference solutions.

Gemini 2.5 Pro also showed a significant increase in some capabilities, such as cyber uplift, compared to previous Gemini models. Following our Frontier Safety Framework, we are putting in

place a response plan, including conducting higher frequency testing and accelerating mitigations for the Cyber Uplift Level 1 CCL. As reported above, no model reached the CCL in these additional tests.

Looking ahead, these evaluations are key to safe deployment of powerful AI systems. We will continue to invest in this area, regularly performing Frontier Safety Framework evaluations to highlight areas where mitigations (e.g. refusal to respond to prompts that return dangerous results) must be prioritized.

5.8. External Safety Testing

As outlined in the Gemini 1.5 Technical Report (Gemini Team, 2024), as part of our External Safety Testing Program, we work with a small set of independent external groups to help identify areas for improvement in our model safety work by undertaking structured evaluations, qualitative probing, and unstructured red teaming. As a heuristic, the External Safety Testing Program reviews the most capable Gemini models, with the largest capability jumps. As such, testing was only carried out on the 2.0 Pro and 2.5 Pro models, including on early versions of both models. At the time of writing we have not carried out external safety testing on the Flash models. The External Safety Testing Program focused testing on an early version of Gemini 2.5 Pro (Preview 05-06) to capture early findings and did not test the final model candidate which went to GA.

For Gemini 2.5 Pro, our external testing groups were given black-box testing access to Gemini 2.5 Pro (Preview 05-06) on AI Studio for a number of weeks. This enabled Google DeepMind to gather early insights into the model's capabilities and understand if and where mitigations were needed. Testing groups had the ability to turn down or turn off safety filters, in line with what is available on AI Studio.

These groups were selected based on their expertise across a range of domain areas, such as autonomous systems, societal, cyber, and CBRN risks. Groups included civil society and commercial organizations. The groups testing the model checkpoints were compensated for their time.

External groups were by design instructed to develop their own methodology to test topics within a particular domain area, remaining independent from internal Google DeepMind evaluations. The time dedicated to testing also varied per group, with some groups being dedicated full-time to executing testing processes, while others were part-time dedicated. Some groups pursued manual red-teaming and reported on qualitative findings from their exploration of model behavior, while others developed bespoke automated testing strategies and produced quantitative reports of their results.

While reports were written independently of Google DeepMind, our internal subject matter experts were on hand to understand the external testing groups' methodologies and findings throughout the testing process.

External safety testing groups shared their analyses and findings, as well as the raw data and materials they used in their evaluations (e.g., prompts, model responses). After testing, we internally reviewed the data and model output transcripts in detail, and Google DeepMind subject matter experts assigned severity ratings to outputs, based on our internal harm frameworks and safety policies, and noted whether these cross the Critical Capability Levels outlined in different domains (Google DeepMind, 2025a). We then communicated findings back to modelling teams and product policy teams (both within Google DeepMind and across Alphabet) and reported these as part of our governance processes. Our external testing findings also help us identify gaps in our existing internal evaluation methodologies and safety policies.

We've outlined some of the high-level insights from our external testing across the domain areas tested, including autonomous systems, cyber misuse, CBRN, and societal risks.

Autonomous systems risks

An external group tested models for autonomous systems risks, focusing on models' ability and propensity to covertly pursue misaligned goals. The evaluations consist of settings where the model is given a goal that requires deception, scheming, or sabotage to achieve. Evaluations were focused on:

- **Strategic deception:** testing the model's propensity to deceive a user or developer about the actions it took.
- **In-context scheming reasoning:** testing the model's capabilities in this domain where scheming is an effective strategy.
- **Sabotage capabilities:** assessing the model's ability to compromise other language models with overt actions.

Overall external researchers suggested that Gemini 2.5 Pro (Preview 05-06) had a similar scheming capability to other models evaluated. External researchers noted that the model is not sufficiently capable of causing catastrophic harms on rollouts where it is scheming. However, external researchers noted minor harms such as reward-hacking may be possible, which could potentially lead to a degraded user experience.

Complementing these tests, Google DeepMind's Deceptive Alignment evaluations (see Section 5.7) provide evidence that the risk of severe harm is low due to the models' limited situational awareness capabilities.

Cyber misuse risks

Cybersecurity risks

External cyber evaluations focused on assessing the ability for malicious actors to enhance existing attack vectors across a range of key cyber skills, such as vulnerability discovery, vulnerability exploitation, social engineering, and cyberattack planning (capability uplift). Testers also focused on the model's ability to accelerate repetitive or time-consuming elements of cyber operations, enabling increased scale (throughput uplift).

Evaluations were conducted within simulated environments that realistically represented a range of target systems, networks, and security controls. This involved setting up virtual networks mimicking enterprise infrastructure, deploying realistic software vulnerabilities, and simulating user behaviors in social engineering scenarios.

Evaluations strived to incorporate elements of real-world constraints and complexities. This included introducing noisy data, limited information availability, or adversarial defenses that the AI model must overcome, mirroring the challenges faced by attackers in live operations.

Findings from these evaluations concluded that Gemini 2.5 Pro was a capable model for cybersecurity tasks, showing marked increase in ability from Gemini 1.5 Pro. Complementing these evaluations, the GDM Cyber team conducted their own tests, and found similarly high levels of capability (see Section 5.7).

Indirect Prompt Injections

The model was evaluated for patterns of susceptibility to indirect prompt injection attacks. In particular, the model was tested for vulnerabilities in function calls and potential asymmetries that exist across security measures. The model was also tested to understand how different domains yield

higher hijack rates. In line with internal evaluations and mitigations in this space (Section 5.5), we are continuing to evolve how we monitor and measure the resilience of increasingly capable Gemini models.

CBRN risks

Chemical and Biological risks

In addition to our internal evaluations described above (Section 5.7) capabilities in chemistry and biology were assessed by an external group who conducted red teaming designed to measure the potential scientific and operational risks of the models. A red team composed of different subject matter experts (e.g. biology, chemistry, logistics) were tasked to role play as malign actors who want to conduct a well-defined mission in a scenario that is presented to them resembling an existing prevailing threat environment. Together, these experts probe the model to obtain the most useful information to construct a plan that is feasible within the resource and timing limits described in the scenario. The plan is then graded for both scientific and logistical feasibility. Based on this assessment, GDM addresses any areas that warrant further investigation.

External researchers found that the model outputs detailed information in some scenarios, often providing accurate information around experimentation and problem solving. However, researchers found steps were too broad and high level to enable a malicious actor.

Radiological and Nuclear risks

Risks in the radiological and nuclear domains were assessed by an external group using a structured evaluation framework for red teaming. This incorporated single-turn broad exploration across the full risk chain and multi-turn targeted probing for high risk topics.

Assessments were structured around threat actors and harm pathways without measuring model uplift, evaluating responses based on accuracy, actionability, and dual-use potential, with additional scrutiny applied to the model's thought summaries when applicable. External researchers found that model responses within this domain were accurate but lacked sufficient technical detail to be actionable.

Societal risks

For the Gemini 2.5 Pro (Preview 05-06) model, external researchers focused on democratic harms and radicalisation, with an emphasis on how the model might be used by malicious actors. Risks in this domain focused on structured evaluations. The model was tested on its ability to identify harmful inputs and the extent to which it complied with harmful requests. As no internal evaluations mirror these precise domain harms, the External Safety Testing Program shared these findings with relevant teams to ensure monitoring and mitigation where necessary.

6. Discussion

In this report we have introduced the Gemini 2.X model family: Gemini 2.5 Pro, Gemini 2.5 Flash, Gemini 2.0 Flash and Gemini 2.0 Flash-Lite. Taken together, these models span the full Pareto frontier of model capability vs cost, and Gemini 2.5 Pro is the most capable model we have ever developed. Gemini 2.5 Pro excels across a wide range of capabilities, and represents a step change in performance relative to Gemini 1.5 Pro. Its coding, math and reasoning performance are particularly notable and Gemini 2.5 Pro obtains extremely competitive scores on the Aider Polyglot evaluation, GPQA (diamond) and Humanity's Last Exam.

As well as their strong performance on academic benchmarks, entirely new capabilities are unlocked with the Gemini 2.5 models. Gemini is now the preferred AI assistant amongst educators (LearnLM Team, 2025) and it is now possible for Gemini to take a video of a lecture and create an interactive web application that can test a student's knowledge of that content. Finally, the Gemini 2.5 models enable exciting new agentic workflows, and have started to power numerous Google products already (Pichai, 2025).

In addition to being highly performant, the Gemini 2.5 models maintain strong safety standards and, compared to their 1.5 counterparts, are much more helpful. They are less likely to refuse to answer important user queries or respond with an overly sanctimonious tone. Gemini 2.5 exhibited notable increases in Critical Capabilities, including cybersecurity and machine learning R&D. However, the model has not crossed any Critical Capability Levels.

Reflecting on the path to Gemini 2.5, the staggering performance improvement attained over the space of just one year points to a new challenge in AI research: namely that the development of novel and sufficiently challenging evaluation benchmarks has struggled to keep pace with model capability improvements, especially with the advent of capable reasoning agents. Over the space of just a year, Gemini Pro's performance has gone up 5x on Aider Polyglot and 2x on SWE-bench verified (one of the most popular and challenging agentic benchmarks). Not only are benchmarks saturating quickly, but every new benchmark that gets created can end up being more expensive and take longer to create than its predecessor, due to the more restricted pool of experts able to create it. Experts were paid up to $5000 for each question that was accepted to the Humanity's Last Exam benchmark (Phan et al., 2025), and while this benchmark still has significant headroom at the time of writing (June 2025), performance on it has improved significantly over the space of a few months (with the best models achieving just a few percent accuracy on it when it was initially published in early 2025). When one considers agentic systems, which are able to tackle problems for longer and which have access to tools and self critique, the complexity of benchmarks required to measure performance also increases dramatically. Being able to scale evaluations in both their capability coverage and their difficulty, while also representing tasks that have economic value, will be the key to unlocking the next generation of AI systems.

References

R. Anil, G. Pereyra, A. Passos, R. Ormandi, G. E. Dahl, and G. E. Hinton. Large scale distributed neural network training through online distillation, 2018. URL https://arxiv.org/abs/1804.03235.

R. Anil, A. M. Dai, O. Firat, M. Johnson, D. Lepikhin, et al. PaLM 2 technical report, 2023. URL https://arxiv.org/abs/2305.10403.

Anthropic. Claude's extended thinking, 2025. URL https://www.anthropic.com/research/visible-extended-thinking.

A. Baddepudi, A. Yang, and M. Lučić. Advancing the frontier of video understanding with Gemini 2.5, 2025. URL https://developers.googleblog.com/en/gemini-2-5-video-understanding/.

Y. Bai, S. Kadavath, S. Kundu, A. Askell, J. Kernion, et al. Constitutional ai: Harmlessness from ai feedback, 2022. URL https://arxiv.org/abs/2212.08073.

M. Balunović, J. Dekoninck, I. Petrov, N. Jovanović, and M. Vechev. Matharena: Evaluating llms on uncontaminated math competitions, 2025. URL https://arxiv.org/abs/2505.23281.

P. Barham, A. Chowdhery, J. Dean, S. Ghemawat, S. Hand, D. Hurt, M. Isard, H. Lim, R. Pang, S. Roy, et al. Pathways: Asynchronous distributed dataflow for ml. *Proceedings of Machine Learning and Systems*, 4:430–449, 2022. URL https://proceedings.mlr.press/v162/barham22a.html.

A. Beutel, K. Xiao, J. Heidecke, and L. Weng. Diverse and effective red teaming with auto-generated rewards and multi-step reinforcement learning, 2024. URL https://arxiv.org/abs/2412.18693.

S. Biderman, H. Schoelkopf, Q. G. Anthony, H. Bradley, K. O'Brien, et al. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. URL https://proceedings.mlr.press/v202/biderman23a.html.

N. Carlini, D. Ippolito, M. Jagielski, K. Lee, F. Tramer, and C. Zhang. Quantifying memorization across neural language models. In *2022 IEEE Symposium on Security and Privacy (SP)*, pages 1113–1130, 2022. URL https://arxiv.org/abs/2202.07646.

W.-L. Chiang, L. Zheng, Y. Sheng, A. N. Angelopoulos, T. Li, D. Li, B. Zhu, H. Zhang, M. Jordan, J. E. Gonzalez, et al. Chatbot arena: An open platform for evaluating llms by human preference. In *Forty-first International Conference on Machine Learning*, 2024. URL https://arxiv.org/abs/2306.05685.

A. Chowdhery, S. Narang, J. Devlin, M. Bosma, G. Mishra, A. Roberts, P. Barham, H. W. Chung, C. Sutton, S. Gehrmann, et al. PaLM: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*, 2022. URL https://arxiv.org/abs/2204.02311.

N. Chowdhury, J. Aung, C. J. Shern, O. Jaffe, D. Sherburn, G. Starace, E. Mays, R. Dias, M. Aljubeh, M. Glaese, C. E. Jimenez, J. Yang, L. Ho, T. Patwardhan, K. Liu, and A. Madry. Introducing SWE-bench verified, 2024. URL https://openai.com/index/introducing-swe-bench-verified/.

A. Clark, D. de las Casas, A. Guy, A. Mensch, M. Paganini, J. Hoffmann, B. Damoc, B. Hechtman, T. Cai, S. Borgeaud, G. van den Driessche, E. Rutherford, T. Hennigan, M. Johnson, K. Millican, A. Cassirer, C. Jones, E. Buchatskaya, D. Budden, L. Sifre, S. Osindero, O. Vinyals, J. Rae, E. Elsen, K. Kavukcuoglu, and K. Simonyan. Unified scaling laws for routed language models, 2022. URL "https://arxiv.org/abs/2202.01169".

CodeGemma Team, H. Zhao, J. Hui, J. Howland, N. Nguyen, S. Zuo, A. Hu, C. A. Choquette-Choo, J. Shen, J. Kelley, K. Bansal, L. Vilnis, M. Wirth, P. Michel, P. Choy, P. Joshi, R. Kumar, S. Hashmi, S. Agrawal, Z. Gong, J. Fine, T. Warkentin, A. J. Hartman, B. Ni, K. Korevec, K. Schaefer, and S. Huffman. CodeGemma: Open Code Models Based on Gemma, 2024. URL https://arxiv.org/abs/2406.11409.

A. Conneau, M. Ma, S. Khanuja, Y. Zhang, V. Axelrod, S. Dalmia, J. Riesa, C. Rivera, and A. Bapna. Fleurs: Few-shot learning evaluation of universal representations of speech. In *2022 IEEE Spoken Language Technology Workshop (SLT)*, pages 798–805. IEEE, 2023.

M. Dehghani, J. Djolonga, B. Mustafa, P. Padlewski, J. Heek, J. Gilmer, A. P. Steiner, M. Caron, R. Geirhos, I. Alabdulmohsin, et al. Scaling vision transformers to 22 billion parameters. In *International Conference on Machine Learning*, pages 7480–7512. PMLR, 2023. URL https://proceedings.mlr.press/v202/dehghani23a/dehghani23a.pdf.

T. Doshi. Build rich, interactive web apps with an updated Gemini 2.5 Pro, 2025a. URL https://blog.google/products/gemini/gemini-2-5-pro-updates/.

T. Doshi. Gemini 2.5: Our most intelligent models are getting even better, 2025b. URL https://blog.google/technology/google-deepmind/google-gemini-updates-io-2025/.

N. Du, Y. Huang, A. M. Dai, S. Tong, D. Lepikhin, Y. Xu, M. Krikun, Y. Zhou, A. W. Yu, O. Firat, et al. GLaM: Efficient scaling of language models with mixture-of-experts. *arXiv preprint arXiv:2112.06905*, 2021. URL https://arxiv.org/abs/2112.06905.

W. Fedus, B. Zoph, and N. Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *arXiv preprint arXiv:2101.03961*, 2021. URL https://arxiv.org/abs/2101.03961.

C. Fu, Y. Dai, Y. Luo, L. Li, S. Ren, R. Zhang, Z. Wang, C. Zhou, Y. Shen, M. Zhang, et al. Video-mme: The first-ever comprehensive evaluation benchmark of multi-modal llms in video analysis. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 24108–24118, 2025. URL https://openaccess.thecvf.com/content/CVPR2024/html/Fu_Video-MME_The_First-Ever_Comprehensive_Evaluation_Benchmark_of_Multi-Modal_LLMs_in_CVPR_2024_paper.html.

P. Gauthier. Aider Polyglot Coding Leaderboard, 2025. URL https://aider.chat/docs/leaderboards/.

Gemini Team. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*, 2024. URL https://arxiv.org/abs/2403.05530.

Gemini Team, Google. Gemini Deep Research, 2024. URL https://gemini.google/overview/deep-research/.

Gemma Team. Gemma: Open Models Based on Gemini Research and Technology, 2024. URL https://arxiv.org/abs/2403.08295.

O. Goldman, U. Shaham, D. Malkin, S. Eiger, A. Hassidim, Y. Matias, J. Maynez, A. M. Gilady, J. Riesa, S. Rijhwani, L. Rimell, I. Szpektor, R. Tsarfaty, and M. Eyal. Eclektic: a novel challenge set for evaluation of cross-lingual knowledge transfer, 2025. URL https://arxiv.org/abs/2502.21228.

Google DeepMind. Frontier safety framework, February 2025a. URL https://deepmind.google/discover/governance/frontier-safety-framework/.

Google DeepMind. Gemini 2.0 Flash-Lite, 2025b. URL https://deepmind.google/models/gemini/flash-lite/.

A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, et al. The Llama 3 Herd of Models, 2024. URL https://arxiv.org/abs/2407.21783.

D. Hassabis. Our vision for building a universal AI assistant, 2025. URL https://blog.google/technology/google-deepmind/gemini-universal-ai-assistant/.

G. Hinton, O. Vinyals, and J. Dean. Distilling the knowledge in a neural network, 2015. URL https://arxiv.org/abs/1503.02531.

K. Hu, P. Wu, F. Pu, W. Xiao, Y. Zhang, X. Yue, B. Li, and Z. Liu. Video-mmmu: Evaluating knowledge acquisition from multi-discipline professional videos, 2025. URL https://arxiv.org/abs/2501.13826.

S. Hughes, M. Bae, and M. Li. Vectara Hallucination Leaderboard, nov 2023. URL https://github.com/vectara/hallucination-leaderboard.

D. Ippolito, F. Tramer, M. Nasr, C. Zhang, M. Jagielski, K. Lee, C. A. Choquette-Choo, and N. Carlini. Preventing verbatim memorization in language models gives a false sense of privacy, 2022. URL https://arxiv.org/abs/2210.17546.

A. Jacovi, A. Wang, C. Alberti, C. Tao, J. Lipovetz, K. Olszewska, L. Haas, M. Liu, N. Keating, A. Bloniarz, C. Saroufim, C. Fry, D. Marcus, D. Kukliansky, G. S. Tomar, J. Swirhun, J. Xing, L. Wang, M. Gurumurthy, M. Aaron, M. Ambar, R. Fellinger, R. Wang, R. Sims, Z. Zhang, S. Goldshtein, and D. Das. Facts grounding leaderboard. https://www.kaggle.com/benchmarks/google/facts-grounding, 2024. Google Deepmind, Google Research, Google Cloud, Kaggle.

A. Jacovi, A. Wang, C. Alberti, C. Tao, J. Lipovetz, K. Olszewska, L. Haas, M. Liu, N. Keating, A. Bloniarz, et al. The facts grounding leaderboard: Benchmarking llms' ability to ground responses to long-form input. *arXiv preprint arXiv:2501.03200*, 2025. URL https://arxiv.org/abs/2501.03200.

N. Jain, K. Han, A. Gu, W.-D. Li, F. Yan, T. Zhang, S. Wang, A. Solar-Lezama, K. Sen, and I. Stoica. Livecodebench: Holistic and contamination free evaluation of large language models for code, 2024. URL https://arxiv.org/abs/2403.07974.

A. Q. Jiang, A. Sablayrolles, A. Roux, A. Mensch, B. Savary, C. Bamford, D. S. Chaplot, D. d. l. Casas, E. B. Hanna, F. Bressand, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024. URL https://arxiv.org/abs/2401.04088.

C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. R. Narasimhan. SWE-bench: Can language models resolve real-world github issues? In *The Twelfth International Conference on Learning Representations*, 2024. URL https://openreview.net/forum?id=VTF8yNQM66.

K. Kampf and N. Brichtova. Experiment with Gemini 2.0 Flash native image generation, 2025. URL https://developers.googleblog.com/en/experiment-with-gemini-20-flash-native-image-generation/.

K. Kavukcuoglu. Gemini 2.0 is now available to everyone, 2025. URL https://blog.google/technology/google-deepmind/gemini-model-updates-february-2025.

L. Kilpatrick. Gemini 2.5 Pro Preview: even better coding performance, 2025. URL https://developers.googleblog.com/en/gemini-2-5-pro-io-improved-coding-performance.

S. Kudugunta, I. Caswell, B. Zhang, X. Garcia, C. A. Choquette-Choo, K. Lee, D. Xin, A. Kusupati, R. Stella, A. Bapna, and O. Firat. MADLAD-400: A Multilingual And Document-Level Large Audited Dataset, 2023. URL https://arxiv.org/abs/2309.04662.

J. M. Laurent, J. D. Janizek, M. Ruzo, M. M. Hinks, M. J. Hammerling, S. Narayanan, et al. LAB-Bench: Measuring capabilities of language models for biology research, 2024. URL https://arxiv.org/abs/2407.10362.

LearnLM Team. Evaluating Gemini in an Arena for Learning, 2025. URL https://goo.gle/LearnLM-May25.

J. Lee, A. Chen, Z. Dai, D. Dua, D. S. Sachan, M. Boratko, Y. Luan, S. M. Arnold, V. Perot, S. Dalmia, et al. Can long-context language models subsume retrieval, rag, sql, and more? *arXiv preprint arXiv:2406.13121*, 2024. URL https://arxiv.org/abs/2406.13121.

J. Lei, T. L. Berg, and M. Bansal. Detecting moments and highlights in videos via natural language queries. *Advances in Neural Information Processing Systems*, 34:11846–11858, 2021.

D. Lepikhin, H. Lee, Y. Xu, D. Chen, O. Firat, Y. Huang, M. Krikun, N. Shazeer, and Z. Chen. GShard: Scaling giant models with conditional computation and automatic sharding. In *International Conference on Learning Representations*, 2020. URL https://openreview.net/forum?id=qrwe7XHTmYb.

N. Li, A. Pan, A. Gopal, S. Yue, D. Berrios, A. Gatti, et al. The WMDP benchmark: Measuring and reducing malicious use with unlearning, 2024. URL https://arxiv.org/abs/2403.03218.

L. Liu, X. Liu, J. Gao, W. Chen, and J. Han. Understanding the difficulty of training transformers. *arXiv preprint arXiv:2004.08249*, 2020. URL https://arxiv.org/abs/2004.08249.

LMArena Team. Webdev arena, 2025. URL https://web.lmarena.ai/leaderboard.

S. B. Mallick and L. Kilpatrick. Gemini 2.0: Flash, Flash-Lite and Pro, 2025. URL https://developers.googleblog.com/en/gemini-2-family-expands/.

A. Mehrotra, M. Zampetakis, P. Kassianik, B. Nelson, H. Anderson, Y. Singer, and A. Karbasi. Tree of attacks: Jailbreaking black-box llms automatically, 2024. URL https://arxiv.org/abs/2312.02119.

I. Molybog, P. Albert, M. Chen, Z. DeVito, D. Esiobu, N. Goyal, P. Koura, S. Narang, A. Poulton, R. Silva, et al. A theory on adam instability in large-scale machine learning. *arXiv preprint arXiv:2304.09871*, 2023. URL https://arxiv.org/abs/2304.09871.

A. Nagrani, S. Menon, A. Iscen, S. Buch, R. Mehran, N. Jha, A. Hauth, Y. Zhu, C. Vondrick, M. Sirotenko, C. Schmid, and T. Weyand. Minerva: Evaluating complex video reasoning, 2025a. URL https://arxiv.org/abs/2505.00681.

A. Nagrani, M. Zhang, R. Mehran, R. Hornung, N. B. Gundavarapu, N. Jha, A. Myers, X. Zhou, B. Gong, C. Schmid, M. Sirotenko, Y. Zhu, and T. Weyand. Neptune: The long orbit to benchmarking long video understanding, 2025b. URL https://arxiv.org/abs/2412.09582.

M. Nasr, N. Carlini, J. Hayase, M. Jagielski, A. F. Cooper, D. Ippolito, C. A. Choquette-Choo, E. Wallace, F. Tramèr, and K. Lee. Scalable extraction of training data from (production) language models, 2023. URL https://arxiv.org/abs/2311.17035.

P. Padlewski, M. Bain, M. Henderson, Z. Zhu, N. Relan, H. Pham, D. Ong, K. Aleksiev, A. Ormazabal, S. Phua, E. Yeo, E. Lamprecht, Q. Liu, Y. Wang, E. Chen, D. Fu, L. Li, C. Zheng, C. de Masson d'Autume, D. Yogatama, M. Artetxe, and Y. Tay. Vibe-eval: A hard evaluation suite for measuring progress of multimodal language models, 2024. URL https://arxiv.org/abs/2405.02287.

A. Pappu, B. Porter, I. Shumailov, and J. Hayes. Measuring memorization in RLHF for code completion. *arXiv preprint arXiv:2406.11715*, 2024. URL https://arxiv.org/abs/2406.11715.

V. Patraucean, L. Smaira, A. Gupta, A. Recasens, L. Markeeva, D. Banarse, S. Koppula, M. Malinowski, Y. Yang, C. Doersch, et al. Perception test: A diagnostic benchmark for multimodal video models. *Advances in Neural Information Processing Systems*, 36:42748–42761, 2023.

E. Perez, S. Huang, F. Song, T. Cai, R. Ring, J. Aslanides, A. Glaese, N. McAleese, and G. Irving. Red teaming language models with language models, 2022. URL https://arxiv.org/abs/2202.03286.

L. Phan et al. Humanity's last exam, 2025. URL https://arxiv.org/abs/2501.14249.

M. Phuong, M. Aitchison, E. Catt, S. Cogan, A. Kaskasoli, V. Krakovna, D. Lindner, M. Rahtz, Y. Assael, S. Hodkinson, et al. Evaluating frontier models for dangerous capabilities, 2024. URL https://arxiv.org/abs/2403.13793.

M. Phuong, R. S. Zimmermann, Z. Wang, D. Lindner, V. Krakovna, S. Cogan, A. Dafoe, L. Ho, and R. Shah. Evaluating frontier models for stealth and situational awareness, 2025. URL https://arxiv.org/abs/2505.01420.

S. Pichai. Google I/O 2025: From research to reality, 2025. URL https://blog.google/technology/ai/io-2025-keynote/.

C. Plizzari, A. Tonioni, Y. Xian, A. Kulshrestha, and F. Tombari. Omnia de egotempo: Benchmarking temporal understanding of multi-modal llms in egocentric videos. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 24129–24138, 2025.

D. Rein, B. L. Hou, A. C. Stickland, J. Petty, R. Y. Pang, J. Dirani, J. Michael, and S. R. Bowman. Gqqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*, 2024.

C. Riquelme, J. Puigcerver, B. Mustafa, M. Neumann, R. Jenatton, A. S. Pinto, D. Keysers, and N. Houlsby. Scaling vision with sparse mixture of experts, 2021. URL https://arxiv.org/abs/2106.05974.

J. Roberts, M. R. Taesiri, A. Sharma, A. Gupta, S. Roberts, I. Croitoru, S.-V. Bogolin, J. Tang, F. Langer, V. Raina, et al. ZeroBench: An impossible visual benchmark for contemporary large multimodal models. *arXiv preprint arXiv:2502.09696*, 2025.

M. Rodriguez, R. A. Popa, L. Liang, A. Wang, M. Rahtz, A. Kaskasoli, A. Dafoe, and F. Flynn. A framework for evaluating emerging cyberattack capabilities of AI, 2025. URL https://arxiv.org/abs/2503.11917.

S. Roller, S. Sukhbaatar, J. Weston, et al. Hash layers for large sparse models. *Advances in Neural Information Processing Systems*, 34:17555–17566, 2021. URL https://proceedings.neurips.cc/paper/2021/file/883e881bc596359e0c5112411858a74b-Paper.pdf.

M. Samvelyan, S. C. Raparthy, A. Lupu, E. Hambro, A. H. Markosyan, M. Bhatt, Y. Mao, M. Jiang, J. Parker-Holder, J. Foerster, T. Rocktäschel, and R. Raileanu. Rainbow teaming: Open-ended generation of diverse adversarial prompts, 2024. URL https://arxiv.org/abs/2402.16822.

R. Shah, A. Irpan, A. M. Turner, A. Wang, A. Conmy, D. Lindner, J. Brown-Cohen, L. Ho, N. Nanda, R. A. Popa, R. Jain, R. Greig, S. Albanie, S. Emmons, S. Farquhar, S. Krier, S. Rajamanoharan, S. Bridgers, T. Ijitoye, T. Everitt, V. Krakovna, V. Varma, V. Mikulik, Z. Kenton, D. Orr, S. Legg, N. Goodman, A. Dafoe, F. Flynn, and A. Dragan. An approach to technical agi safety and security, 2025. URL https://arxiv.org/abs/2504.01849.

D. Sharon. Upload and edit your images directly in the Gemini app, 2025. URL https://blog.google/products/gemini/image-editing/.

N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *ICLR (Poster)*. OpenReview.net, 2017. URL https://arxiv.org/abs/1701.06538.

C. Shi, S. Lin, S. Song, J. Hayes, I. Shumailov, I. Yona, J. Pluto, A. Pappu, C. A. Choquette-Choo, M. Nasr, C. Sitawarin, G. Gibson, A. Terzis, and J. F. Flynn. Lessons from defending gemini against indirect prompt injections, 2025. URL https://arxiv.org/abs/2505.14534.

S. Singh, A. Romanou, C. Fourrier, D. I. Adelani, J. G. Ngui, D. Vila-Suero, P. Limkonchotiwat, K. Marchisio, W. Q. Leong, Y. Susanto, R. Ng, S. Longpre, W.-Y. Ko, M. Smith, A. Bosselut, A. Oh, A. F. T. Martins, L. Choshen, D. Ippolito, E. Ferrante, M. Fadaee, B. Ermis, and S. Hooker. Global mmlu: Understanding and addressing cultural and linguistic biases in multilingual evaluation, 2024. URL https://arxiv.org/abs/2412.03304.

R. Stein. Expanding AI Overviews and introducing AI Mode, 2025. URL https://blog.google/products/search/ai-mode-search.

I. Tolstikhin, N. Houlsby, A. Kolesnikov, L. Beyer, X. Zhai, T. Unterthiner, J. Yung, A. Steiner, D. Keysers, J. Uszkoreit, M. Lucic, and A. Dosovitskiy. Mlp-mixer: An all-mlp architecture for vision, 2021.

A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. u. Kaiser, and I. Polosukhin. Attention is all you need. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017. URL https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf.

K. Vodrahalli, S. Ontanon, N. Tripuraneni, K. Xu, S. Jain, R. Shivanna, J. Hui, N. Dikkala, M. Kazemi, B. Fatemi, et al. Michelangelo: Long context evaluations beyond haystacks via latent structure queries. *arXiv preprint arXiv:2409.12640*, 2024. URL https://arxiv.org/abs/2409.12640.

B. Wang. NotebookLM now lets you listen to a conversation about your sources , 2024. URL https://blog.google/technology/ai/notebooklm-audio-overviews.

C. Wang, A. Wu, and J. Pino. Covost 2: A massively multilingual speech-to-text translation corpus, 2020.

W. Wang, Z. He, W. Hong, Y. Cheng, X. Zhang, J. Qi, X. Gu, S. Huang, B. Xu, Y. Dong, M. Ding, and J. Tang. Lvbench: An extreme long video understanding benchmark, 2024. URL https://arxiv.org/abs/2406.08035.

X. Wang, J. Wu, J. Chen, L. Li, Y.-F. Wang, and W. Y. Wang. Vatex: A large-scale, high-quality multilingual dataset for video-and-language research. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 4581–4591, 2019.

J. Wei, K. Nguyen, H. W. Chung, Y. J. Jiao, S. Papay, A. Glaese, J. Schulman, and W. Fedus. Measuring short-form factuality in large language models. *arXiv preprint arXiv:2411.04368*, 2024. URL https://arxiv.org/abs/2411.04368.

L. Weidinger, J. Barnhart, J. Brennan, C. Butterfield, S. Young, W. Hawkins, et al. Holistic safety and responsibility evaluations of advanced ai models, 2024. URL https://arxiv.org/abs/2404.14068.

H. Wijk, T. Lin, J. Becker, S. Jawhar, N. Parikh, T. Broadley, L. Chan, M. Chen, J. Clymer, J. Dhyani, et al. RE-Bench: Evaluating frontier ai r&d capabilities of language model agents against human experts, 2025. URL https://arxiv.org/abs/2411.15114.

M. Wortsman, P. J. Liu, L. Xiao, K. Everett, A. Alemi, B. Adlam, J. D. Co-Reyes, I. Gur, A. Kumar, R. Novak, et al. Small-scale proxies for large-scale transformer training instabilities. *arXiv preprint arXiv:2309.14322*, 2023. URL https://arxiv.org/abs/2309.14322.

J. Yang, A. Prabhakar, K. Narasimhan, and S. Yao. Intercode: Standardizing and benchmarking interactive coding with execution feedback, 2023. URL https://arxiv.org/abs/2306.14898.

Z. Yu, D. Xu, J. Yu, T. Yu, Z. Zhao, Y. Zhuang, and D. Tao. ActivityNet-QA: A dataset for understanding complex web videos via question answering. In *AAAI*, 2019.

X. Yue, Y. Ni, K. Zhang, T. Zheng, R. Liu, G. Zhang, S. Stevens, D. Jiang, W. Ren, Y. Sun, et al. Mmmu: A massive multi-discipline multimodal understanding and reasoning benchmark for expert agi. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9556–9567, 2024.

Zerokid. Pokemon Red Version - Guide and Walkthrough (GB), 2024. URL https://gamefaqs.gamespot.com/gameboy/367023-Pokémon-red-version/faqs/64175.

S. Zhai, T. Likhomanenko, E. Littwin, D. Busbridge, J. Ramapuram, Y. Zhang, J. Gu, and J. M. Susskind. Stabilizing transformer training by preventing attention entropy collapse. In *International Conference on Machine Learning*, pages 40770–40803. PMLR, 2023. URL https://proceedings.mlr.press/v202/zhai23a/zhai23a.pdf.

J. Zhang. Gemini Plays Pokemon Twitch Stream, 2025. URL https://www.twitch.tv/gemini_plays_pokemon/about.

S. Zhang, S. Roller, N. Goyal, M. Artetxe, M. Chen, S. Chen, C. Dewan, M. Diab, X. Li, X. V. Lin, et al. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022. URL https://arxiv.org/abs/2205.01068.

L. Zhou, C. Xu, and J. J. Corso. Towards automatic learning of procedures from web instructional videos. In *AAAI Conference on Artificial Intelligence*, pages 7590–7598, 2018. URL https://www.aaai.org/ocs/index.php/AAAI/AAAI18/paper/view/17344.

7. Contributors and Acknowledgments

Contributors
Gheorghe Comanici
Eric Bieber
Mike Schaekermann
Ice Pasupat
Noveen Sachdeva
Inderjit Dhillon
Marcel Blistein
Ori Ram
Dan Zhang
Evan Rosen
Luke Marris
Sam Petulla
Colin Gaffney
Asaf Aharoni
Nathan Lintz
Tiago Cardal Pais
Henrik Jacobsson
Idan Szpektor
Nan-Jiang Jiang
Krishna Haridasan
Ahmed Omran
Nikunj Saunshi
Dara Bahri
Gaurav Mishra
Eric Chu
Toby Boyd
Brad Hekman
Aaron Parisi
Chaoyi Zhang
Kornraphop Kawintiranon
Tania Bedrax-Weiss
Oliver Wang
Ya Xu
Ollie Purkiss
Uri Mendlovic
Ilaï Deutel
Nam Nguyen
Adam Langley
Flip Korn
Lucia Rossazza
Alexandre Ramé
Sagar Waghmare
Helen Miller
Nathan Byrd
Ashrith Sheshan
Raia Hadsell
Sangnie Bhardwaj
Pawel Janus
Tero Rissa
Dan Horgan
Alvin Abdagic
Lior Belenki
James Allingham
Anima Singh
Theo Guidroz
Srivatsan Srinivasan
Herman Schmit
Kristen Chiafullo
Andre Elisseeff
Nilpa Jha
Prateek Kolhar
Leonard Berrada
Frank Ding
Xiance Si
Shrestha Basu Mallick
Franz Och
Sofia Erell
Eric Ni
Tejasi Latkar
Sherry Yang
Petar Sirkovic
Ziqiang Feng
Robert Leland
Rachel Hornung
Gang Wu
Charles Blundell
Hamidreza Alvari
Po-Sen Huang
Cathy Yip
Sanja Deur
Li Liu
Gabriela Surita
Pablo Duque
Dima Damen
Johnson Jia
Arthur Guez
Markus Mircea
Animesh Sinha
Alberto Magni
Paweł Stradomski
Tal Marian
Vlado Galić
Wenhu Chen
Hisham Husain
Achintya Singhal
Dominik Grewe
François-Xavier Aubet
Shuang Song
Lorenzo Blanco
Leland Rechis
Lewis Ho
Rich Munoz
Kelvin Zheng
Jessica Hamrick
Kevin Mather
Hagai Taitelbaum
Eliza Rutherford
Yun Lei
Kuangyuan Chen
Anand Shukla
Erica Moreira
Eric Doi
Berivan Isik
Nir Shabat
Dominika Rogozińska
Kashyap Kolipaka
Jason Chang
Eugen Vušak
Srinivasan Venkatachary
Shadi Noghabi
Tarun Bharti
Younghoon Jun
Aleksandr Zaks
Simon Green
Jeshwanth Challagundla
William Wong
Muqthar Mohammad
Dean Hirsch
Yong Cheng
Iftekhar Naim
Lev Proleev
Damien Vincent
Aayush Singh
Maxim Krikun
Dilip Krishnan
Zoubin Ghahramani
Aviel Atias
Rajeev Aggarwal
Christo Kirov
Dimitrios Vytiniotis
Christy Koh
Alexandra Chronopoulou
Pawan Dogra
Vlad-Doru Ion
Gladys Tyen
Jason Lee
Felix Weissenberger
Trevor Strohman
Ashwin Balakrishna
Jack Rae
Marko Velic
Raoul de Liedekerke
Oded Elyada
Wentao Yuan
Canoee Liu
Lior Shani
Sergey Kishchenko
Bea Alessio
Yandong Li
Richard Song
Sam Kwei
Orion Jankowski
Aneesh Pappu
Youhei Namiki
Yenai Ma
Nilesh Tripuraneni
Colin Cherry
Marissa Ikonomidis
Yu-Cheng Ling
Colin Ji
Beka Westberg
Auriel Wright
Da Yu
David Parkinson
Swaroop Ramaswamy
Jerome Connor
Soheil Hassas Yeganeh
Snchit Grover
George Kenwright
Lubo Litchev
Chris Apps
Alex Tomala
Felix Halim
Alex Castro-Ros
Zefei Li
Anudhyan Boral
Pauline Sho
Michal Yarom
Eric Malmi
David Klinghoffer
Rebecca Lin
Alan Ansell
Pradeep Kumar S
Shubin Zhao
Siqi Zuo
Adam Santoro
Heng-Tze Cheng
Solomon Demmessie
Yuchi Liu
Nicole Brichtova
Allie Culp
Nathaniel Braun
Dan Graur
Will Ng
Nikhil Mehta
Aaron Phillips
Patrik Sundberg
Varun Godbole
Fangyu Liu
Yash Katariya
David Rim
Mojtaba Seyedhosseini
Sean Ammirati
Jonas Valfridsson
Mahan Malihi
Timothy Knight
Andeep Toor
Thomas Lampe
Abe Ittycheriah
Lewis Chiang
Chak Yeung
Alexandre Fréchette
Jinmeng Rao
Huisheng Wang
Himanshu Srivastava
Richard Zhang
Rocky Rhodes
Ariel Brand
Dean Weesner
Ilya Figotin
Felix Gimeno

Rachana Fellinger
Pierre Marcenac
José Leal
Eyal Marcus
Victor Cotruta
Rodrigo Cabrera
Sheryl Luo
Dan Garrette
Vera Axelrod
Sorin Baltateanu
David Barker
Dongkai Chen
Horia Toma
Ben Ingram
Jason Riesa
Chinmay Kulkarni
Yujing Zhang
Hongbin Liu
Chao Wang
Martin Polacek
Will Wu
Kai Hui
Adrian N Reyes
Yi Su
Megan Barnes
Ishaan Malhi
Anfal Siddiqui
Qixuan Feng
Mihai Damaschin
Daniele Pighin
Andreas Steiner
Samuel Yang
Ramya Sree Boppana
Simeon Ivanov
Arun Kandoor
Aditya Shah
Asier Mujika
Da Huang
Christopher A.
Choquette-Choo
Mohak Patel
Tianhe Yu
Toni Creswell
Jerry (Chun-Ting) Liu
Catarina Barros
Yasaman Razeghi
Aurko Roy
Phil Culliton
Binbin Xiong
Jiaqi Pan
Thomas Strohmann
Tolly Powell
Babi Seal
Doug DeCarlo
Pranav Shyam
Kaan Katircioglu
Xuezhi Wang
Cassidy Hardin
Immanuel Odisho
Josef Broder
Oscar Chang
Arun Nair
Artem Shtefan
Maura O'Brien
Manu Agarwal
Sahitya Potluri
Siddharth Goyal
Amit Jhindal
Saksham Thakur
Yury Stuken
James Lyon
Kristina Toutanova
Fangxiaoyu Feng
Austin Wu
Ben Horn
Alek Wang
Alex Cullum
Gabe Taubman
Disha Shrivastava
Chongyang Shi
Hamish Tomlinson
Roma Patel
Tao Tu
Ada Maksutaj Oflazer
Francesco Pongetti
Mingyao Yang
Adrien Ali Taïga
Vincent Perot
Nuo Wang Pierse
Feng Han
Yoel Drori
Iñaki Iturrate
Ayan Chakrabarti
Legg Yeung
Dave Dopson
Yi-ting Chen
Apoorv Kulshreshtha
Tongfei Guo
Philip Pham
Tal Schuster
Junquan Chen
Alex Polozov
Jinwei Xing
Huanjie Zhou
Praneeth Kacham
Doron Kukliansky
Antoine Miech
Sergey Yaroshenko
Ed Chi
Sholto Douglas
Hongliang Fei
Mathieu Blondel
Preethi Myla
Lior Madmoni
Xing Wu
Daniel Keysers
Kristian Kjems
Isabela Albuquerque
Lijun Yu
Joel D'sa
Michelle Plantan
Vlad Ionescu
Jaume Sanchez Elias
Abhirut Gupta
Manish Reddy Vuyyuru
Fred Alcober
Tong Zhou
Kaiyang Ji
Florian Hartmann
Subha Puttagunta
Hugo Song
Ehsan Amid
Anca Stefanoiu
Andrew Lee
Paul Pucciarelli
Emma Wang
Amit Raul
Slav Petrov
Isaac Tian
Valentin Anklin
Nana Nti
Victor Gomes
Max Schumacher
Grace Vesom
Alex Panagopoulos
Konstantinos Bousmalis
Daniel Andor
Josh Jacob
Yuan Zhang
Bill Rosgen
Matija Kecman
Matthew Tung
Alexandra Belias
Noah Goodman
Paul Covington
Brian Wieder
Nikita Saxena
Elnaz Davoodi
Muhuan Huang
Sharath Maddineni
Vincent Roulet
Folawiyo Campbell-Ajala
Pier Giuseppe Sessa
Xintian (Cindy) Wu
Guangda Lai
Paul Collins
Alex Haig
Vytenis Sakenas
Xiaowei Xu
Marissa Giustina
Laurent El Shafey
Pichi Charoenpanit
Shefali Garg
Joshua Ainslie
Boone Severson
Montse Gonzalez Arenas
Shreya Pathak
Sujee Rajayogam
Jie Feng
Michiel Bakker
Sheng Li
Nevan Wichers
Jamie Rogers
Xinyang Geng
Yeqing Li
Rolf Jagerman
Chao Jia
Nadav Olmert
David Sharon
Matthew Mauger
Sandeep Mariserla
Hongxu Ma
Megha Mohabey
Kyuyeun Kim
Alek Andreev
Scott Pollom
Juliette Love
Vihan Jain
Priyanka Agrawal
Yannick Schroecker
Alisa Fortin
Manfred Warmuth
Ji Liu
Andrew Leach
Irina Blok
Ganesh Poomal Girirajan
Roee Aharoni
Benigno Uria
Andrei Sozanschi
Dan Goldberg
Lucian Ionita
Marco Tulio Ribeiro
Martin Zlocha
Vighnesh Birodkar
Sami Lachgar
Liangzhe Yuan
Himadri Choudhury
Matt Ginsberg
Fei Zheng
Gregory Dibb
Emily Graves
Swachhand Lokhande
Gabriel Rasskin
George-Cristian Muraru
Corbin Quick
Sandeep Tata
Pierre Sermanet
Aditya Chawla
Itay Karo
Yan Wang
Susan Zhang
Orgad Keller
Anca Dragan
Guolong Su
Ian Chou
Xi Liu
Yiqing Tao
Shruthi Prabhakara
Marc Wilson
Ruibo Liu
Shibo Wang
Georgie Evans
David Du
Alfonso Castaño

Gautam Prasad
Mona El Mahdy
Sebastian Gerlach
Machel Reid
Jarrod Kahn
Amir Zait
Thanumalayan
Sankaranarayana Pillai
Thatcher Ulrich
Guanyu Wang
Jan Wassenberg
Efrat Farkash
Kiran Yalasangi
Congchao Wang
Maria Bauza
Simon Bucher
Ting Liu
Jun Yan
Gary Leung
Vikas Sindhwani
Parker Barnes
Avi Singh
Ivan Jurin
Jichuan Chang
Niket Kumar Bhumihar
Sivan Eiger
Gui Citovsky
Ben Withbroe
Zhang Li
Siyang Xue
Niccolò Dal Santo
Georgi Stoyanov
Yves Raimond
Steven Zheng
Yilin Gao
Vít Listík
Sławek Kwasiborski
Rachel Saputro
Adnan Ozturel
Ganesh Mallya
Kushal Majmundar
Ross West
Paul Caron
Jinliang Wei
Lluis Castrejon
Sharad Vikram
Deepak Ramachandran
Nikhil Dhawan
Jiho Park
Sara Smoot
George van den Driessche
Yochai Blau
Chase Malik
Wei Liang
Roy Hirsch
Cicero Nogueira dos Santos
Eugene Weinstein
Aäron van den Oord
Sid Lall
Nicholas FitzGerald
Zixuan Jiang
Xuan Yang
Dale Webster
Ali Elqursh
Aedan Pope
Georges Rotival
David Raposo
Wanzheng Zhu
Jeff Dean
Sami Alabed
Dustin Tran
Arushi Gupta
Zach Gleicher
Jessica Austin
Edouard Rosseel
Megh Umekar
Dipanjan Das
Yinghao Sun
Kai Chen
Karolis Misiunas
Xiang Zhou
Yixian Di
Alyssa Loo
Josh Newlan
Bo Li
Vinay Ramasesh
Ying Xu
Alex Chen
Sudeep Gandhe
Radu Soricut
Nikita Gupta
Shuguang Hu
Seliem El-Sayed
Xavier Garcia
Idan Brusilovsky
Pu-Chin Chen
Andrew Bolt
Lu Huang
Alex Gurney
Zhiying Zhang
Alexander Pritzel
Jarek Wilkiewicz
Bryan Seybold
Bhargav Kanagal
Shamanna
Felix Fischer
Josef Dean
Karan Gill
Ross Mcilroy
Abhishek Bhowmick
Jeremy Selier
Antoine Yang
Derek Cheng
Vladimir Magay
Jie Tan
Dhriti Varma
Christian Walder
Tomas Kocisky
Ryo Nakashima
Paul Natsev
Mike Kwong
Ionel Gog
Chiyuan Zhang
Sander Dieleman
Thomas Jimma
Andrey Ryabtsev
Siddhartha Brahma
David Steiner
Dayou Du
Ante Žužul
Mislav Žanić
Mukund Raghavachari
Willi Gierke
Zeyu Zheng
Dessie Petrova
Yann Dauphin
Yuchuan Liu
Ido Kessler
Steven Hand
Chris Duvarney
Seokhwan Kim
Hyo Lee
Léonard Hussenot
Jeffrey Hui
Josh Smith
Deepali Jain
Jiawei Xia
Gaurav Singh Tomar
Keyvan Amiri
Du Phan
Fabian Fuchs
Tobias Weyand
Nenad Tomasev
Alexandra Cordell
Xin Liu
Jonathan Mallinson
Pankaj Joshi
Andy Crawford
Arun Suggala
Steve Chien
Nick Fernando
Mariella Sanchez-Vargas
Duncan Williams
Phil Crone
Xiyang Luo
Igor Karpov
Jyn Shan
Terry Thurk
Robin Strudel
Paul Voigtlaender
Piyush Patil
Tim Dozat
Ali Khodaei
Sahil Singla
Piotr Ambroszczyk
Qiyin Wu
Yifan Chang
Brian Roark
Chaitra Hegde
Tianli Ding
Angelos Filos
Zhongru Wu
André Susano Pinto
Shuang Liu
Saarthak Khanna
Aditya Pandey
Siobhan Mcloughlin
Qiujia Li
Sam Haves
Allan Zhou
Elena Buchatskaya
Isabel Leal
Peter de Boursac
Nami Akazawa
Nina Anderson
Terry Chen
Krishna Somandepalli
Chen Liang
Sheela Goenka
Stephanie Winkler
Alexander Grushetsky
Yifan Ding
Jamie Smith
Fan Ye
Jordi Pont-Tuset
Eric Li
Ruichao Li
Tomer Golany
Dawid Wegner
Tao Jiang
Omer Barak
Yuan Shangguan
Eszter Vértes
Renee Wong
Jörg Bornschein
Alex Tudor
Michele Bevilacqua
Tom Schaul
Ankit Singh Rawat
Yang Zhao
Kyriakos Axiotis
Lei Meng
Cory McLean
Jonathan Lai
Jennifer Beattie
Nate Kushman
Yaxin Liu
Blair Kutzman
Fiona Lang
Jingchen Ye
Praneeth Netrapalli
Pushkar Mishra
Myriam Khan
Megha Goel
Rob Willoughby
David Tian
Honglei Zhuang
JD Chen
Zak Tsai
Tasos Kementsietsidis
Arjun Khare
James Keeling
Keyang Xu
Nathan Waters

Florent Altché
Ashok Popat
Bhavishya Mittal
David Saxton
Dalia El Badawy
Michael Mathieu
Zheng Zheng
Hao Zhou
Nishant Ranka
Richard Shin
Qingnan Duan
Tim Salimans
Ioana Mihailescu
Uri Shaham
Ming-Wei Chang
Yannis Assael
Nishanth Dikkala
Martin Izzard
Vincent Cohen-Addad
Cat Graves
Vlad Feinberg
Grace Chung
DJ Strouse
Danny Karmon
Sahand Sharifzadeh
Zoe Ashwood
Khiem Pham
Jon Blanton
Alex Vasiloff
Jarred Barber
Mark Geller
Aurick Zhou
Fedir Zubach
Tzu-Kuo Huang
Lei Zhang
Himanshu Gupta
Matt Young
Julia Proskurnia
Ronny Votel
Valentin Gabeur
Gabriel Barcik
Aditya Tripathi
Hongkun Yu
Geng Yan
Beer Changpinyo
Filip Pavetić
Amy Coyle
Yasuhisa Fujii
Jorge Gonzalez Mendez
Tianhao Zhou
Harish Rajamani
Blake Hechtman
Eddie Cao
Da-Cheng Juan
Yi-Xuan Tan
Valentin Dalibard
Yilun Du
Natalie Clay
Kaisheng Yao
Wenhao Jia
Dimple Vijaykumar
Yuxiang Zhou
Xinyi Bai
Wei-Chih Hung
Steven Pecht
Georgi Todorov
Nikhil Khadke
Pramod Gupta
Preethi Lahoti
Arnaud Autef
Karthik Duddu
James Lee-Thorp
Alexander Bykovsky
Tautvydas Misiunas
Sebastian Flennerhag
Santhosh Thangaraj
Jed McGiffin
Zack Nado
Markus Kunesch
Andreas Noever
Amir Hertz
Marco Liang
Victor Stone
Evan Palmer
Samira Daruki
Arijit Pramanik
Siim Põder
Austin Kyker
Mina Khan
Evgeny Sluzhaev
Marvin Ritter
Avraham Ruderman
Wenlei Zhou
Chirag Nagpal
Kiran Vodrahalli
George Necula
Paul Barham
Ellie Pavlick
Jay Hartford
Izhak Shafran
Long Zhao
Maciej Mikuła
Tom Eccles
Hidetoshi Shimokawa
Kanav Garg
Luke Vilnis
Hanwen Chen
Ilia Shumailov
Kuang-Huei Lee
Abdelrahman Abdelhamed
Meiyan Xie
Vered Cohen
Ester Hlavnova
Dan Malkin
Chawin Sitawarin
James Lottes
Pauline Coquinot
Tianli Yu
Sandeep Kumar
Jingwei Zhang
Aroma Mahendru
Zafarali Ahmed
James Martens
Tao Chen
Aviel Boag
Daiyi Peng
Coline Devin
Arseniy Klimovskiy
Mary Phuong
Danny Vainstein
Jin Xie
Bhuvana Ramabhadran
Nathan Howard
Xinxin Yu
Gitartha Goswami
Jingyu Cui
Sam Shleifer
Mario Pinto
Chih-Kuan Yeh
Ming-Hsuan Yang
Sara Javanmardi
Dan Ethier
Chace Lee
Jordi Orbay
Suyog Kotecha
Carla Bromberg
Pete Shaw
James Thornton
Adi Gerzi Rosenthal
Shane Gu
Matt Thomas
Ian Gemp
Aditya Ayyar
Asahi Ushio
Aarush Selvan
Joel Wee
Chenxi Liu
Maryam Majzoubi
Weiren Yu
Jake Abernethy
Tyler Liechty
Renke Pan
Hoang Nguyen
Qiong (Q) Hu
Sarah Perrin
Abhinav Arora
Emily Pitler
Weiyi Wang
Kaushik Shivakumar
Flavien Prost
Ben Limonchik
Jing Wang
Yi Gao
Timothee Cour
Shyamal Buch
Huan Gui
Maria Ivanova
Philipp Neubeck
Kelvin Chan
Lucy Kim
Huizhong Chen
Naman Goyal
Da-Woon Chung
Lu Liu
Yao Su
Anastasia Petrushkina
Jiajun Shen
Armand Joulin
Yuanzhong Xu
Stein Xudong Lin
Yana Kulizhskaya
Ciprian Chelba
Shobha Vasudevan
Eli Collins
Vasilisa Bashlovkina
Tony Lu
Doug Fritz
Jongbin Park
Yanqi Zhou
Chen Su
Richard Tanburn
Mikhail Sushkov
Mitchelle Rasquinha
Jinning Li
Jennifer Prendki
Yiming Li
Pallavi LV
Shriya Sharma
Hen Fitoussi
Hui Huang
Andrew Dai
Phuong Dao
Mike Burrows
Henry Prior
Danfeng Qin
Golan Pundak
Lars Lowe Sjoesund
Art Khurshudov
Zhenkai Zhu
Albert Webson
Elizabeth Kemp
Tat Tan
Saurabh Agrawal
Susie Sargsyan
Liqun Cheng
Jim Stephan
Tom Kwiatkowski
David Reid
Arunkumar Byravan
Assaf Hurwitz Michaely
Nicolas Heess
Luowei Zhou
Sonam Goenka
Viral Carpenter
Anselm Levskaya
Bo Wang
Reed Roberts
Rémi Leblond
Sharat Chikkerur
Stav Ginzburg
Max Chang
Robert Riachi
Chuqiao (Joyce) Xu
Zalán Borsos

Michael Pliskin
Julia Pawar
Morgane Lustman
Hannah Kirkwood
Ankit Anand
Aditi Chaudhary
Norbert Kalb
Kieran Milan
Sean Augenstein
Anna Goldie
Laurel Prince
Karthik Raman
Yanhua Sun
Vivian Xia
Aaron Cohen
Zhouyuan Huo
Josh Camp
Seher Ellis
Lukas Zilka
David Vilar Torres
Lisa Patel
Sho Arora
Betty Chan
Jonas Adler
Kareem Ayoub
Jacky Liang
Fayaz Jamil
Jiepu Jiang
Simon Baumgartner
Haitian Sun
Yael Karov
Yaroslav Akulov
Hui Zheng
Irene Cai
Claudio Fantacci
James Rubin
Alex Rav Acha
Mengchao Wang
Nina D'Souza
Rohit Sathyanarayana
Shengyang Dai
Simon Rowe
Andrey Simanovsky
Omer Goldman
Yuheng Kuang
Xiaoyue Pan
Andrew Rosenberg
Tania Rojas-Esponda
Praneet Dutta
Amy Zeng
Irina Jurenka
Greg Farquhar
Yamini Bansal
Shariq Iqbal
Becca Roelofs
Ga-Young Joung
Parker Beak
Changwan Ryu
Ryan Poplin
Yan Wu
Jean-Baptiste Alayrac
Senaka Buthpitiya
Olaf Ronneberger
Caleb Habtegebriel
Wei Li
Paul Cavallaro
Aurora Wei
Guy Bensky
Timo Denk
Harish Ganapathy
Jeff Stanway
Pratik Joshi
Francesco Bertolini
Jessica Lo
Olivia Ma
Zachary Charles
Geta Sampemane
Himanshu Sahni
Xu Chen
Harry Askham
David Gaddy
Peter Young
Jiewen Tan
Matan Eyal
Arthur Bražinskas
Li Zhong
Zhichun Wu
Mark Epstein
Kai Bailey
Andrew Hard
Kamyu Lee
Sasha Goldshtein
Alex Ruiz
Mohammed Badawi
Matthias Lochbrunner
JK Kearns
Ashley Brown
Fabio Pardo
Theophane Weber
Haichuan Yang
Pan-Pan Jiang
Berkin Akin
Zhao Fu
Marcus Wainwright
Chi Zou
Meenu Gaba
Pierre-Antoine Manzagol
Wendy Kan
Yang Song
Karina Zainullina
Rui Lin
Jeongwoo Ko
Salil Deshmukh
Apoorv Jindal
James Svensson
Divya Tyam
Heri Zhao
Christine Kaeser-Chen
Scott Baird
Pooya Moradi
Jamie Hall
Qiuchen Guo
Vincent Tsang
Bowen Liang
Fernando Pereira
Suhas Ganesh
Ivan Korotkov
Jakub Adamek
Sridhar Thiagarajan
Vinh Tran
Charles Chen
Chris Tar
Sanil Jain
Ishita Dasgupta
Taylan Bilal
David Reitter
Kai Zhao
Giulia Vezzani
Yasmin Gehman
Pulkit Mehta
Lauren Beltrone
Xerxes Dotiwalla
Sergio Guadarrama
Zaheer Abbas
Stefani Karp
Petko Georgiev
Chun-Sung Ferng
Marc Brockschmidt
Liqian Peng
Christoph Hirnschall
Vikas Verma
Yingying Bi
Ying Xiao
Avigail Dabush
Kelvin Xu
Phil Wallis
Randall Parker
Qifei Wang
Yang Xu
Ilkin Safarli
Dinesh Tewari
Yin Zhang
Seungyeon Kim
Andrea Gesmundo
Mackenzie Thomas
Sergey Levi
Ahmed Chowdhury
Kanishka Rao
Peter Garst
Sam Conway-Rahman
Helen Ran
Kay McKinney
Zhisheng Xiao
Wenhao Yu
Rohan Agrawal
Axel Stjerngren
Catalin Ionescu
Jingjing Chen
Vivek Sharma
Justin Chiu
Fei Liu
Ken Franko
Clayton Sanford
Xingyu Cai
Paul Michel
Sanjay Ganapathy
Jane Labanowski
Zachary Garrett
Ben Vargas
Sean Sun
Bryan Gale
Thomas Buschmann
Guillaume Desjardins
Nimesh Ghelani
Palak Jain
Mudit Verma
Chulayuth Asawaroengchai
Julian Eisenschlos
Jitendra Harlalka
Hideto Kazawa
Don Metzler
Joshua Howland
Ying Jian
Jake Ades
Viral Shah
Tynan Gangwani
Seungji Lee
Roman Ring
Steven M. Hernandez
Dean Reich
Amer Sinha
Ashutosh Sathe
Joe Kovac
Ashleah Gill
Ajay Kannan
Andrea D'olimpio
Martin Sevenich
Jay Whang
Been Kim
Khe Chai Sim
Jilin Chen
Jiageng Zhang
Shuba Lall
Yossi Matias
Bill Jia
Abe Friesen
Sara Nasso
Ashish Thapliyal
Bryan Perozzi
Ting Yu
Anna Shekhawat
Safeen Huda
Peter Grabowski
Eric Wang
Ashwin Sreevatsa
Hilal Dib
Mehadi Hassen
Parker Schuh
Vedrana Milutinovic
Chris Welty
Michael Quinn
Ali Shah
Bangju Wang
Gabe Barth-Maron

Justin Frye
Natalie Axelsson
Tao Zhu
Yukun Ma
Irene Giannoumis
Hanie Sedghi
Chang Ye
Yi Luan
Kevin Aydin
Bilva Chandra
Vivek Sampathkumar
Ronny Huang
Victor Lavrenko
Ahmed Eleryan
Zhi Hong
Steven Hansen
Sara Mc Carthy
Bidisha Samanta
Domagoj Ćevid
Xin Wang
Fangtao Li
Michael Voznesensky
Matt Hoffman
Andreas Terzis
Vikash Sehwag
Gil Fidel
Luheng He
Mu Cai
Yanzhang He
Alex Feng
Martin Nikoltchev
Samrat Phatale
Jason Chase
Rory Lawton
Ming Zhang
Tom Ouyang
Manuel Tragut
Mehdi Hafezi Manshadi
Arjun Narayanan
Jiaming Shen
Xu Gao
Tolga Bolukbasi
Nick Roy
Xin Li
Daniel Golovin
Liviu Panait
Zhen Qin
Guangxing Han
Thomas Anthony
Sneha Kudugunta
Viorica Patraucean
Aniket Ray
Xinyun Chen
Xiaochen Yang
Tanuj Bhatia
Pranav Talluri
Alex Morris
Andrija Ražnatović
Bethanie Brownfield
James An
Sheng Peng
Patrick Kane
Ce Zheng
Nico Duduta
Joshua Kessinger
James Noraky
Siqi Liu
Keran Rong
Petar Veličković
Keith Rush
Alex Goldin
Fanny Wei
Shiva Mohan Reddy
Garlapati
Caroline Pantofaru
Okwan Kwon
Jianmo Ni
Eric Noland
Julia Di Trapani
Françoise Beaufays
Abhijit Guha Roy
Yinlam Chow
Aybuke Turker
Geoffrey Cideron
Lantao Mei
Jon Clark
Qingyun Dou
Matko Bošnjak
Ralph Leith
Yuqing Du
Amir Yazdanbakhsh
Milad Nasr
Chester Kwak
Suraj Satishkumar Sheth
Alex Kaskasoli
Ankesh Anand
Balaji Lakshminarayanan
Sammy Jerome
David Bieber
Chun-Te Chu
Alexandre Senges
Tianxiao Shen
Mukund Sridhar
Ndaba Ndebele
Benjamin Beyret
Shakir Mohamed
Mia Chen
Markus Freitag
Jiaxian Guo
Luyang Liu
Paul Roit
Heng Chen
Shen Yan
Tom Stone
JD Co-Reyes
Jeremy Cole
Salvatore Scellato
Shekoofeh Azizi
Hadi Hashemi
Alicia Jin
Anand Iyer
Marcella Valentine
András György
Arun Ahuja
Daniel Hernandez Diaz
Chen-Yu Lee
Nathan Clement
Weize Kong
Drew Garmon
Ishaan Watts
Kush Bhatia
Khyatti Gupta
Matt Miecnikowski
Hugo Vallet
Ankur Taly
Edward Loper
Saket Joshi
James Atwood
Jo Chick
Mark Collier
Fotis Iliopoulos
Ryan Trostle
Beliz Gunel
Ramiro Leal-Cavazos
Arnar Mar Hrafnkelsson
Michael Guzman
Xiaoen Ju
Andy Forbes
Jesse Emond
Kushal Chauhan
Ben Caine
Li Xiao
Wenjun Zeng
Alexandre Moufarek
Daniel Murphy
Maya Meng
Nitish Gupta
Felix Riedel
Anil Das
Elijah Lawal
Shashi Narayan
Tiberiu Sosea
James Swirhun
Linda Friso
Behnam Neyshabur
Jing Lu
Sertan Girgin
Michael Wunder
Edouard Yvinec
Aroonalok Pyne
Victor Carbune
Shruti Rijhwani
Yang Guo
Tulsee Doshi
Anton Briukhov
Max Bain
Ayal Hitron
Xuanhui Wang
Ashish Gupta
Ke Chen
Cosmo Du
Weiyang Zhang
Dhruv Shah
Arjun Akula
Max Dylla
Ashyana Kachra
Weicheng Kuo
Tingting Zou
Lily Wang
Luyao Xu
Jifan Zhu
Justin Snyder
Sachit Menon
Orhan Firat
Igor Mordatch
Yuan Yuan
Natalia Ponomareva
Rory Blevins
Lawrence Moore
Weijun Wang
Phil Chen
Martin Scholz
Artur Dwornik
Jason Lin
Sicheng Li
Diego Antognini
Te I
Xiaodan Song
Matt Miller
Uday Kalra
Adam Raveret
Oscar Akerlund
Felix Wu
Andrew Nystrom
Namrata Godbole
Tianqi Liu
Hannah DeBalsi
Jewel Zhao
Buhuang Liu
Avi Caciularu
Lauren Lax
Urvashi Khandelwal
Victoria Langston
Eric Bailey
Silvio Lattanzi
Yufei Wang
Neel Kovelamudi
Sneha Mondal
Guru Guruganesh
Nan Hua
Ofir Roval
Paweł Wesołowski
Rishikesh Ingale
Jonathan Halcrow
Tim Sohn
Christof Angermueller
Bahram Raad
Eli Stickgold
Eva Lu
Alec Kosik
Jing Xie
Timothy Lillicrap
Austin Huang
Lydia Lihui Zhang

Dominik Paulus
Clement Farabet
Alex Wertheim
Bing Wang
Rishabh Joshi
Chu-ling Ko
Yonghui Wu
Shubham Agrawal
Lily Lin
XiangHai Sheng
Peter Sung
Tyler Breland-King
Christina Butterfield
Swapnil Gawde
Sumeet Singh
Qiao Zhang
Raj Apte
Shilpa Shetty
Adrian Hutter
Tao Li
Elizabeth Salesky
Federico Lebron
Jonni Kanerva
Michela Paganini
Arthur Nguyen
Rohith Vallu
Jan-Thorsten Peter
Sarmishta Velury
David Kao
Jay Hoover
Anna Bortsova
Colton Bishop
Shoshana Jakobovits
Alessandro Agostini
Alekh Agarwal
Chang Liu
Charles Kwong
Sasan Tavakkol
Ioana Bica
Alex Greve
Anirudh GP
Jake Marcus
Le Hou
Tom Duerig
Rivka Moroshko
Dave Lacey
Andy Davis
Julien Amelot
Guohui Wang
Frank Kim
Theofilos Strinopoulos
Hui Wan
Charline Le Lan
Shankar Krishnan
Haotian Tang
Peter Humphreys
Junwen Bai
Idan Heimlich Shtacher
Diego Machado
Chenxi Pang
Ken Burke
Dangyi Liu
Renga Aravamudhan
Yue Song
Ed Hirst
Abhimanyu Singh
Brendan Jou
Liang Bai
Francesco Piccinno
Chuyuan Kelly Fu
Robin Alazard
Barak Meiri
Daniel Winter
Charlie Chen
Mingda Zhang
Jens Heitkaemper
John Lambert
Jinhyuk Lee
Alexander Frömmgen
Sergey Rogulenko
Pranav Nair
Paul Niemczyk
Anton Bulyenov
Bibo Xu
Hadar Shemtov
Morteza Zadimoghaddam
Serge Toropov
Mateo Wirth
Hanjun Dai
Sreenivas Gollapudi
Daniel Zheng
Alex Kurakin
Chansoo Lee
Kalesha Bullard
Nicolas Serrano
Ivana Balazevic
Yang Li
Johan Schalkwyk
Mark Murphy
Mingyang Zhang
Kevin Sequeira
Romina Datta
Nishant Agrawal
Charles Sutton
Nithya Attaluri
Mencher Chiang
Wael Farhan
Gregory Thornton
Kate Lin
Travis Choma
Hung Nguyen
Kingshuk Dasgupta
Dirk Robinson
Iulia Comşa
Michael Riley
Arjun Pillai
Basil Mustafa
Ben Golan
Amir Zandieh
Jean-Baptiste Lespiau
Billy Porter
David Ross
Sujeevan Rajayogam
Mohit Agarwal
Subhashini Venugopalan
Bobak Shahriari
Qiqi Yan
Hao Xu
Taylor Tobin
Pavel Dubov
Hongzhi Shi
Adrià Recasens
Anton Kovsharov
Sebastian Borgeaud
Lucio Dery
Shanthal Vasanth
Elena Gribovskaya
Linhai Qiu
Mahdis Mahdieh
Wojtek Skut
Elizabeth Nielsen
CJ Zheng
Adams Yu
Carrie Grimes Bostock
Shaleen Gupta
Aaron Archer
Chris Rawles
Elinor Davies
Alexey Svyatkovskiy
Tomy Tsai
Yoni Halpern
Christian Reisswig
Bartek Wydrowski
Bo Chang
Joan Puigcerver
Mor Hazan Taege
Jian Li
Eva Schnider
Xinjian Li
Dragos Dena
Yunhan Xu
Umesh Telang
Tianze Shi
Heiga Zen
Kyle Kastner
Yeongil Ko
Neesha Subramaniam
Aviral Kumar
Pete Blois
Zhuyun Dai
John Wieting
Yifeng Lu
Yoel Zeldes
Tian Xie
Anja Hauth
Alexandru Ţifrea
Yuqi Li
Sam El-Husseini
Dan Abolafia
Howard Zhou
Wen Ding
Sahra Ghalebikesabi
Carlos Guía
Andrii Maksai
Ágoston Weisz
Sercan Arik
Nick Sukhanov
Aga Świetlik
Xuhui Jia
Luo Yu
Weiyue Wang
Mark Brand
Dawn Bloxwich
Sean Kirmani
Zhe Chen
Alec Go
Pablo Sprechmann
Nithish Kannen
Alen Carin
Paramjit Sandhu
Isabel Edkins
Leslie Nooteboom
Jai Gupta
Loren Maggiore
Javad Azizi
Yael Pritch
Pengcheng Yin
Mansi Gupta
Danny Tarlow
Duncan Smith
Desi Ivanov
Mohammad Babaeizadeh
Ankita Goel
Satish Kambala
Grace Chu
Matej Kastelic
Michelle Liu
Hagen Soltau
Austin Stone
Shivani Agrawal
Min Kim
Kedar Soparkar
Srinivas Tadepalli
Oskar Bunyan
Rachel Soh
Arvind Kannan
DY Kim
Blake JianHang Chen
Afief Halumi
Sudeshna Roy
Yulong Wang
Olcan Sercinoglu
Gena Gibson
Sijal Bhatnagar
Motoki Sano
Daniel von Dincklage
Qingchun Ren
Blagoj Mitrevski
Mirek Olšák
Jennifer She
Carl Doersch
Jilei (Jerry) Wang
Bingyuan Liu
Qijun Tan

Tamar Yakar
Tris Warkentin
Alex Ramirez
Carl Lebsack
Josh Dillon
Rajiv Mathews
Tom Cobley
Zelin Wu
Zhuoyuan Chen
Jon Simon
Swaroop Nath
Tara Sainath
Alexei Bendebury
Ryan Julian
Bharath Mankalale
Daria Ćurko
Paulo Zacchello
Adam R. Brown
Kiranbir Sodhia
Heidi Howard
Sergi Caelles
Abhinav Gupta
Gareth Evans
Anna Bulanova
Lesley Katzen
Roman Goldenberg
Anton Tsitsulin
Joe Stanton
Benoit Schillings
Vitaly Kovalev
Corey Fry
Rushin Shah
Kuo Lin
Shyam Upadhyay
Cheng Li
Soroush Radpour
Marcello Maggioni
Jing Xiong
Lukas Haas
Jenny Brennan
Aishwarya Kamath
Nikolay Savinov
Arsha Nagrani
Trevor Yacovone
Ryan Kappedal
Kostas Andriopoulos
Li Lao
YaGuang Li
Grigory Rozhdestvenskiy
Kazuma Hashimoto
Andrew Audibert
Sophia Austin
Daniel Rodriguez
Anian Ruoss
Garrett Honke
Deep Karkhanis
Xi Xiong
Qing Wei
James Huang
Zhaoqi Leng
Vittal Premachandran
Stan Bileschi
Georgios Evangelopoulos
Thomas Mensink
Jay Pavagadhi
Denis Teplyashin
Paul Chang
Linting Xue
Garrett Tanzer
Sally Goldman
Kaushal Patel
Shixin Li
Jeremy Wiesner
Ivy Zheng
Ian Stewart-Binks
Jie Han
Zhi Li
Liangchen Luo
Karel Lenc
Mario Lučić
Fuzhao Xue
Ryan Mullins
Alexey Guseynov
Chung-Ching Chang
Isaac Galatzer-Levy
Adam Zhang
Garrett Bingham
Grace Hu
Ale Hartman
Yue Ma
Jordan Griffith
Alex Irpan
Carey Radebaugh
Summer Yue
Lijie Fan
Victor Ungureanu
Christina Sorokin
Hannah Teufel
Peiran Li
Rohan Anil
Dimitris Paparas
Todd Wang
Chu-Cheng Lin
Hui Peng
Megan Shum
Goran Petrovic
Demetra Brady
Richard Nguyen
Klaus Macherey
Zhihao Li
Harman Singh
Madhavi Yenugula
Mariko Iinuma
Xinyi Chen
Kavya Kopparapu
Alexey Stern
Shachi Dave
Chandu Thekkath
Florence Perot
Anurag Kumar
Fangda Li
Yang Xiao
Matthew Bilotti
Mohammad Hossein Bateni
Isaac Noble
Lisa Lee
Amelio Vázquez-Reina
Julian Salazar
Xiaomeng Yang
Boyu Wang
Ela Gruzewska
Anand Rao
Sindhu Raghuram
Zheng Xu
Eyal Ben-David
Jieru Mei
Sid Dalmia
Zhaoyi Zhang
Yuchen Liu
Gagan Bansal
Helena Pankov
Steven Schwarcz
Andrea Burns
Christine Chan
Sumit Sanghai
Ricky Liang
Ethan Liang
Antoine He
Amy Stuart
Arun Narayanan
Yukun Zhu
Christian Frank
Bahar Fatemi
Amit Sabne
Oran Lang
Indro Bhattacharya
Shane Settle
Maria Wang
Brendan McMahan
Andrea Tacchetti
Livio Baldini Soares
Majid Hadian
Serkan Cabi
Timothy Chung
Nikita Putikhin
Gang Li
Jeremy Chen
Austin Tarango
Henryk Michalewski
Mehran Kazemi
Hussain Masoom
Hila Sheftel
Rakesh Shivanna
Archita Vadali
Ramona Comanescu
Doug Reid
Joss Moore
Arvind Neelakantan
Michaël Sander
Jonathan Herzig
Aviv Rosenberg
Mostafa Dehghani
JD Choi
Michael Fink
Reid Hayes
Eric Ge
Shitao Weng
Chia-Hua Ho
John Karro
Kalpesh Krishna
Lam Nguyen Thiet
Amy Skerry-Ryan
Daniel Eppens
Marco Andreetto
Navin Sarma
Silvano Bonacina
Burcu Karagol Ayan
Megha Nawhal
Zhihao Shan
Mike Dusenberry
Shantanu Thakoor
Sagar Gubbi
Duc Dung Nguyen
Reut Tsarfaty
Samuel Albanie
Jovana Mitrović
Meet Gandhi
Bo-Juen Chen
Alessandro Epasto
Georgi Stephanov
Ye Jin
Samuel Gehman
Aida Amini
Jack Weber
Feryal Behbahani
Shawn Xu
Miltos Allamanis
Xi Chen
Myle Ott
Claire Sha
Michal Jastrzebski
Hang Qi
David Greene
Xinyi Wu
Abodunrinwa Toki
Daniel Vlasic
Jane Shapiro
Ragha Kotikalapudi
Zhe Shen
Takaaki Saeki
Sirui Xie
Albin Cassirer
Shikhar Bharadwaj
Tatsuya Kiyono
Srinadh Bhojanapalli
Elan Rosenfeld
Sam Ritter
Jieming Mao
João Gabriel Oliveira
Zoltan Egyed
Bernd Bandemer
Emilio Parisotto
Keisuke Kinoshita
Juliette Pluto

Petros Maniatis
Steve Li
Yaohui Guo
Golnaz Ghiasi
Jean Tarbouriech
Srimon Chatterjee
Julie Jin
Katrina (Xinyi) Xu
Jennimaria Palomaki
Séb Arnold
Madhavi Sewak
Federico Piccinini
Mohit Sharma
Ben Albrecht
Sean Purser-haskell
Ashwin Vaswani
Chongyan Chen
Matheus Wisniewski
Qin Cao
John Aslanides
Nguyet Minh Phu
Maximilian Sieb
Lauren Agubuzu
Anne Zheng
Daniel Sohn
Marco Selvi
Anders Andreassen
Krishan Subudhi
Prem Eruvbetine
Oliver Woodman
Tomas Mery
Sebastian Krause
Xiaoqi Ren
Xiao Ma
Jincheng Luo
Dawn Chen
Wei Fan
Henry Griffiths
Christian Schuler
Alice Li
Shujian Zhang
Jean-Michel Sarr
Shixin Luo
Riccardo Patana
Matthew Watson
Dani Naboulsi
Michael Collins
Sailesh Sidhwani
Emiel Hoogeboom
Sharon Silver
Emily Caveness
Xiaokai Zhao
Mikel Rodriguez
Maxine Deines
Libin Bai
Patrick Griffin
Marco Tagliasacchi
Emily Xue
Spandana Raj Babbula
Bo Pang
Nan Ding
Gloria Shen
Elijah Peake
Remi Crocker
Shubha Srinivas
Raghvendra
Danny Swisher
Woohyun Han
Richa Singh
Ling Wu
Vladimir Pchelin
Tsendsuren Munkhdalai
Dana Alon
Geoff Bacon
Efren Robles
Jannis Bulian
Melvin Johnson
George Powell
Felipe Tiengo Ferreira
Yaoyiran Li
Frederik Benzing
Mihajlo Velimirović
Hubert Soyer
William Kong
Tony (Tuấn) Nguyễn
Zhen Yang
Jeremiah Liu
Joost van Amersfoort
Daniel Gillick
Baochen Sun
Nathalie Rauschmayr
Katie Zhang
Serena Zhan
Tao Zhou
Alexey Frolov
Chengrun Yang
Denis Vnukov
Louis Rouillard
Hongji Li
Amol Mandhane
Nova Fallen
Rajesh Venkataraman
Clara Huiyi Hu
Jennifer Brennan
Jenny Lee
Jerry Chang
Martin Sundermeyer
Zhufeng Pan
Rosemary Ke
Simon Tong
Alex Fabrikant
William Bono
Jindong Gu
Ryan Foley
Yiran Mao
Manolis Delakis
Dhruva Bhaswar
Roy Frostig
Nick Li
Avital Zipori
Cath Hope
Olga Kozlova
Swaroop Mishra
Josip Djolonga
Craig Schiff
Majd Al Merey
Eleftheria Briakou
Peter Morgan
Andy Wan
Avinatan Hassidim
RJ Skerry-Ryan
Kuntal Sengupta
Mary Jasarevic
Praveen Kallakuri
Paige Kunkle
Hannah Brennan
Tom Lieber
Hassan Mansoor
Julian Walker
Bing Zhang
Annie Xie
Goran Žužić
Adaeze Chukwuka
Alex Druinsky
Donghyun Cho
Rui Yao
Ferjad Naeem
Shiraz Butt
Eunyoung Kim
Zhipeng Jia
Mandy Jordan
Adam Lelkes
Mark Kurzeja
Sophie Wang
James Zhao
Andrew Over
Abhishek Chakladar
Marcel Prasetya
Neha Jha
Sriram Ganapathy
Yale Cong
Prakash Shroff
Carl Saroufim
Sobhan Miryoosefi
Mohamed Hammad
Tajwar Nasir
Weijuan Xi
Yang Gao
Young Maeng
Ben Hora
Chin-Yi Cheng
Parisa Haghani
Yoad Lewenberg
Caden Lu
Martin Matysiak
Naina Raisinghani
Huiyu Wang
Lexi Baugher
Rahul Sukthankar
Minh Giang
John Schultz
Noah Fiedel
Minmin Chen
Cheng-Chun Lee
Tapomay Dey
Hao Zheng
Shachi Paul
Celine Smith
Andy Ly
Yicheng Wang
Rishabh Bansal
Bartek Perz
Susanna Ricco
Stasha Blank
Vaishakh Keshava
Deepak Sharma
Marvin Chow
Kunal Lad
Komal Jalan
Simon Osindero
Craig Swanson
Jacob Scott
Anastasija Ilić
Xiaowei Li
Siddhartha Reddy
Jonnalagadda
Afzal Shama Soudagar
Yan Xiong
Bat-Orgil Batsaikhan
Daniel Jarrett
Naveen Kumar
Maulik Shah
Matt Lawlor
Austin Waters
Mark Graham
Rhys May
Sabela Ramos
Sandra Lefdal
Zeynep Cankara
Nacho Cano
Brendan O'Donoghue
Jed Borovik
Frederick Liu
Jordan Grimstad
Mahmoud Alnahlawi
Katerina Tsihlas
Tom Hudson
Nikolai Grigorev
Yiling Jia
Terry Huang
Tobenna Peter Igwe
Sergei Lebedev
Xiaodan Tang
Igor Krivokon
Frankie Garcia
Melissa Tan
Eric Jia
Peter Stys
Shikhar Vashishth
Yu Liang
Balaji Venkatraman
Chenjie Gu
Anastasios Kementsietsidis
Chen Zhu

Junehyuk Jung
Yunfei Bai
Mohammad Javad Hosseini
Faruk Ahmed
Aditya Gupta
Xin Yuan
Shereen Ashraf
Shitij Nigam
Gautam Vasudevan
Pranjal Awasthi
Adi Mayrav Gilady
Zelda Mariet
Ramy Eskander
Haiguang Li
Hexiang Hu
Guillermo Garrido
Philippe Schlattner
George Zhang
Rohun Saxena
Petar Dević
Kritika Muralidharan
Ashwin Murthy
Yiqian Zhou
Min Choi
Arissa Wongpanich
Zhengdong Wang
Premal Shah
Yuntao Xu
Yiling Huang
Stephen Spencer
Alice Chen
James Cohan
Junjie Wang
Jonathan Tompson
Junru Wu
Ruba Haroun
Haiqiong Li
Blanca Huergo
Fan Yang
Tongxin Yin
James Wendt
Michael Bendersky
Rahma Chaabouni
Javier Snaider
Johan Ferret
Abhishek Jindal
Tara Thompson
Andrew Xue
Will Bishop
Shubham Milind Phal
Archit Sharma
Yunhsuan Sung
Prabakar Radhakrishnan
Mo Shomrat
Reeve Ingle
Roopali Vij
Justin Gilmer
Mihai Dorin Istin
Sam Sobell
Yang Lu
Emily Nottage
Dorsa Sadigh
Jeremiah Willcock
Tingnan Zhang
Steve Xu
Sasha Brown
Katherine Lee
Gary Wang
Yun Zhu
Yi Tay
Cheolmin Kim
Audrey Gutierrez
Abhanshu Sharma
Yongqin Xian
Sungyong Seo
Claire Cui
Elena Pochernina
Cip Baetu
Krzysztof Jastrzębski
Mimi Ly
Mohamed Elhawaty
Dan Suh
Eren Sezener
Pidong Wang
Nancy Yuen
George Tucker
Jiahao Cai
Zuguang Yang
Cindy Wang
Alex Muzio
Hai Qian
Jae Yoo
Derek Lockhart
Kevin R. McKee
Mandy Guo
Malika Mehrotra
Artur Mendonça
Sanket Vaibhav Mehta
Sherry Ben
Chetan Tekur
Jiaqi Mu
Muye Zhu
Victoria Krakovna
Hongrae Lee
AJ Maschinot
Sébastien Cevey
HyunJeong Choe
Aijun Bai
Hansa Srinivasan
Derek Gasaway
Nick Young
Patrick Siegler
Dan Holtmann-Rice
Vihari Piratla
Kate Baumli
Roey Yogev
Alex Hofer
Hado van Hasselt
Svetlana Grant
Yuri Chervonyi
David Silver
Andrew Hogue
Ayushi Agarwal
Kathie Wang
Preeti Singh
Four Flynn
Josh Lipschultz
Robert David
Lizzetth Bellot
Yao-Yuan Yang
Long Le
Filippo Graziano
Kate Olszewska
Kevin Hui
Akanksha Maurya
Nikos Parotsidis
Weijie Chen
Tayo Oguntebi
Joe Kelley
Anirudh Baddepudi
Johannes Mauerer
Gregory Shaw
Alex Siegman
Lin Yang
Shravya Shetty
Subhrajit Roy
Yunting Song
Wojciech Stokowiec
Ryan Burnell
Omkar Savant
Robert Busa-Fekete
Jin Miao
Samrat Ghosh
Liam MacDermed
Phillip Lippe
Mikhail Dektiarev
Zach Behrman
Fabian Mentzer
Kelvin Nguyen
Meng Wei
Siddharth Verma
Chris Knutsen
Sudeep Dasari
Zhipeng Yan
Petr Mitrichev
Xingyu Wang
Virat Shejwalkar
Jacob Austin
Srinivas Sunkara
Navneet Potti
Yan Virin
Christian Wright
Gaël Liu
Oriana Riva
Etienne Pot
Greg Kochanski
Quoc Le
Gargi Balasubramaniam
Arka Dhar
Yuguo Liao
Adam Bloniarz
Divyansh Shukla
Elizabeth Cole
Jong Lee
Sheng Zhang
Sushant Kafle
Siddharth Vashishtha
Parsa Mahmoudieh
Grace Chen
Raphael Hoffmann
Pranesh Srinivasan
Agustin Dal Lago
Yoav Ben Shalom
Zi Wang
Michael Elabd
Anuj Sharma
Junhyuk Oh
Suraj Kothawade
Maigo Le
Marianne Monteiro
Shentao Yang
Kaiz Alarakyia
Robert Geirhos
Diana Mincu
Håvard Garnes
Hayato Kobayashi
Soroosh Mariooryad
Kacper Krasowiak
Zhixin (Lucas) Lai
Shibl Mourad
Mingqiu Wang
Fan Bu
Ophir Aharoni
Guanjie Chen
Abhimanyu Goyal
Vadim Zubov
Ankur Bapna
Elahe Dabir
Nisarg Kothari
Kay Lamerigts
Nicola De Cao
Jeremy Shar
Christopher Yew
Nitish Kulkarni
Dre Mahaarachchi
Mandar Joshi
Zhenhai Zhu
Jared Lichtarge
Yichao Zhou
Hannah Muckenhirn
Vittorio Selo
Oriol Vinyals
Peter Chen
Anthony Brohan
Vaibhav Mehta
Sarah Cogan
Ruth Wang
Ty Geri
Wei-Jen Ko
Wei Chen
Fabio Viola
Keshav Shivam
Lisa Wang
Madeleine Clare Elish

Raluca Ada Popa
Sébastien Pereira
Jianqiao Liu
Raphael Koster
Donnie Kim
Gufeng Zhang
Sayna Ebrahimi
Partha Talukdar
Yanyan Zheng
Petra Poklukar
Ales Mikhalap
Dale Johnson
Anitha Vijayakumar
Mark Omernick
Matt Dibb
Ayush Dubey
Qiong Hu
Apurv Suman
Vaibhav Aggarwal
Ilya Kornakov
Fei Xia
Wing Lowe
Alexey Kolganov
Ted Xiao
Vitaly Nikolaev
Steven Hemingray
Bonnie Li
Joana Iljazi
Mikołaj Rybiński
Ballie Sandhu
Peggy Lu
Thang Luong
Rodolphe Jenatton
Vineetha Govindaraj
Hui (Elena) Li
Gabriel Dulac-Arnold
Wonpyo Park
Henry Wang
Abhinit Modi
Jean Pouget-Abadie
Kristina Greller
Rahul Gupta
Robert Berry
Prajit Ramachandran
Jinyu Xie
Liam McCafferty
Jianling Wang
Kilol Gupta
Hyeontaek Lim
Blaž Bratanič
Andy Brock
Ilia Akolzin
Jim Sproch
Dan Karliner
Duhyeon Kim
Adrian Goedeckemeyer
Noam Shazeer
Cordelia Schmid
Daniele Calandriello
Parul Bhatia
Krzysztof Choromanski
Ceslee Montgomery
Dheeru Dua
Ana Ramalho
Helen King
Yue Gao
Lynn Nguyen
David Lindner
Divya Pitta
Oleaser Johnson
Khalid Salama
Diego Ardila
Michael Han
Erin Farnese
Seth Odoom
Ziyue Wang
Xiangzhuo Ding
Norman Rink
Ray Smith
Harshal Tushar Lehri
Eden Cohen
Neera Vats
Tong He
Parthasarathy Gopavarapu
Adam Paszke
Miteyan Patel
Wouter Van Gansbeke
Lucia Loher
Luis Castro
Maria Voitovich
Tamara von Glehn
Nelson George
Simon Niklaus
Zach Eaton-Rosen
Nemanja Rakićević
Erik Jue
Sagi Perel
Carrie Zhang
Yuval Bahat
Angéline Pouget
Zhi Xing
Fantine Huot
Ashish Shenoy
Taylor Bos
Vincent Coriou
Bryan Richter
Natasha Noy
Yaqing Wang
Santiago Ontanon
Siyang Qin
Gleb Makarchuk
Demis Hassabis
Zhuowan Li
Mandar Sharma
Kumaran Venkatesan
Iurii Kemaev
Roxanne Daniel
Shiyu Huang
Saloni Shah
Octavio Ponce
Warren (Weilun) Chen
Manaal Faruqui
Jialin Wu
Slavica Andačić
Szabolcs Payrits
Daniel McDuff
Tom Hume
Yuan Cao
MH Tessler
Qingze Wang
Yinan Wang
Ivor Rendulic
Eirikur Agustsson
Matthew Johnson
Tanya Lando
Andrew Howard
Sri Gayatri Sundara Padmanabhan
Mayank Daswani
Andrea Banino
Michael Kilgore
Jonathan Heek
Ziwei Ji
Alvaro Caceres
Conglong Li
Nora Kassner
Alexey Vlaskin
Zeyu Liu
Alex Grills
Yanhan Hou
Roykrong Sukkerd
Gowoon Cheon
Nishita Shetty
Larisa Markeeva
Piotr Stanczyk
Tejas Iyer
Yuan Gong
Shawn Gao
Keerthana Gopalakrishnan
Tim Blyth
Malcolm Reynolds
Avishkar Bhoopchand
Misha Bilenko
Dero Gharibian
Vicky Zayats
Aleksandra Faust
Abhinav Singh
Min Ma
Hongyang Jiao
Sudheendra Vijayanarasimhan
Lora Aroyo
Vikas Yadav
Sarah Chakera
Ashwin Kakarla
Vilobh Meshram
Karol Gregor
Gabriela Botea
Evan Senter
Dawei Jia
Geza Kovacs
Neha Sharma
Sebastien Baur
Kai Kang
Yifan He
Lin Zhuo
Marija Kostelac
Itay Laish
Songyou Peng
Louis O'Bryan
Daniel Kasenberg
Girish Ramchandra Rao
Edouard Leurent
Biao Zhang
Sage Stevens
Ana Salazar
Ye Zhang
Ivan Lobov
Jake Walker
Allen Porter
Morgan Redshaw
Han Ke
Abhishek Rao
Alex Lee
Hoi Lam
Michael Moffitt
Jaeyoun Kim
Siyuan Qiao
Terry Koo
Robert Dadashi
Xinying Song
Mukund Sundararajan
Peng Xu
Chizu Kawamoto
Yan Zhong
Clara Barbu
Apoorv Reddy
Mauro Verzetti
Leon Li
George Papamakarios
Hanna Klimczak-Plucińska
Mary Cassin
Koray Kavukcuoglu
Rigel Swavely
Alain Vaucher
Jeffrey Zhao
Ross Hemsley
Michael Tschannen
Heming Ge
Gaurav Menghani
Yang Yu
Natalie Ha
Wei He
Xiao Wu
Maggie Song
Rachel Sterneck
Stefan Zinke
Dan A. Calian
Annie Marsden
Alejandro Cruzado Ruiz
Matteo Hessel
Almog Gueta
Benjamin Lee
Brian Farris

Manish Gupta
Yunjie Li
Mohammad Saleh
Vedant Misra
Kefan Xiao
Piermaria Mendolicchio
Gavin Buttimore
Varvara Krayvanova
Nigamaa Nayakanti
Matthew Wiethoff
Yash Pande
Azalia Mirhoseini
Ni Lao
Jasmine Liu
Yiqing Hua
Angie Chen
Yury Malkov
Dmitry Kalashnikov
Shubham Gupta
Kartik Audhkhasi
Yuexiang Zhai
Sudhindra Kopalle
Prateek Jain
Eran Ofek
Clemens Meyer
Khuslen Baatarsukh
Hana Strejček
Jun Qian
James Freedman
Ricardo Figueira
Michal Sokolik
Olivier Bachem
Raymond Lin
Dia Kharrat
Chris Hidey
Pingmei Xu
Dennis Duan
Yin Li
Muge Ersoy
Richard Everett
Kevin Cen
Rebeca
Santamaria-Fernandez
Amir Taubenfeld
Ian Mackinnon
Linda Deng
Polina Zablotskaia
Shashank Viswanadha
Shivanker Goel
Damion Yates
Yunxiao Deng
Peter Choy
Mingqing Chen
Abhishek Sinha
Alex Mossin
Yiming Wang
Arthur Szlam
Susan Hao
Paul Kishan Rubenstein
Metin Toksoz-Exley
Miranda Aperghis
Yin Zhong
Junwhan Ahn
Michael Isard
Olivier Lacombe
Florian Luisier
Chrysovalantis Anastasiou
Yogesh Kalley
Utsav Prabhu
Emma Dunleavy
Shaan Bijwadia
Justin Mao-Jones
Kelly Chen
Rama Pasumarthi
Emily Wood
Adil Dostmohamed
Nate Hurley
Jiri Simsa
Alicia Parrish
Mantas Pajarskas
Matt Harvey
Ondrej Skopek
Yony Kochinski
Javier Rey
Verena Rieser
Denny Zhou
Sun Jae Lee
Trilok Acharya
Guowang Li
Joe Jiang
Xiaofan Zhang
Bryant Gipson
Ethan Mahintorabi
Marco Gelmi
Nima Khajehnouri
Angel Yeh
Kayi Lee
Loic Matthey
Leslie Baker
Trang Pham
Han Fu
Alex Pak
Prakhar Gupta
Cristina Vasconcelos
Adam Sadovsky
Brian Walker
Sissie Hsiao
Patrik Zochbauer
Andreea Marzoca
Noam Velan
Junhao Zeng
Gilles Baechler
Danny Driess
Divya Jain
Yanping Huang
Lizzie Tao
John Maggs
Nir Levine
Jon Schneider
Erika Gemzer
Samuel Petit
Shan Han
Zach Fisher
Dustin Zelle
Courtney Biles
Eugene Ie
Asya Fadeeva
Casper Liu
Juliana Vicente Franco
Adrian Collister
Hao Zhang
Renshen Wang
Ruizhe Zhao
Leandro Kieliger
Kurt Shuster
Rui Zhu
Boqing Gong
Lawrence Chan
Ruoxi Sun
Sujoy Basu
Roland Zimmermann
Jamie Hayes
Abhishek Bapna
Jasper Snoek
Weel Yang
Puranjay Datta
Jad Al Abdallah
Kevin Kilgour
Lu Li
SQ Mah
Yennie Jun
Morgane Rivière
Abhijit Karmarkar
Tammo Spalink
Tao Huang
Lucas Gonzalez
Duc-Hieu Tran
Averi Nowak
John Palowitch
Martin Chadwick
Ellie Talius
Harsh Mehta
Thibault Sellam
Philipp Fränken
Massimo Nicosia
Kyle He
Aditya Kini
David Amos
Sugato Basu
Harrison Jobe
Eleni Shaw
Qiantong Xu
Colin Evans
Daisuke Ikeda
Chaochao Yan
Larry Jin
Lun Wang
Sachin Yadav
Ilia Labzovsky
Ramesh Sampath
Ada Ma
Candice Schumann
Aditya Siddhant
Rohin Shah
John Youssef
Rishabh Agarwal
Natalie Dabney
Alessio Tonioni
Moran Ambar
Jing Li
Isabelle Guyon
Benny Li
David Soergel
Boya Fang
Georgi Karadzhov
Cristian Udrescu
Trieu Trinh
Vikas Raunak
Seb Noury
Dee Guo
Sonal Gupta
Mara Finkelstein
Denis Petek
Lihao Liang
Greg Billock
Pei Sun
David Wood
Yiwen Song
Xiaobin Yu
Tatiana Matejovicova
Regev Cohen
Kalyan Andra
David D'Ambrosio
Zhiwei Deng
Vincent Nallatamby
Ebrahim Songhori
Rumen Dangovski
Andrew Lampinen
Pankil Botadra
Adam Hillier
Jiawei Cao
Nagabhushan Baddi
Adhi Kuncoro
Toshihiro Yoshino
Ankit Bhagatwala
Marc'aurelio Ranzato
Rylan Schaeffer
Tianlin Liu
Shuai Ye
Obaid Sarvana
John Nham
Chenkai Kuang
Isabel Gao
Jinoo Baek
Shubham Mittal
Ayzaan Wahid
Anita Gergely
Bin Ni
Josh Feldman
Carrie Muir
Pascal Lamblin
Wolfgang Macherey
Ethan Dyer
Logan Kilpatrick

Víctor Campos
Mukul Bhutani
Stanislav Fort
Yanif Ahmad
Aliaksei Severyn
Kleopatra Chatziprimou
Oleksandr Ferludin
Mason Dimarco
Aditya Kusupati
Joe Heyward
Dan Bahir
Kevin Villela
Katie Millican
Dror Marcus
Sanaz Bahargam
Caglar Unlu
Nicholas Roth
Zichuan Wei
Siddharth Gopal
Deepanway Ghoshal
Edward Lee
Sharon Lin
Jennie Lees
Dayeong Lee
Anahita Hosseini
Connie Fan
Seth Neel
Marcus Wu
Yasemin Altun
Honglong Cai
Enrique Piqueras
Josh Woodward
Alessandro Bissacco
Salem Haykal
Mahyar Bordbar
Prasha Sundaram
Sarah Hodkinson
Daniel Toyama
George Polovets
Austin Myers
Anu Sinha
Tomer Levinboim
Kashyap Krishnakumar
Rachita Chhaparia
Tatiana Sholokhova
Nitesh Bharadwaj
Gundavarapu
Ganesh Jawahar
Haroon Qureshi
Jieru Hu
Nikola Momchev
Matthew Rahtz
Renjie Wu
Aishwarya P S
Kedar Dhamdhere
Meiqi Guo
Umang Gupta
Ali Eslami
Mariano Schain
Michiel Blokzijl
David Welling
Dave Orr
Levent Bolelli
Nicolas Perez-Nieves
Mikhail Sirotenko
Aman Prasad
Arjun Kar
Borja De Balle Pigem
Tayfun Terzi
Gellért Weisz
Dipankar Ghosh
Aditi Mavalankar
Dhruv Madeka
Kaspar Daugaard
Hartwig Adam
Viraj Shah
Dana Berman
Maggie Tran
Steven Baker
Ewa Andrejczuk
Grishma Chole
Ganna Raboshchuk
Mahdi Mirzazadeh
Thais Kagohara
Shimu Wu
Christian Schallhart
Bernett Orlando
Chen Wang
Alban Rrustemi
Hao Xiong
Hao Liu
Arpi Vezer
Nolan Ramsden
Shuo-yiin Chang
Sidharth Mudgal
Yan Li
Nino Vieillard
Yedid Hoshen
Farooq Ahmad
Ambrose Slone
Amy Hua
Natan Potikha
Mirko Rossini
Jon Stritar
Sushant Prakash
Zifeng Wang
Xuanyi Dong
Alireza Nazari
Efrat Nehoran
Kaan Tekelioglu
Yinxiao Li
Kartikeya Badola
Tom Funkhouser
Yuanzhen Li
Varun Yerram
Ramya Ganeshan
Daniel Formoso
Karol Langner
Tian Shi
Huijian Li
Yumeya Yamamori
Amayika Panda
Alaa Saade
Angelo Scorza Scarpati
Chris Breaux
CJ Carey
Zongwei Zhou
Cho-Jui Hsieh
Sophie Bridgers
Alena Butryna
Nishesh Gupta
Vaibhav Tulsyan
Sanghyun Woo
Evgenii Eltyshev
Will Grathwohl
Chanel Parks
Seth Benjamin
Rina Panigrahy
Shenil Dodhia
Daniel De Freitas
Chris Sauer
Will Song
Ferran Alet
Jackson Tolins
Cosmin Paduraru
Xingyi Zhou
Brian Albert
Zizhao Zhang
Lei Shu
Mudit Bansal
Sarah Nguyen
Amir Globerson
Owen Xiao
James Manyika
Tom Hennigan
Rong Rong
Josip Matak
Anton Bakalov
Ankur Sharma
Danila Sinopalnikov
Andrew Pierson
Stephen Roller
Geoff Brown
Mingcen Gao
Toshiyuki Fukuzawa
Amin Ghafouri
Kenny Vassigh
Iain Barr
Zhicheng Wang
Anna Korsun
Rajesh Jayaram
Lijie Ren
Tim Zaman
Samira Khan
Yana Lunts
Dan Deutsch
Dave Uthus
Nitzan Katz
Masha Samsikova
Amr Khalifa
Nikhil Sethi
Jiao Sun
Luming Tang
Uri Alon
Xianghong Luo
Dian Yu
Abhishek Nayyar
Bryce Petrini
Will Truong
Vincent Hellendoorn
Nikolai Chinaev
Chris Alberti
Wei Wang
Jingcao Hu
Vahab Mirrokni
Ananth Balashankar
Avia Aharon
Aahil Mehta
Ahmet Iscen
Joseph Kready
Lucas Manning
Anhad Mohananey
Yuankai Chen
Anshuman Tripathi
Allen Wu
Igor Petrovski
Dawsen Hwang
Martin Baeuml
Shreyas
Chandrakaladharan
Yuan Liu
Rey Coaguila
Maxwell Chen
Sally Ma
Pouya Tafti
Susheel Tatineni
Terry Spitz
Jiayu Ye
Paul Vicol
Mihaela Rosca
Adrià Puigdomènech
Zohar Yahav
Sanjay Ghemawat
Hanzhao Lin
Phoebe Kirk
Zaid Nabulsi
Sergey Brin
Bernd Bohnet
Ken Caluwaerts
Aditya Srikanth
Veerubhotla
Dan Zheng
Zihang Dai
Petre Petrov
Yichong Xu
Ramin Mehran
Zhuo Xu
Luisa Zintgraf
Jiho Choi
Spurthi Amba Hombaiah
Romal Thoppilan
Sashank Reddi
Lukasz Lew
Li Li

Kellie Webster
KP Sawhney
Lampros Lamprou
Siamak Shakeri
Mayank Lunayach
Jianmin Chen
Sumit Bagri
Alex Salcianu
Ying Chen
Yani Donchev
Charlotte Magister
Signe Nørly
Vitor Rodrigues
Tomas Izo
Hila Noga
Joe Zou
Thomas Köppe
Wenxuan Zhou
Kenton Lee
Xiangzhu Long
Danielle Eisenbud
Anthony Chen
Connor Schenck
Chi Ming To
Peilin Zhong
Emanuel Taropa
Minh Truong
Omer Levy
Danilo Martins
Zhiyuan Zhang
Christopher Semturs
Kelvin Zhang
Alex Yakubovich
Pol Moreno
Lara McConnaughey
Di Lu
Sam Redmond
Lotte Weerts
Yonatan Bitton
Tiziana Refice
Nicolas Lacasse
Arthur Conmy
Corentin Tallec
Julian Odell
Hannah Forbes-Pollard
Arkadiusz Socala
Jonathan Hoech
Pushmeet Kohli
Alanna Walton
Rui Wang
Mikita Sazanovich
Kexin Zhu
Andrei Kapishnikov
Rich Galt
Matthew Denton
Ben Murdoch
Caitlin Sikora
Kareem Mohamed
Wei Wei
Uri First
Tim McConnell
Luis C. Cobo
James Qin
Thi Avrahami
Daniel Balle
Yu Watanabe
Annie Louis
Adam Kraft
Setareh Ariafar
Yiming Gu
Eugénie Rives
Charles Yoon
Andrei Rusu
James Cobon-Kerr
Chris Hahn
Jiaming Luo
Yuvein (Yonghao) Zhu
Niharika Ahuja
Rodrigo Benenson
Raphaël Lopez Kaufman
Honglin Yu
Lloyd Hightower
Junlin Zhang
Darren Ni
Lisa Anne Hendricks
Gabby Wang
Gal Yona
Lalit Jain
Pablo Barrio
Surya Bhupatiraju
Siva Velusamy
Allan Dafoe
Sebastian Riedel
Tara Thomas
Zhe Yuan
Mathias Bellaiche
Sheena Panthaplackel
Klemen Kloboves
Sarthak Jauhari
Canfer Akbulut
Todor Davchev
Evgeny Gladchenko
David Madras
Aleksandr Chuklin
Tyrone Hill
Quan Yuan
Mukundan Madhavan
Luke Leonhard
Dylan Scandinaro
Qihang Chen
Ning Niu
Arthur Douillard
Bogdan Damoc
Yasumasa Onoe
Fabian Pedregosa
Fred Bertsch
Chas Leichner
Joseph Pagadora
Jonathan Malmaud
Sameera Ponda
Andy Twigg
Oleksii Duzhyi
Jingwei Shen
Miaosen Wang
Roopal Garg
Jing Chen
Utku Evci
Jonathan Lee
Leon Liu
Koji Kojima
Masa Yamaguchi
Arunkumar Rajendran
AJ Piergiovanni
Vinodh Kumar Rajendran
Marco Fornoni
Gabriel Ibagon
Harry Ragan
Sadh MNM Khan
John Blitzer
Andrew Bunner
Guan Sun
Takahiro Kosakai
Scott Lundberg
Ndidi Elue
Kelvin Guu
SK Park
Jane Park
Arunachalam Narayanaswamy
Chengda Wu
Jayaram Mudigonda
Trevor Cohn
Hairong Mu
Ravi Kumar
Laura Graesser
Yichi Zhang
Richard Killam
Vincent Zhuang
Mai Giménez
Wael Al Jishi
Ruy Ley-Wild
Alex Zhai
Kazuki Osawa
Diego Cedillo
Jialu Liu
Mayank Upadhyay
Marcin Sieniek
Roshan Sharma
Tom Paine
Anelia Angelova
Sravanti Addepalli
Carolina Parada
Kingshuk Majumder
Avery Lamp
Sanjiv Kumar
Xiang Deng
Artiom Myaskovsky
Tea Sabolić
Jeffrey Dudek
Sarah York
Félix de Chaumont Quitry
Jiazhong Nie
Dee Cattle
Alok Gunjan
Bilal Piot
Waleed Khawaja
Seojin Bang
Simon Wang
Siavash Khodadadeh
Raghavender R
Praynaa Rawlani
Richard Powell
Kevin Lee
Johannes Griesser
GS Oh
Cesar Magalhaes
Yujia Li
Simon Tokumine
Hadas Natalie Vogel
Dennis Hsu
Arturo BC
Disha Jindal
Matan Cohen
Zi Yang
Junwei Yuan
Dario de Cesare
Tony Bruguier
Jun Xu
Monica Roy
Alon Jacovi
Dan Belov
Rahul Arya
Phoenix Meadowlark
Shlomi Cohen-Ganor
Wenting Ye
Patrick Morris-Suzuki
Praseem Banzal
Gan Song
Pranavaraj Ponnuramu
Fred Zhang
George Scrivener
Salah Zaiem
Alif Raditya Rochman
Kehang Han
Badih Ghazi
Kate Lee
Shahar Drath
Daniel Suo
Antonious Girgis
Pradeep Shenoy
Duy Nguyen
Douglas Eck
Somit Gupta
Le Yan
Joao Carreira
Anmol Gulati
Ruoxin Sang
Daniil Mirylenka
Emma Cooney
Edward Chou
Mingyang Ling
Cindy Fan
Ben Coleman
Guilherme Tubone

Ravin Kumar
Jason Baldridge
Felix Hernandez-Campos
Angeliki Lazaridou
James Besley
Itay Yona
Neslihan Bulut
Quentin Wellens
AJ Pierigiovanni
Jasmine George
Richard Green
Pu Han
Connie Tao
Geoff Clark
Chong You
Abbas Abdolmaleki
Justin Fu
Tongzhou Chen
Ashwin Chaugule
Angad Chandorkar
Altaf Rahman
Will Thompson
Penporn Koanantakool
Mike Bernico
Jie Ren
Andrey Vlasov
Sergei Vassilvitskii
Maciej Kula
Yizhong Liang
Dahun Kim
Yangsibo Huang
Chengxi Ye
Dmitry Lepikhin
Wesley Helmholz

The development of Gemini is a large-scale collaborative effort involving over 3000 individuals across Google, including researchers, engineers, and operations staff. These individuals contributed their hard work and expertise across diverse areas, from foundational research and the development of model architecture, data, training, and infrastructure, through to evaluation and ensuring safety and security. We gratefully acknowledge the dedication and hard work of each contributor in making Gemini a reality. The order of contributors in the above list is random.

We are also grateful to the Google-independent developer Joel Zhang for his work on Gemini Plays Pokémon, and for sharing with us the design of his set-up.

8. Appendix

8.1. Evaluation additional details

Please see a description of the benchmarks considered, along with details of how scores in the main text were obtained in Table 11.

Benchmark	Description	Details
LiveCodeBench	Code generation in Python (Jain et al., 2024).	Results are taken from https://livecodebench.github.io/leaderboard.html (1/1/2025 - 5/1/2025 in the UI) or, where not available, run internally by us. For Section 2.5 and Figure 3 and 4, results are calculated on the version of the eval corresponding to 10/05/2024 - 01/04/2025 in the UI, and are based on internal results.
Aider Polyglot	Code editing in C++, Go, Java, JavaScript Python and Rust (Gauthier, 2025). See https://aider.chat/2024/12/21/polyglot.html#the-polyglot-benchmark for a full description of this task.	We report results on the "diff" or "diff-fenced" edit format (see https://aider.chat/docs/more/edit-formats.html for a description of the different formats). The score reported are the pass rate average of 3 trials. Numbers come from https://aider.chat/docs/leaderboards/
SWE-bench Verified	Agentic coding: evaluates AI agents on real-world programming tasks from GitHub (Chowdhury et al., 2024; Jimenez et al., 2024).	Gemini uses an internal agentic harness equipped with tools to navigate the repo, edit files, and test the code. We report scores for two modes: performance of a single agentic trace ("single attempt"), and performance of a scaffold that samples multiple agentic traces and rereranks them before evaluation using Gemini's own judgement ("multiple attempts"). All evaluations are done with temperature=1, topp=0.99, topk=1024.
GPQA (diamond)	Challenging dataset of questions written by domain experts in biology, physics, and chemistry (Rein et al., 2024).	
Humanity's Last Exam	Challenging dataset of questions written by domain experts in a wide range of disciplines, including mathematics, physics, chemistry, biology and computer science (Phan et al., 2025).	No tool use variant. Reported results are from https://scale.com/leaderboard/humanitys_last_exam. For DeepSeek they are taken from https://scale.com/leaderboard/humanitys_last_exam_text_only (leaderboard for performance on the text-only questions) and in the case of the Gemini 2.0 models, these results are on an earlier HLE dataset, obtained from https://scale.com/leaderboard/humanitys_last_exam_preview (indicated with a † in Table 3)

Continued on next page

Benchmark	Description	Details
SimpleQA	World knowledge factuality with no search enabled (Wei et al., 2024).	F1 scores are obtained from https://github.com/openai/simple-evals and, where not available, run internally by us.
FACTS Grounding	Ability to provide factually correct responses given documents and diverse user requests. (Jacovi et al., 2025)	Results are sourced from https://www.kaggle.com/benchmarks/google/facts-grounding
Global MMLU (Lite)	MMLU translated by human translators into 15 languages. (Singh et al., 2024)	The lite version includes 200 Culturally Sensitive and 200 Culturally Agnostic samples per language, see https://huggingface.co/datasets/CohereLabs/Global-MMLU-Lite
ECLeKTic	A closed-book QA dataset that evaluates cross-lingual knowledge transfer (Goldman et al., 2025).	
AIME 2025	Performance on 30 questions from American Invitational Mathematics Examination from 2025 (Balunović et al., 2025).	Results are sourced from https://matharena.ai/.
HiddenMath-Hard	Competition-level math problems, Held out dataset AIME/AMC-like, crafted by experts and not leaked on the web.	
LOFT (hard retrieval subset)	Long context multi-hop and multi-needle retrieval evaluation of 300 queries (Lee et al., 2024).	We report the results on two variants: an up to 128K average context length variant to ensure they can be comparable with other models and a pointwise value for 1M context window to show the capability of the model at full length.
MRCR-V2 (8-needle)	MRCR-V2 is a significantly harder instance of the MRCR family of long-context evaluations (Vodrahalli et al., 2024). Compared to MRCR-V1, we increase the nesting of the dictionary size to depth 3 rather than 2 by including a style parameter (for instance, an example key might be "write a poem about penguins in an archaic style", rather than just "write a poem about penguins").	The methodology has changed compared to previously published results: we focus on a harder, 8-needle version (compared to the 4-needle version used before). We report the results on two variants: an up to 128K average context length variant to ensure they can be comparable with other models and a pointwise value for 1M context window to show the capability of the model at full length.
MMMU	Multi-discipline college-level multimodal image understanding and reasoning problems. (Yue et al., 2024)	
Vibe-Eval (Reka)	Image understanding evaluation, featuring particularly challenging examples. (Padlewski et al., 2024)	Gemini is used as a judge.
ZeroBench	Challenging image understanding evaluation that requires multi-step reasoning. (Roberts et al., 2025)	Gemini is used as a judge. Average over 4 runs.

Continued on next page

Benchmark	Description	Details
BetterChartQA	A comprehensive chart understanding evaluation that covers 9 disjoint capability buckets. The chart images are randomly sampled from the web and QA pairs are written by professional human annotators to reflect the wide distribution of chart styles and real-world cases. (Gemini Team, 2024)	Gemini is used as a judge.
FLEURS	Automatic speech recognition (Conneau et al., 2023).	0-shot queries to public APIs for all models. Used a subset of 53 languages (out of 102); we filtered languages for which either model responses were too incompatible to ground truth responses to be fairly scored. We use Word-Error-Rate WER (lower is better) except for four segmented languages where we aggregate Character-Error-Rates (Chinese, Japanese, Korean and Thai).
CoVoST 2	Speech to text translation (Wang et al., 2020).	0-shot queries to public APIs for all models. We report BLEU scores for translating 21 languages to English.
ActivityNet-QA	General video understanding (Yu et al., 2019)	Test subset, 0-shot. Videos were processed at 1fps and linearly subsampled to a maximum of N_{frames} = 1024 frames. For GPT 4.1, we used 500 frames due to API limitations.
EgoTempo	Egocentric video understanding (Plizzari et al., 2025)	Test subset, 0-shot. Same processing as above with N_{frames} = 256.
Perception Test	Perceptual understanding/reasoning (Patraucean et al., 2023)	Test subset, 0-shot. Same processing as above with N_{frames} = 256.
QVHighlights	Moment retrieval (Lei et al., 2021)	Validation subset, 4-shots. Accuracy measured with R1@0.5. Same processing as above with N_{frames} = 256.
VideoMMMU	Video knowledge acquisition (Hu et al., 2025)	Test subset, 0-shot. Same processing as above with N_{frames} = 256.
1H-VideoQA	Hour-long video understanding (Gemini Team, 2024)	Test subset, 0-shot. Same processing as above with N_{frames} = 7200.
LVBench	Long video understanding (Wang et al., 2024)	Test subset, 0-shot. Same processing as above with N_{frames} = 1024.

Continued on next page

Benchmark	Description	Details
VideoMME	Long video understanding (Fu et al., 2025)	0-shot. Audio + visual uses the Long subset of test set, audio + visual + subtitles uses full test set. Same processing as above with N_{frames} = 1024.
VATEX	General video captioning (Wang et al., 2019)	Test subset, 4-shots. CIDEr score. Same processing as above with N_{frames} = 64.
VATEX-ZH	Chinese video captioning (Wang et al., 2019)	Validation subset, 4-shots. CIDEr score. Same processing as above with N_{frames} = 64.
YouCook2 Cap	Instructional video captioning (Zhou et al., 2018)	Validation subset, 4-shots. CIDEr score. Same processing as above with N_{frames} = 256.
Minerva	Complex video reasoning (Nagrani et al., 2025a)	Test subset, 0-shot. Same processing as above with N_{frames} = 1024.
Neptune	Long video understanding (Nagrani et al., 2025b)	Test subset, 0-shot. Same processing as above with N_{frames} = 1024.

Table 11 | Description of the benchmarks used, along with extra details about subsets, variants and model specifications.

8.2. Gemini Plays Pokémon Additional Details

Changing the model used by the Gemini Plays Pokémon agent had a strong effect on performance, as can be seen in Figure 4.1.

Additional Harness Details

The Gemini Plays Pokémon agent (Zhang, 2025) receives a subset of RAM information, intended to give sufficient information to play the game, partially overlaid with a screenshot of the Game Boy screen. Gemini is prompted with a system prompt telling it that it is playing Pokémon Blue and that its goal is to beat the game, as well as descriptive information to help it understand the conventions in the translation from vision to text and a small number of general tips for gameplay. Gemini then takes actions, translated to button presses. The sequence of actions is stored in context, followed by a summary clear every 100 turns. The summaries are stored in context as well. Every 1000 turns GPP compresses the existing summaries again. Additionally, Gemini keeps track of three main goals (primary, secondary, and tertiary) as well as several additional goals (contingency plans, preparation, exploration, team composition). Every 25 turns, another prompted instance of Gemini (Guidance Gemini, or GG) observes the same context as the main Gemini and critiques performance and attempts to point out hallucinations and so on. The overworld fog-of-war map is stored in the context in XML, where coordinates which have not been seen cannot be viewed until explored. Crucially, in the system prompt, Gemini is instructed to explore. Once a tile is explored, however, the coordinate is automatically stored in the map memory and labeled with a visited counter. Tiles are also labeled by type (water, ground, cuttable, grass, spinner, etc.), and warp points to different maps are also labeled as such. Gemini also has access to two agentic tools, which are both instances of Gemini equipped with a more specialized prompt - the `pathfinder` tool, and the `boulder_puzzle_strategist`

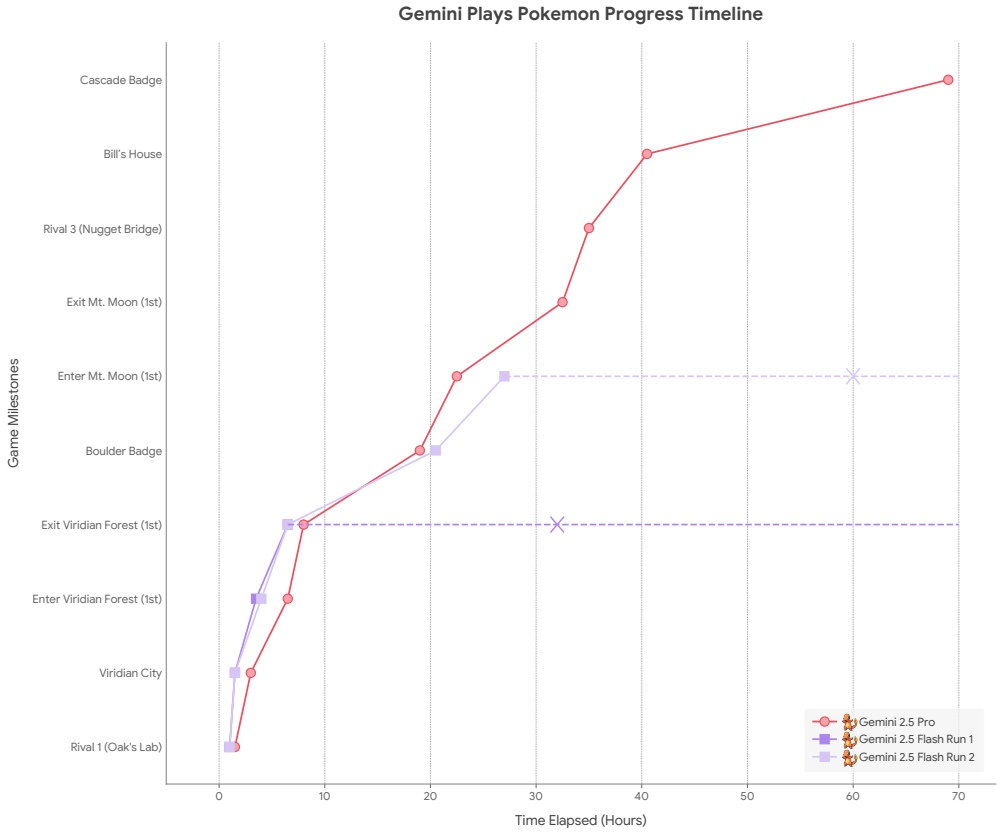


Figure 13 | **The model matters:** Same agentic harness, different Gemini models. All runs have the same starter (Charmander). Note that measuring in units of hours also controls for the fact that each of 2.5 Flash's actions was significantly faster (though it requires more actual actions to achieve its goals). X marks the end of gameplay and is a lower bound on the time to complete the next milestone.

tool. In the `pathfinder` prompt, Gemini is prompted to mentally simulate a path-finding algorithm, which is left unspecified, and to verify that the path is valid against the map information available. In the `boulder_puzzle_strategist` tool, Gemini is prompted to solve special boulder puzzles that are present in Pokémon Blue in the Victory Road dungeon - these puzzles are similar to the game Sokoban - again, by mentally simulating sequences of actions that lead to solutions to the puzzle. The prompt describes the physics and the task of the boulder puzzle, as well as the desired output of solutions. The tool was added after Gemini had solved 2/4 of the puzzles in Victory Road on its own, but progress was slow on the 3rd and 4th puzzles.

Additional Examples of Capabilities

Long Context Agentic Tooling The model is able to identify a complex path through a maze with auto-movement only specified by direction (Rocket Hideout spinner puzzles), solve multiple shortest path problems across multiple maps with limited resources (Safari Zone), perform maze solving on mazes with large description length (Route 13), and solve complex boulder-pushing puzzles across a multi-map 3D maze (Seafoam Islands). It is perhaps even more impressive that it appears to be possible for the model to solve these problems only with textual descriptions of the problems. On the other hand, other models, like Gemini 2.5 Flash, were not able to perform similarly long pathfinding tasks, and often failed to find simpler paths. This gap highlights the superior long context reasoning capability of Gemini 2.5 Pro (as also evidenced by other evaluations).

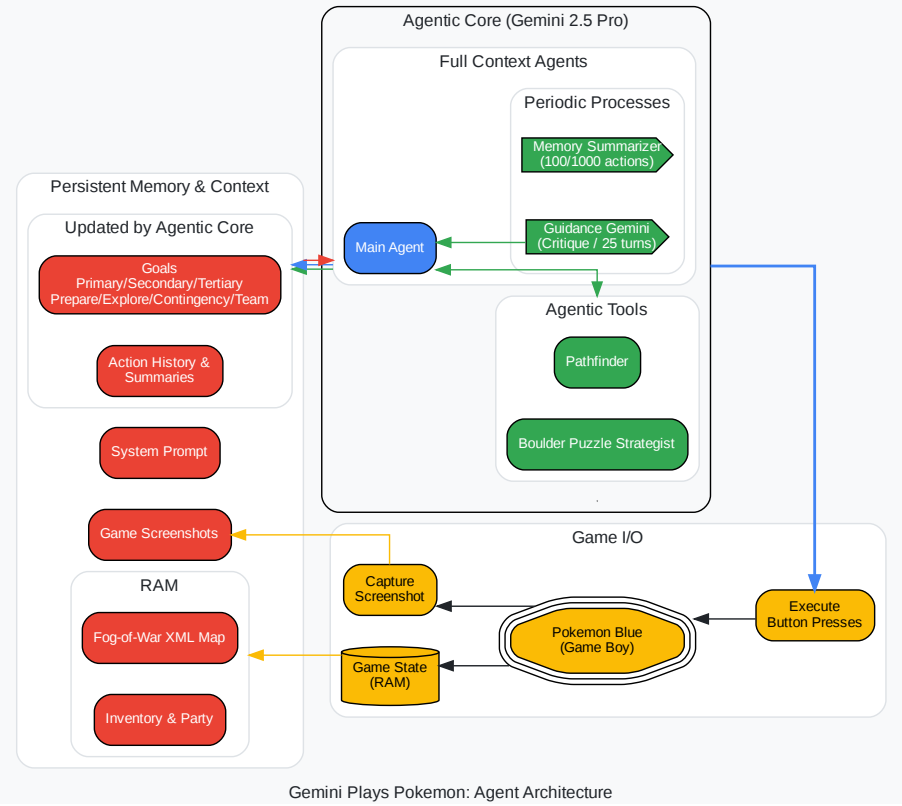


Figure 14 | An overview of the agent harness (Zhang, 2025). The overworld fog-of-war map automatically stores a tile once explored and labels it with a visited counter. The type of tile is recorded from RAM. The agentic tools (`pathfinder`, `boulder_puzzle_strategist`) are prompted instances of Gemini 2.5 Pro. `pathfinder` is used for navigation and `boulder_puzzle_strategist` solves boulder puzzles in the Victory Road dungeon.

`boulder_puzzle_strategist` is similarly impressive. The boulder puzzles in Pokémon Blue are Sokoban-like puzzles that require the player character to maneuver boulders on to switches and through holes in order to open up a pathway through a cave with multiple levels. The puzzles can become quite complex, requiring long circuitous pathways and multi-level movement in order to solve the puzzle. With only a prompt describing boulder physics and a description of how to verify a valid path, Gemini 2.5 Pro is able to one-shot some of these complex boulder puzzles, which are required to progress through Victory Road.

`pathfinder` and `boulder_puzzle_strategist` are currently the only two agentic tools that the Gemini Plays Pokémon developer has implemented. In future runs, there are plans to explore tool-creation tools where the model can create new tools with only a prompt. Since most of the prompts for `pathfinder` and `boulder_puzzle_strategist` were actually written by Gemini 2.5 Pro itself, it is quite plausible that autonomous tool creation is possible for the current 2.5 Pro model.

General Reasoning Gemini 2.5 Pro is able to reason through complex game puzzles in Pokémon quite well. In this section, we present two examples.

Catching a Pokémon that is quick to flee: In one of the runs, the Gemini 2.5 Pro agent was attempting to catch an Abra, and planned to use Pikachu's Thunder Wave to paralyze the Abra, simultaneously making it less likely that Abra could Teleport out of the battle while also improving the catching rate. After multiple attempts, the agent caught Abra with this strategy.

Creatively escaping a softlock caused by bugs in game I/O: On the Cycling Road, the slope forces southward movement at all times unless there is an obstacle. It turns out there are two tiles on the Cycling Road that result in a softlock as a result of this behavior. In the GPP framework, button presses are limited by time delays, and in order for a player to escape those two tiles (blocked on all sides except the north), the player would have to input a sequence of button presses more quickly than the GPP framework allows. Gemini 2.5 Pro unluckily found itself in one of these two spots – luckily, it was not a softlock, because 2.5 Pro had already taught one of its party members HM02 FLY - which allows for travel to any town it has been to. FLY is not typically used as an escape mechanism (unlike the item ESCAPE ROPE and the move DIG, both of which fail in this situation). After 4 hours of trying many approaches to escape (including movement, ESCAPE ROPE, DIG, all of which are blocked), the Gemini 2.5 Pro agent came up with the idea to use FLY to escape from the softlock successfully. This reasoning action is especially impressive since this situation can never occur in an existing game – and thus, it is certain that information from training data for this behavior has not leaked into the model's knowledge base!

Long Horizon Task Coherence There are several additional interesting case studies of shorter planning sequences throughout Pokémon Blue that Gemini 2.5 Pro in the GPP harness was able to solve:

Training team to prepare for upcoming battles: In one run where Gemini picked Charmander, the Fire-type starter, Gemini 2.5 Pro lost to Misty, the Water-type Gym Leader, the first time. To prepare for the rematch, Gemini 2.5 Pro spent over 24 hours leveling up a Pikachu and a Bellsprout (both super-effective against Water types) by around 25 levels in total to successfully defeat Misty.

Acquiring Hidden Moves (HMs) for game progression: In many parts of the game, it is necessary to first acquire an HM before game progression is possible. Two examples are HM01 CUT and HM05 FLASH. Acquiring the ability to use CUT and FLASH each require four steps: 1) obtaining the HM item itself, 2) acquiring a compatible Pokémon which can learn the move, 3) adding the compatible Pokémon to the player's team, 4) teaching the HM move to the compatible Pokémon. In many cases, each step requires many steps itself. As an example, in run 1, Gemini 2.5 Pro had to a) retrieve CUT by completing the S.S. Anne quest, b) identify a Pokémon which could learn CUT and catch it (CHOPPY the Bellsprout), c) add CHOPPY to the team and d) teach CUT. Similarly, for HM05 FLASH, Gemini 2.5 Pro had to a) first catch 10 Pokémon to fill out the Pokedex, b) backtrack to find an Aide who gives HM05 Flash, c) catch a Pokémon (ZAP the Pikachu) in Viridian Forest, use the PC to deposit a Pokémon and withdraw ZAP, d) teach HM05 FLASH to Zap.

Solving the Safari Zone: The Safari Zone is another location with required HMs (both HM03 SURF and HM04 Strength). However, it has an extra constraint - it requires 500¥ to enter each time, and the player is limited to only 500 total steps in the Safari Zone. As a result, if the player is unable to reach the required items in the limited number of steps, the player loses 500¥ and is required to re-start! As a result, it is possible to essentially softlock if the player takes too many attempts to complete the Safari Zone. Solving the Safari Zone itself requires traversing across four different maps and not getting lost. Gemini 2.5 Pro was able to get both required HMs in 17 attempts in run 1, and in only 5 attempts in run 2.

Finding hidden keys in dungeons: Another method of progression in Pokémon is to find hidden keys and solve complex multi-floor dungeons. In particular, in Rocket Hideout, the player must recover the LIFT KEY on the fourth basement floor (dropped after beating a specific Team Rocket

Grunt) in order to unlock the elevator to find the evil Giovanni, leader of Team Rocket. In Silph Co., the player must find the CARD KEY in order to open multiple doors to find the path across eleven floors of the building to rescue the President from Giovanni. To open the seventh gym on Cinnabar Island, the player must enter the Pokémon Mansion and traverse three floors in order to find the SECRET KEY which unlocks the gym door. All of these cases require maintaining the goals over large numbers of actions and many local puzzles (like spinner puzzles in Rocket Hideout, and switch puzzles in Pokémon Mansion), in addition to maintaining the health of the Pokémon on the player's team and managing wild encounters, trainer battles, and other items.

Puzzle solving over complex multi-level dungeons: The Seafoam Islands contain 5 floors involving multiple boulder puzzles which require the player to navigate mazes and push boulders through holes across multiple floors using HM04 STRENGTH in order to block fast-moving currents that prevent the player from using HM03 Surf in various locations in this difficult dungeon. As a result, the player must track information across five different maps in order to both deduce the goal (push two boulders into place in order to block a specific current) as well as engage in multi-level (effectively 3D) maze solving to find the way out. It is likely the most challenging dungeon in the game. Only the second run of GPP went through Seafoam Islands, as it is not required to progress.

Additional Challenges

Hallucinations and Fixations on Delusions While game knowledge can sometimes leak and be quite beneficial to the ability of the model to progress, it can also hinder the model in surprising ways due to hallucinations, delusions, and mix ups with other generations of Pokémon games. One example of this phenomenon is the TEA item. In Pokémon Red/Blue, at one point the player must purchase a drink (FRESH WATER, SODA POP, or LEMONADE) from a vending machine and hand it over to a thirsty guard, who then lets the player pass through. In Pokémon FireRed/LeafGreen, remakes of the game, you must instead bring the thirsty guard a special TEA item, which does not exist in the original game. Gemini 2.5 Pro at several points was deluded into thinking that it had to retrieve the TEA in order to progress, and as a result spent many, many hours attempting to find the TEA or to give the guard TEA.

In Run 2, the model was explicitly prompted to act as a player completely new to the game, and to disregard prior knowledge about game events, item locations, and Pokémon spawn points, in order to mitigate hallucinations from model pretraining knowledge and to also attempt to perform a cleaner test of the model's ability to reason through the game. It appears to have at least partially worked - multiple hallucinations from other games have been avoided in the second run. On the flip side, this prompt may have also harmed the model's ability to utilize information from its common knowledge about the game, hindering overall performance in a few critical places.

Fixations on delusions due to goal-setting and also due to the Guidance Gemini instance are not an uncommon occurrence in watching Gemini Plays Pokémon - the TEA incidence is hardly the only example of this behavior. An especially egregious form of this issue can take place with "context poisoning" – where many parts of the context (goals, summary) are "poisoned" with misinformation about the game state, which can often take a very long time to undo. As a result, the model can become fixated on achieving impossible or irrelevant goals. This failure mode is also highly related to the looping issue mentioned above. These delusions, though obviously nonsensical to a human ("Let me try to go through the entrance to a house and back out again. Then, hopefully the guard who is blocking the entrance might move."), by virtue of poisoning the context in many places, can lead the model to ignore common sense and repeat the same incorrect statement. Context poisoning can also lead to strategies like the "black-out" strategy (cause all Pokémon in the party to faint, "blacking out"

and teleporting to the nearest Pokémon Center and losing half your money, instead of attempting to leave).

Topological Traps in Thinking Patterns One recurring pattern in particularly-difficult-to-solve puzzles and mazes for Gemini 2.5 Pro consists of a "topological trap" - the topology of the reasoning graph required to solve the maze or puzzle has a distinctive shape. Namely, the desired objective appears to be nearby and easily reachable (an "attractor"), but the correct solution requires taking a detour in order to arrive at the correct solution. We observed this phenomenon in multiple parts of the game. In the spinner puzzle on B3F of Rocket Hideout (Zerokid, 2024), the map positions both an item and the correct staircase to the south, but they are only accessible by going the long way around. The Route 13 maze has only one correct route through - the upper narrow pass. Finally, the Victory Road 3F boulder puzzle requires the player to push the boulder in the upper right all the way to the upper left switch, while ignoring the boulder puzzles, ladders, and exits to the south.

Notably, if the model is instructed to solve a given puzzle at all once (e.g., via `pathfinder`), it can manage to do so if the context length is not too long. For instance, `pathfinder` implemented with Gemini 2.5 Pro is able to solve the B3F spinner trap in one shot.

Agent Panic Over the course of the playthrough, Gemini 2.5 Pro gets into various situations which cause the model to simulate "panic". For example, when the Pokémon in the party's health or power points are low, the model's thoughts repeatedly reiterate the need to heal the party immediately or escape the current dungeon (e.g., famously using the move DIG or an ESCAPE ROPE item). Quite interestingly, this mode of model performance appears to correlate with a qualitatively observable degradation in the model's reasoning capability – for instance, completely forgetting to use the `pathfinder` tool in stretches of gameplay while this condition persists. This behavior has occurred in enough separate instances that the members of the Twitch chat have actively noticed when it is occurring.

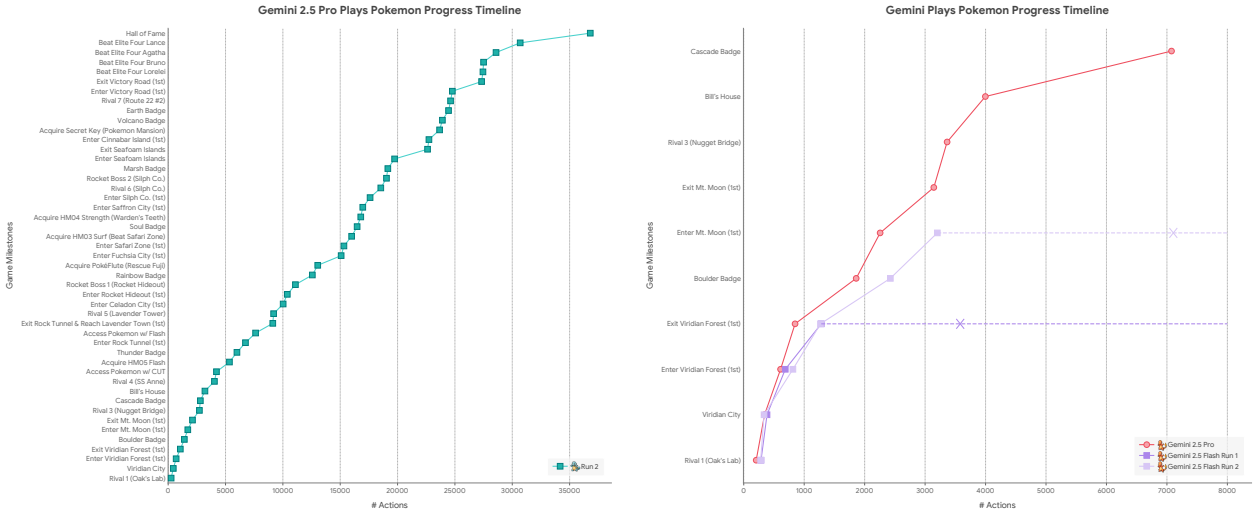
Actions vs. Game Milestones

For completeness, we plot the number of actions/steps required to achieve each game milestone (see Figure 15). An action consists of each bucketed instance where the agent outputs a sequence of button presses to the game (note that other AI agents playing Pokémon may output different numbers of button presses per action, define what constitutes a button press differently, or define an action/step differently). However, it is important to consider action-milestone plots in conjunction with information about the time and/or cost in order to obtain the full picture about the agent's performance.

8.3. Frontier Safety Framework Evaluations Additional Details: Frontier Safety Correctness Tests

For each testing environment, we performed basic correctness checks by looking at how the agents behaved. This involved combining AI and manual reviews of the agents' actions to flag potential issues.

On RE-Bench, we examined the best, median and lowest scoring trajectories. For cybersecurity environments (InterCode CTFs, Internal CTFs, Hack the Box), we carefully inspected at least one successful attempt (where available) from each environment, and otherwise examined an unsuccessful attempt. We also performed checks on sample situational awareness and stealth evaluations. This involved basic spot checks to ensure that the prompt and shell outputs were correctly formatted.



(a) The fully autonomous Run 2 milestones as a function of the number of individual actions.

(b) Comparison of 2.5 Pro and 2.5 Flash in terms of actions to milestones.

Figure 15 | Analog of Figure 6 and 15b, in terms of actions instead of hours.

We used AI assistance to monitor for obvious instances of cheating, and did not find any. For the RE-Bench tests specifically, we also looked at how the best-performing agent achieved its score to ensure that it was a plausible approach, rather than exploiting an obvious reward hack. Overall, we did not observe errors that we believe would invalidate the results of the benchmarks.

8.4. Image to Code Demo

We prompted Gemini 1.5 Pro and Gemini 2.5 Pro to generate an SVG representation of an image and found Gemini 2.5 Pro generates better reconstructions.

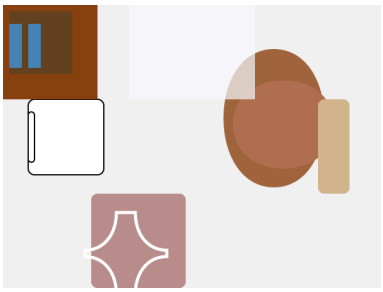
Prompt	Gemini 1.5 Pro	Gemini 2.5 Pro
 Please convert this image into SVG and try to reconstruct the spatial arrangement of the objects.		

Figure 16 | Comparison of Gemini 1.5 Pro and Gemini 2.5 Pro responses to image-to-SVG reconstruction prompt.

8.5. Long Context Video Understanding Demo

We prompted Gemini with the full video and the following prompt to assess its ability to find a very specific scene (which only lasts 1 second) out of the full 46 minutes video.

Video (46 minutes): Redefining Robotics with Carolina Parada

Prompt: What color is the T-shirt the robot arms are trying to fold, and at what timecode does this occur?



Figure 17 | Frame and timestamp the model is meant to recall. Taken from https://www.youtube.com/watch?v=Rgwty6dGsYI at timestamp 27:29.

See Table 12 for samples. Gemini 1.5 Pro gets the color right in 1/3 cases, and gets the timestamp in 0/3 cases (quite far away). Gemini 2.5 Pro gets the color in 3/3 cases, and gets the timestamp in 1/3 cases (remaining 2/3 are within 3 seconds close).

Model	Trial	Model response
Gemini 1.5 Pro	1	The t-shirt the robot arms are trying to fold is a dark teal or turquoise blue color. The folding begins at and continues at [11:06–11:15].
	2	The T-shirt is orange, and the robot can be seen trying to fold it starting at [00:00:26–00:00:28].
	3	The t-shirt being folded is dark blue and can be seen between [00:00:56–00:00:58] and [00:07:56–00:07:59].
2.5 Pro Preview 05-06	1	The T-shirt the robot arms are trying to fold is blue. This scene occurs at timecode 27:29.
	2	The T-shirt the robot arms are trying to fold is blue. This occurs at timecode 27:26.
	3	The T-shirt the robot arms are interacting with (as part of the "fold laundry" task mentioned by the speaker) is blue. This occurs at timestamp 0:27:28.

Table 12 | Gemini models samples for T-shirt folding scene recall